PitchBook

The $100 Billion Entry Fee

The Cost Architecture That Wall Street Isn't Pricing

Harrison Rolfes, Senior Research Analyst, PitchBook

September 9, 2026

v5.0

Anthropic and OpenAI on the 5-Layer Cost Stack, 2023A-2030E, with a greenfield entry-cost section. Institutional research; figures carry ledger row references (L-xxx) and tiers throughout.

Contents

§0 The entry fee, restated

Two numbers in this report will survive the S-1s unchanged. Anthropic has raised $124.3B of equity ($124,254M at the Series H close on May 28, 2026; estimated, summed from ten PitchBook equity rounds; L-007). OpenAI has raised $181.2B ($181,216.5M at the March 31, 2026 close; estimated, same method; L-022). Both exclude debt, grants and the lease structures that sit off both balance sheets; they are what the two companies have taken from shareholders to stand where they stand. Two other numbers will not survive: Anthropic's $65B annualized run-rate (end-July 2026, confirmed, gross of cloud-partner resale; L-032, CNBC and Bloomberg) and OpenAI's run-rate of more than $40B (July 2026 performance, confirmed, basis not stated; L-065, Bloomberg). Both are company-to-investor communications relayed by the press. Neither has been audited, and the audited versions will arrive on a basis the companies choose, in prospectuses that are not yet public (EDGAR, September 9, 2026: no S-1 filed by either registrant; confirmed, T1; L-001, L-002).

v4 of this report argued that the obligation stack exceeded cumulative burn by an order of magnitude. That framing is retired. It compared a spend plan with a burn plan, two objects of different scope and vintage, and §3 shows why the comparison fails. v5 says the entry fee has three prices, and that until now only the middle one has been counted.

The first price is the burn to a first positive operating result. For Anthropic that is roughly $12-20B (derived: a FY2024 net loss of −$5.3B or −$8.3B depending on the PitchBook vintage, C-05, L-012; FY2025 burn of $5.6B from a January 2026 plan, recalled, T4, L-120; and earlier years with no row). The range is wide because the inputs are weak; the figure is printed to be corrected by the S-1, not relied on.

The second price is the capital raised to get there: $124.3B and $181.2B.

The third price is the contracts signed to keep the seat. On the documented count, Anthropic has signed $324.1B of compute contracts with eight counterparties and OpenAI $480.4B with four (derived from the T1 and T2 rows in §3 plus one T3 row, Volta's $10B; press-valued items marked). Neither figure includes the gigawatts no filing prices: 5 GW of Google TPU capacity under contract, 10 GW more in Broadcom's line of sight, 2 GW of AMD silicon, and 8.75 GW-IT of Ohio and Texas leases signed by OpenAI at undisclosed rents. One figure exists for the Google leg, about $200B over five years (The Information via Reuters, May 5, 2026; estimated, T3; L-155), and it carries a [VERIFY] flag everywhere this report uses it: the page could not be opened, and Alphabet's own 10-Q for the June quarter mentions Anthropic zero times (confirmed, T1; L-153). If the reported figure is true, Anthropic's contract stack is $524.1B rather than $324.1B, and the Google leg alone is larger than everything else combined. The S-1 commitments note will be the first document to settle it.

Where the thesis moved: it got weaker at the operating line. In the second quarter of 2026 Anthropic reported positive adjusted operating income on quarterly revenue of $11.5-11.6B (preliminary, confirmed; L-033, L-151), up from $4,730M in the first quarter (confirmed; L-151). That is a public counterexample to "structurally incompatible with returns", and it is weaker than the phrase "operating profit" would suggest. The measure is non-GAAP. The company has not published its definition. The relays that describe it say it includes training cost and excludes stock-based compensation; the Wall Street Journal wrote that "the methodology behind the adjusted profitability measure is unclear"; and the $559M the press attached to it was a company projection on $10.9B of revenue, about 5.1% (Reuters, August 14; amount and definition estimated, T3; L-034, L-152). Real evidence, narrower than it sounds. Do not read it as "no profit" either: the sign is confirmed by three outlets.

It got stronger at the commitment line. OpenAI lost $12.3B at the operating level in the second quarter on $6.7B of revenue, and the loss widened from $9.3B in the first quarter (WSJ, August 19; confirmed, T2; L-066). Its compute plan through 2030 rose from about $600B in May to about $750B in July (L-078, L-062). Anthropic added roughly $99B of neocloud contracts in 27 days of August (§3). The marks did not move: $965B for Anthropic since May 28 (L-005) and $852B for OpenAI, re-affirmed by an employee tender at the same price on August 10 (L-019, L-067). Documented commitments more than doubled in 2026 while the prices stood still. That is the doctrine test this report has to pass before it counts as analysis rather than commentary, and it passes: the cost architecture is not in the prices.

The next mispricing is specific, and it is inside Anthropic's own plan. The 77% gross margin the company has shown investors for 2028, the $22B training budget for that year, and the take-or-pay status of its gigawatt contracts cannot all be true at once (§5). The market prices the margin and the gigawatts as if both were real. Exhibit 1 sets the two companies side by side as they stand on September 9.

Exhibit 1. The entry fee at a glance, both companies (Workbook A / 08_Output / A1:H14)

Item	Anthropic	OpenAI
Standing mark (post-money)	$965,000M (Series H, May 28, 2026; confirmed; L-005)	$852,000M (Mar 31, 2026; employee tender at the same price Aug 10; confirmed; L-019, L-067)
Equity raised, equity-only	$124,254M (estimated, sum of ten PB rounds; L-007)	$181,216.5M (estimated, sum of PB rounds; L-022)
Run-rate (basis; as-of)	$65,000M, GROSS, end-Jul 2026 (confirmed; L-032)	>$40,000M, basis unstated, Jul 2026 (confirmed; L-065)
Q2-2026 recognized revenue	$11,500-11,600M, gross presumed; Q1 $4,730M (confirmed; L-033, L-151)	$6,700M, net presumed; Q1 $5,700M (confirmed; L-066)
Q2-2026 operating result	Positive ADJUSTED operating income, non-GAAP, definition unpublished; projected $559M ≈ 5.1% (L-034, L-151, L-152)	Operating loss $12,300M incl. SBC; Q1 $9,300M (confirmed; L-066)
FY2026E recognized (derived)	$57,580-62,580M gross; $48,830M if flat from July	$32,400-45,500M, net presumed; PB projection $41,300M
Documented-$ compute contracts	$324,100M, eight counterparties; $358,600M incl. the $34,500M TPU lease SPV	$480,400M, four counterparties (Oracle, AWS, CoreWeave, Cerebras)
Reported, not verified	Google Cloud ~$200,000M over 5 yr (T3 [VERIFY]; L-155): reported-$ $524,100M	Azure $250,000M, Broadcom $350,000M, AMD $90,000M (could-not-verify; L-121, L-122, L-123)
Unpriced gigawatts	Google 5 GW contracted, 10 GW line of sight; AMD 2 GW	SB Energy 8.75 GW-IT leases, rent undisclosed; AMD 6 GW; Broadcom >5 GW
IPO status, Sep 9, 2026	Confidential S-1 Jun 1; no public S-1; mid-October reported (T3; L-001, L-003, L-040)	Confidential S-1 Jun 8; no public S-1; "public company in 2027" per the CFO (L-002, L-004, L-068)

Note: $M throughout. Documented-$ counts the T1 and T2 rows plus one T3 row (Volta, $10,000M), with press-valued items marked in Exhibits 6 and 7; the reported Google value never enters documented-$.

§1 The 5-Layer Cost Stack: definition, placement rule, and the table

The v4 outline that defined the five layers could not be located in any store reachable for this run, so the definition is printed here and is the one v5 uses. Layers 1 and 2 hold commitments: what is contracted or leased, in dollars and gigawatts, by counterparty type. Layers 3 to 5 hold consumption: what the P&L absorbs in a year. One amendment to the reconstruction: a placement rule. A contract value goes to Layer 1 or 2 exactly once, an annual spend to Layer 3, 4 or 5 exactly once, and the obligation stack in §3 reconciles the two against cash. Without the rule the SpaceX contract would appear as a $45B commitment in Layer 2 and again as $15B a year of inference cost in Layer 4, and a reader adding the columns would count it twice.

Exhibit 2. The 5-Layer Cost Stack: definition and seed placement (Workbook A / 00_Assumptions / A40:F60)

Layer	Name	Placement rule (counterparty or nature)	What sits here	Unit
L1	Silicon	Chip-vendor counterparties: Nvidia, AMD, Broadcom, Google (TPU as chip via SPV), Cerebras; chip-lease SPVs; custom-silicon programs	Chip purchases, chip-lease SPV obligations, GW-denominated chip agreements, ASIC programs, warrants issued to chip vendors	contracted $, GW
L2	Capacity	Cloud, neocloud, site and power counterparties: AWS, Google Cloud, Azure, Oracle, CoreWeave, SpaceX, Nscale, Lambda, Volta, Riot, Fluidstack, SB Energy	Cloud consumption contracts, full-stack GPU leases, powered-shell leases, site leases, PPAs, network and storage inside those leases	contracted $, GW/MW, $ per MW-year
L3	Training	Annual compute consumed by training, experiments and post-training (R&D compute)	Annual training-compute budgets; per-run estimates are a different basis (C-11) and are shown beside them, never summed	$ per year
L4	Inference	Annual compute consumed by serving (COGS); on Anthropic's gross basis this line also carries cloud-partner payouts	Inference COGS, gross margin, $ per million tokens	$ per year, %
L5	People and everything else	All other annual opex	Comp and SBC, data licensing and labeling, safety and evals, legal and regulatory, GTM, G&A	$ per year

Seed placement, each figure once: the cloud, neocloud and site contracts (AWS, Google Cloud, Azure, SpaceX, Fluidstack, the four August leases) sit in Layer 2; the TPU lease SPV, AMD, Broadcom's schedules and Cerebras in Layer 1; training plans in Layer 3; margin paths, inference cost and list prices in Layer 4; comp, copyright and data in Layer 5. Nvidia's $30B equity into OpenAI is financing, not cost; its $105B guarantee is a backstop on a Layer 2 lease. OpenAI's $50B of 2026 compute spans Layers 3 and 4 and is printed once, in Layer 3, with the split shown.

Exhibit 3 is the stack itself, both companies, 2023A to 2030E. Most historical cells are holes, printed as holes; the S-1s will fill most of them within weeks.

Exhibit 3a. Layer × year, Anthropic, 2023A-2030E ($M; Workbook A / 03_Costs / A5:K40 for L3-L5 and 09_Obligations / A60:K75 for L1-L2)

Layer	2023A	2024A	2025A	2026E	2027E	2028E	2029E	2030E
L1 Silicon	HOLE	HOLE	HOLE	TPU lease SPV $34,500, 5-yr lease from mid-2026 (T3) ≈ $6,900/yr principal; 2026 half-year ≈ $3,450. Broadcom 1 GW Ironwood deploying, $ HOLE (T1)	SPV ≈ $6,900. +5 GW TPU v8i, $ HOLE (T1). AMD first GW MI450 from H1-2027, $ HOLE (T1)	SPV ≈ $6,900. +10 GW "line of sight", $ HOLE (vendor guidance). AMD 2 GW cumulative, $ HOLE	SPV ≈ $6,900. Others HOLE	SPV ≈ $6,900 (lease to ~mid-2031). Others HOLE

Layer	2023A	2024A	2025A	2026E	2027E	2028E	2029E	2030E
L2 Capacity: contracts signed; annualized priced run	HOLE	HOLE	Signed: Azure $30,000, up to 1 GW (T1); Fluidstack $50,000 DC build (T1). Annual spend HOLE	Signed: AWS >$100,000/10 yr, 5 GW (T1); Google Cloud 5 GW, $ undisclosed (T2); SpaceX $1,250/month to May-2029 (T1); Volta $10,000/6 yr (T3); Riot $9,100/20 yr (T2); Nscale ~$45,000/6 yr (T2); Lambda $35,000 (T2). Priced run ≈ $24,700 (derived; Azure at an AJ 5-yr term)	Priced run ≈ $40,800 (derived; Lambda at an AJ 6-yr term), plus Google 5 GW ($ HOLE)	Priced run ≈ $46,500, plus Google 10 GW and AMD 2 GW ($ HOLE)	≈ $37,700 (SpaceX ends May-2029)	≈ $31,500 (AWS runs to ~2036; Riot to 2048)
L3 Training (annual budget, C-11)	HOLE	HOLE	HOLE	$7,000 (recalled, T4)	$14,000 (recalled, T4)	$22,000 (recalled, T4)	HOLE (scenario input)	HOLE (scenario input)
L4 Inference / COGS (gross basis, incl. partner payouts)	HOLE (revenue $100)	$1,940 (derived: revenue $1,000 at GM −94%, T4)	$6,000 at 40% GM on PB $10,000, or $2,700-3,600 on the implied $4,500-6,000 (C-03 switch)	$23,000-35,000 (derived: FY2026E gross recognized $57,580-62,580 at GM 44-60%, T4)	$52,000-56,000 (derived: revenue $140,000-150,000 AJ at GM 63%, T4)	$44,000-46,000 (derived: revenue $190,000-200,000 at GM 77%, T4)	HOLE	HOLE

Layer	2023A	2024A	2025A	2026E	2027E	2028E	2029E	2030E
L5 People and other	HOLE	HOLE	HOLE (Bartz settlement $1,500 agreed; cash 2026)	Comp ≈ $2,500-3,750 cash (AJ: 5,000 heads × $500-750K, anchored on H-1B base pay, T3); SBC HOLE; copyright $1,500 (T2); data licensing ≤ low single-digit $B; safety $55-115 (could-not-verify)	HOLE	HOLE	HOLE	HOLE
Cross-layer check		COGS $1,940 on revenue $1,000; net loss −$5,300 or −$8,300 (C-05)	No verified loss; FY2025 burn $5,600 (T4)	Priced commitments ≈ $28,200 (L1 + L2) inside the COGS + training envelope of $30,000-42,000: consistent. Loss/burn frozen (C-12)	Priced ≈ $47,700 vs envelope ≈ $65,800-69,500: consistent before Google and AMD (§5)	Priced ≈ $53,400 vs envelope ≈ $65,700-68,000: consistent before Google and AMD (§5)		

Note: rows L-044, L-049, L-043, L-074, L-076, L-072, L-073, L-046, L-048, L-136, L-135, L-137, L-125, L-010, L-012, L-119, L-151, L-033, L-032, L-035, L-036, L-008, L-165, L-086, L-113, L-114, L-117, L-120; AJ = analyst judgment, no ledger anchor. Revenue is gross (L-127); PB TTM fields are projections and appear nowhere in this table (Ruling 4).

Exhibit 3b. Layer × year, OpenAI, 2023A-2030E ($M; Workbook A / 03_Costs / A5:K40 for L3-L5 and 09_Obligations / A60:K75 for L1-L2)

Layer	2023A	2024A	2025A	2026E	2027E	2028E	2029E	2030E
L1 Silicon	HOLE	HOLE	Signed: AMD up to 6 GW with a 160M-share warrant at $0.01 (T1), $ undisclosed ($90,000 could-not-verify); Broadcom custom XPU program (T1), $ undisclosed ($350,000 could-not-verify); Nvidia $100,000 LOI retired (C-20)	Cerebras >$20,000 multi-year, 750 MW in tranches 2026-2028 (T1); $1,000 working-capital loan advanced BY OpenAI (T1). Nvidia $30,000 equity into OpenAI is an inflow, not cost	Broadcom Jalapeno 1.3 GW (T1), $ HOLE; AMD first GW, $ HOLE; Cerebras ≈ $5,000/yr (AJ phasing)	Broadcom >5 GW (T1), $ HOLE; Cerebras 750 MW complete	HOLE	Cerebras option window (1.25 GW) ends

Layer	2023A	2024A	2025A	2026E	2027E	2028E	2029E	2030E
L2 Capacity	Azure exclusive; $ HOLE	$ HOLE; $4,000 revolver	Microsoft $17,200 paid in CY2025 (recalled, T4, calendar basis: do not trend). Signed: Oracle $300,000/ 5 yr from 2027 (T2 press; Oracle discloses only RPO $638,000, T1); AWS $38,000 (T2); CoreWeave $6,500 order form (T1), ~$22,400 cumulative (press)	Microsoft FY2026 revenue from OpenAI arrangements $24,100, A/R $6,000 (T1); any cloud permitted from Apr-2026 (T1); AWS expanded to ~$138,000 (T2); SB Energy 8.0 GW-IT 20-yr leases + Milam 753 MW 15-yr, rent undisclosed, Nvidia RVG $105,000 (T1); Azure $250,000 could-not-verify	Priced run ≈ $110,000 (derived: Oracle $60,000 + AWS $17,250 + CoreWeave $3,733 + Cerebras $5,000 + Azure at the FY26 run $24,100); PORTS rent HOLE	≈ $110,000 priced + Stargate completions Q4-2028 (T3) + PORTS rent from ready-for-service (HOLE; scale proxy ≈ $198,000 of lease value from the RVG)	≈ $110,000 priced + leases (HOLE)	≈ $110,000 priced + leases (HOLE)
L3 Training (annual budget, C-11)	HOLE	HOLE	HOLE	$25,000 (recalled, T4). Anchor: total 2026 compute spend $50,000 (sworn, T2); $50,000 − $25,000 − $14,100 = $10,900 unexplained	$60,000 (T4)	$112,000 (T4)	$120,000 (T4, peak)	HOLE (scenario input)

Layer	2023A	2024A	2025A	2026E	2027E	2028E	2029E	2030E
L4 Inference / COGS (net revenue basis)	HOLE (revenue $2,000)	HOLE (revenue $6,000; NI −$5,000, T4)	Inference $8,400 (T3); total COGS ≈ $8,777 (derived: revenue $13,100 at 33% GM, T3); plan was 46%	Inference $14,100 projected (T3) on FY2026E $32,400-45,500: 31-44% of revenue	HOLE	HOLE	HOLE	HOLE
L5 People and other	HOLE	HOLE	Retention ~$1,500 per person × ~1,000 over 2 yrs (T3; verbatim could-not-verify); headcount 4,500	Headcount plan 4,500 → 8,000 (T3); comp ≈ $2,700-8,000 (AJ); H1-2026 operating loss $21,600 incl. SBC bounds L3 + L4 + L5 less gross profit (T2); ads run-rate $1,000 is revenue, not cost	HOLE	HOLE	HOLE	HOLE

Layer	2023A	2024A	2025A	2026E	2027E	2028E	2029E	2030E
Cross-layer check			Revenue $13,100 (T2); operating loss $20,920 and group loss $60,350 incl. a $41,550 fair-value swing (recalled, T3); cash ~$40,000 at year-end (T3)	H1 revenue $12,400, H1 operating loss $21,600 (1.74x revenue, widening: Q1 $9,300 → Q2 $12,300); burn plan $25,000-27,000 (T3, two vintages); Q1 burn $3,700 (T3); the gap between annualized loss and burn is SBC, vendor financing and prepayment timing: HOLE	Priced ≈ $110,000 + training $60,000 (part of the same compute) vs burn plan $57,000-63,000; ~$175,000/yr of compute 2027-2030 on the $750,000 plan	Training $112,000 (T4) alone exceeds the July run-rate		Positive cash flow $39,000 in the Feb-2026 plan (T3)

Note: rows L-060, L-123, L-050, L-122, L-075, L-149, L-023, L-054, L-064, L-056, L-058, L-051, L-061, L-121, L-125, L-078, L-083, L-063, L-025, L-026, L-093, L-166, L-027, L-090, L-066, L-081, L-082, L-142, L-062, L-079. Revenue is net of the Microsoft share by register convention; the July run-rate's basis is unstated (C-16).

Walk the 2026E column for Anthropic. In Layer 1 the $34.5B TPU lease is the only priced chip item, about $6.9B a year of principal over five years before the tranche coupons (Epoch AI, August 12; estimated, T3; L-044). In Layer 2, seven contracts signed in 2026 produce an annualized priced run of about $24.7B, on two term assumptions the ledger does not settle (Azure at five years, Lambda at six; analyst judgment, flagged). Layer 3 is a $7B training budget from a leaked document nobody in this run could open (recalled, T4; L-125). Layer 4 is a derivation: FY2026E gross recognized revenue of $57.6-62.6B (Q1 $4,730M, Q2 $11,600M, then a straight monthly ramp from the $65B July run-rate to the $100-120B year-end range investors expect; L-151, L-033, L-032, L-035, the last estimated, T3) times one minus a 44-60% gross margin (recalled, T4; L-119) gives cost of revenue of $23-35B. Layer 5 is mostly holes with a $1.5B copyright settlement in it (Bartz, approved July 20; confirmed, T2; L-086). The cross-layer check passes: about $28.2B of priced commitments sit inside a $30-42B envelope. It passes in 2027 and 2028 too, but only before the Google and AMD gigawatts are priced. That is the whole of §5.

The OpenAI column has anchors of a different kind. The only cash figure with an audit behind it is Microsoft's: $24.1B of FY2026 revenue from OpenAI arrangements, including revenue sharing, with $6.0B receivable at June 30 (Microsoft 10-K filed July 29; confirmed, T1; L-054). The only company-stated consumption figure is Greg Brockman's $50B of 2026 compute spend, given under oath on May 5 (Reuters; confirmed, T2; L-078). The leaked split, $25B of training (T4; L-125) and $14.1B of inference (Sacra, estimated, T3; L-083), leaves $10.9B of it unexplained, so the split is not treated as a fact. The cell that will matter most in the S-1 is Layer 2 in 2028: 8.0 GW-IT of twenty-year Ohio leases at an undisclosed rent, with Nvidia guaranteeing residual value up to $105B (SB Energy S-1 filed September 1; confirmed, T1; L-061).

The Layer 4 row across the years is the shape of the business: cost of revenue of $1.94B on $1.0B of revenue in 2024, a −94% gross margin (revenue from PitchBook deal records, T2, L-012; margin recalled, T4, L-119), $23-35B in 2026, and $44-46B on $190-200B of revenue in the company's 2028 plan (L-036). Both ends of that arc rest on T4 margin rows the S-1 will replace.

One structural finding governs how the whole table should be read. Neither lab owns a meaningful GPU fleet on its own balance sheet. The capex sits with Microsoft, Oracle, CoreWeave, SB Energy, AMD and Nvidia-backed special-purpose vehicles; both labs consume compute as lease or cloud opex with prepayments, and OpenAI adds twenty-year site leases as a tenant (SB Energy S-1, SpaceX S-1, Oracle's account of customer prepayments and "bring your own hardware"; confirmed, T1; L-132, L-061, L-046, L-057). The S-1 balance sheets will look light and the commitments notes heavy, and an investor reading the balance sheet first will get the wrong answer.

One number is deliberately absent from Exhibit 3a. No Anthropic 2026 compute-spend figure comparable with OpenAI's $50B exists at any tier this run could reach; Reuters' full text confirms the company gave none (L-150). A "$19B" figure in circulation is a conflation with the April run-rate and is not used (L-128, could-not-verify). The cell prints the derived envelope and a hole.

§2 Revenue on a comparable basis: run-rate, recognized, forward; gross and net

Start with the trap, because the S-1s will not remove it. Every "TTM" revenue field on both PitchBook profiles is a forward projection. Anthropic's profile shows $71,000M for "TTM 4Q2027" (L-009) and OpenAI's $41,300M for "TTM 4Q2026" (L-024); the OpenAI "TTM 4Q2025" field of $20,000M is the end-2025 exit run-rate, not recognized revenue, which was $13.1B (WSJ, July 22; confirmed, T2; L-063; conflict C-06). None of these is a current figure (Ruling 4), and Exhibit 4 prints them in their own block, labelled, so that nobody lifts one into a multiple by accident. The projections themselves moved: PitchBook's Anthropic FY2027 field went from $55B on July 16 to $71B (+29%) and its OpenAI FY2026 field from $30B to $41.3B (+38%) in eight weeks, which says something about the base rate for plans (§10) and nothing about revenue.

Exhibit 4. Revenue ladders: run-rate, recognized, forward, dated and with basis ($M; Workbook A / 02_Revenue / A5:K30)

Company	Type	Date or period	Value	Basis	Status, tier, row
Anthropic	Run-rate	end-Dec 2025	~9,000	GROSS	confirmed, T1 (company, Apr 6), L-139
Anthropic	Run-rate	Apr 6, 2026	>30,000	GROSS	confirmed, T1, L-139
Anthropic	Run-rate	early May 2026	>47,000	GROSS	confirmed, T1 (Series H post), L-037

Company	Type	Date or period	Value	Basis	Status, tier, row
Anthropic	Run-rate	end-Jul 2026	65,000	GROSS	confirmed, T2 (CNBC, Bloomberg, TechCrunch), L-032
Anthropic	Recognized, quarterly	Q2 2025	787	GROSS presumed	confirmed, T2 (Bloomberg), L-151
Anthropic	Recognized, quarterly	Q1 2026	4,730	GROSS presumed	confirmed, T2 (Bloomberg), L-151
Anthropic	Recognized, quarterly	Q2 2026	11,500-11,600 (preliminary)	GROSS presumed	confirmed, T2 (CNBC, WSJ, Bloomberg), L-033, L-151
Anthropic	Recognized, annual	FY2024	1,000	basis unstated	estimated, T2 (PB deal records), L-012
Anthropic	Recognized, annual	FY2025	10,000 (PB) or 4,500-6,000 (implied by the Jan-2026 guidance)	GROSS presumed	C-03 frozen; L-010 (T2), L-120 (T4)
Anthropic	Forward, derived	FY2026E	57,580-62,580; 48,830 if flat from July	GROSS	derived from L-151, L-033, L-032, L-035 (T3)
Anthropic	Forward, investor expectation	YE-2026 run-rate	100,000-120,000	GROSS	estimated, T3 (FT via TechCrunch), L-035
Anthropic	Forward, company forecast	FY2028E	190,000-200,000	GROSS presumed	estimated, T2 (Reuters exclusive), L-036
Anthropic	PB FORWARD PROJECTION (Ruling 4), never current	TTM 4Q2026 / TTM 4Q2027	65,000 / 71,000	unstated	estimated, T2, L-010, L-009
OpenAI	Run-rate	end-2025	>20,000	unstated	confirmed, T2 (Bloomberg), L-065
OpenAI	Monthly revenue	Mar 2026	~2,000 per month	unstated	confirmed, T2 (WSJ), L-063
OpenAI	Run-rate	Jul 2026	>40,000	unstated (NET assumed; C-16)	confirmed, T2 (Bloomberg), L-065
OpenAI	Run-rate growth	quarter to Aug 19	+35% QTD; enterprise +50%	unstated	confirmed, T2 (CFO all-hands via Fortune), L-069
OpenAI	Recognized, annual	FY2025	13,100	NET presumed	confirmed, T2 (WSJ), L-063
OpenAI	Recognized, quarterly	Q1 2026 / Q2 2026	5,700 / 6,700	NET presumed	confirmed, T2 (WSJ), L-066
OpenAI	Forward, derived	FY2026E	32,400-45,500	NET presumed	derived from L-066, L-065 at 0-20% monthly growth
OpenAI	PB FORWARD PROJECTION (Ruling 4), never current	TTM 4Q2025 / TTM 4Q2026	20,000 (run-rate vintage, C-06) / 41,300	unstated	estimated, T2, L-025, L-024

Now the basis. Anthropic reports gross: end-customer spend through AWS Bedrock, Google Vertex and Azure Foundry counts as revenue and the partner's share is booked as cost of revenue (Sacra;

estimated, T3; L-127). OpenAI reports net of the Microsoft share by the convention this register has used since July, though Bloomberg said the two companies "could be using different methodologies" and the basis of the >$40B figure is stated nowhere reachable (could-not-verify; L-158, C-16). Raw, the two run-rates are never comparable, and this report never puts one beside the other without the equalization named.

The internal equalization is 39.75% (Ruling 5; gate-verified on the May 2026 mix; recalled, T4; L-126). The external figure is about 27%: an OpenAI chief revenue officer memo in April claimed Anthropic's then-$30B run-rate was overstated by roughly $8B (adversarial and relayed, T4; L-126). Both are T4 and the dispute is frozen (C-13). Exhibit 5 shows what each does.

Exhibit 5. Gross-to-net restatement at 39.75% and ~27%, with OpenAI's basis switch ($M; Workbook A / 02_Revenue / A35:H50 and 06_Valuation / A5:H20)

Basis	Anthropic net run-rate	Anthropic multiple ($965,000M mark)	Anthropic CE (equity-only $124,254M)	OpenAI run-rate	OpenAI multiple ($852,000M mark)	OpenAI CE (equity-only $181,216.5M)	Relative multiple	CE ratio
Gross, unequalized (display only; forbidden as a comparison, Ruling 5)	65,000	14.8x	0.523x	40,000	21.3x	0.221x	n/a	n/a
39.75% haircut (Ruling 5, base case)	39,162	24.6x	0.315x	40,000 (NET, SW_OAI_BASIS)	21.3x	0.221x	Anthropic +15%	1.43x
33% haircut (grid midpoint)	43,550	22.2x	0.350x	40,000 (NET)	21.3x	0.221x	Anthropic +4%	1.59x
~27% haircut (C-13 external branch)	47,450	20.3x	0.382x	40,000 (NET)	21.3x	0.221x	Anthropic −5%	1.73x
39.75% haircut, OpenAI GROSS case (20% share, T4 only)	39,162	24.6x	0.315x	32,000	26.6x	0.177x	Anthropic −7%	1.78x
Jul-16 reference (superseded)	28,318 (from 47,000 gross)	34.1x	0.228x	~25,000	34.1x	0.138x	0%	1.65x
Q2-2026 recognized	6,989 at 39.75% / 8,468 at 27% (from 11,600 gross)			6,700			Anthropic ahead on either basis	

Note: rows L-032, L-065, L-126, L-005, L-019, L-007, L-022, L-033, L-066, L-133 (Jul-16 reference, recalled, T4). The 20% Microsoft share is T4 and appears only in the gross-case row; it is not a fact.

The finding that replaces the July one. On July 16 both companies sat at 34.1x their equalized run-rates and this report concluded that the market applied no discount to Anthropic's gross basis. That was true on those numbers and it is history. On the September 9 run-rates, with no new mark at either company, the multiples compressed on growth alone to 24.6x and 21.3x, and the sign of the relative multiple now flips inside the haircut range: Anthropic trades at a 15% premium to OpenAI on the internal 39.75% and at a 5% discount on the external 27%. The basis dispute alone is worth about $190B of relative value, twenty percentage points of relative multiple on a $965B mark. Nothing in either company's disclosure lets an investor choose between the two numbers, and the Series H priced at $965B on May 28 with investors who could not choose either. The "no discount" conclusion was not robust to the one input the S-1 revenue-recognition note will settle, and that is the basis risk this report exists to name.

Two comparisons, and a tension between them. On second-quarter recognized revenue Anthropic is ahead of OpenAI on either basis: $11,600M gross becomes $6,989M net at 39.75% or $8,468M at 27%, against OpenAI's $6,700M (L-033, L-066). On July run-rates the two are near parity on a net basis, $39-47B against more than $40B. Both hold together only if OpenAI's July jump, run-rate growth of more than 20% in one month (L-065), is real and sustained; its second quarter annualizes to $26.8B, so July alone had to add a third. If I had to bet, I would bet the recognized quarterly prints over the July run-rates: a completed quarter reported to investors is harder to flatter than one month annualized, and the OpenAI figure asks a single month's growth to hold through the second half. Sarah Friar's "+35% quarter-to-date" on August 19 (L-069) says it may; the third-quarter prints will say whether it did.

Anthropic's own ladder is now T1 or T2 at every step: about $9B at end-2025 and more than $30B on April 6 (company-stated; L-139), more than $47B by early May (L-037), $65B at end-July (L-032); and on the recognized side $787M in Q2 2025, $4,730M in Q1 2026 and more than $11.5B in Q2 2026, 2.45x the prior quarter (Bloomberg's full text, August 14; L-151). H1 2026 recognized revenue is about $16.2-16.3B gross. None of it is audited.

One corollary should be printed as a flag and not as a resolution, because both of its legs are T4. A relay of CNBC's August 15 report puts Anthropic's compute cost at $0.71 per revenue dollar in the first quarter, falling to a projected $0.56 in the second (L-152). If cloud-partner payouts of 39.75% of gross revenue sat in cost of revenue beside that ratio, gross margin would be about 4%, and a positive adjusted operating result on $11.6B of revenue with $3-5B of quarterly opex would be impossible. So one of two things holds: the $0.56 already contains the partner payouts, in which case compute proper is far less than 56% of revenue and the gross-margin band in §4 is understated; or the effective payout share is well below 39.75%, in which case the external 27% branch gains weight and the relative multiple leans toward the discount. The report does not pick. It notes that the one new datapoint since July points against the internal branch, and that the S-1 will decide.

§3 The obligation stack: contracted, cash, cancellable

The tallies in circulation are scopes, not numbers. "At least $175B" (Bloomberg, August 31), "at least $61B" (TechCrunch, August 26), "$1.15T", "$1.4T" and "$750B" all describe real objects, and none of them is the same object. The only figures that survive a filing are contract values documented at T1 or T2, and even those come with two attributes that matter more than the headline: whether the money is cancellable, and who stands behind the counterparty if the lab stops paying. Exhibits 6 and 7 lay each commitment out with those attributes; Exhibit 8 reconciles the tallies.

Exhibit 6. Obligation stack, Anthropic: contracted, cash paid to date, cancellable, backstop ($M; Workbook A / 09_Obligations / A5:J20)

#	Counterparty (layer)	Contracted (documented)	Cash paid to date	Cancellable, contingent or unpriced	Vendor backstop	Rows, tier, as-of
1	AWS (L2)	>100,000 over 10 yrs; up to 5 GW; ~1 GW Trainium online by end-2026	HOLE (AWS does not disclose)	Take-or-pay status not disclosed; treated as contracted. Amazon equity $13,000 invested, up to $20,000 more on milestones (inflow)	Amazon is investor and supplier	L-072, T1, 2026-04-20
2	Google Cloud (L2)	5 GW over 5 yrs; $ undisclosed in any filing: Alphabet 10-Q zero Anthropic mentions, Google release no $, Broadcom 10-Q no customer commitment (T1). Reported ~200,000 over five years, ">40%" of Google's backlog (T3 [VERIFY], page not opened). Broadcom schedule: 1 GW 2026, +5 GW 2027, +10 GW 2028 line of sight	HOLE	2027-2028 GW are vendor guidance, not disclosed contract; take-or-pay undisclosed; Google's further $30,000 equity is contingent (inflow)	Google supports the SPV and the project debt	L-073 T2; L-049 T1 (2026-09-02); L-153, L-154, L-156 T1; L-155 T3
3	TPU lease SPV "AI XPV Platform" (L1)	34,500 of lessor debt financing >1 GW of TPU systems leased to Anthropic for 5 yrs; 30,000 Broadcom-supported, 4,500 unsupported	≤ ~1,700 (derived: ≤ 3 months of a ~6,900/yr lease)	Lease presumed non-cancellable (debt-serviced); PB removed it from Total Raised (C-01)	Broadcom and Google absorb part of the losses if Anthropic or Fluidstack stops paying; tranches T+1pt / 5.75% / 8.5%	L-044, T3, 2026-08-12
4	Azure (L2)	30,000; up to 1 GW; term not stated	HOLE	None disclosed	Microsoft equity up to 5,000 (funded; $3,200 gain at Microsoft)	L-074 T1; L-055 T1

#	Counterparty (layer)	Contracted (documented)	Cash paid to date	Cancellable, contingent or unpriced	Vendor backstop	Rows, tier, as-of
5	SpaceX COLOSSUS / COLOSSUS II (L2)	1,250 per month through May-2029 (≤ ~45,000 derived: 36 months before the ramp discount)	≤ 2,560 in Q2-2026 (SpaceX AI-segment revenue, reduced fee May-June)	Entire contract terminable by either party on 90 days' notice: contracted at any moment ≈ 3,750; ~41,000 cancellable	SpaceX bought ~2,000 of gas turbines for its datacenters	L-046 T1 (2026-05-03); L-047 T2
6	Fluidstack (L2)	50,000 "investment" in US datacenters (Texas, New York); sites online "throughout 2026", "more sites to come"; no GW or $ schedule from either party	HOLE (five project companies raised 15,200 of debt for 1.43 GW: a separate TPU-site program)	Spend commitment, not a customer contract; phasing beyond 2026 undisclosed	Google conditional support on project debt	L-076 T1 (2025-11-12); L-168 T1; L-045 T3
7	Nscale (L2)	~45,000 over 6 yrs; ~460 MW Vera Rubin; capacity from late 2027	0	Terms not published [VERIFY]; 460 MW is T3/T4	Not disclosed	L-135, T2, 2026-08-26
8	Lambda (L2)	35,000; ~350 MW; Beacon Point, Nueces County TX; Anthropic-Lambda term not stated. Landlord level: Hut 8 second 352 MW IT lease, 15 yrs, 704 MW tenant capacity, campus base-term value 19,600 (50,200 with renewals) ≈ $1.86M per MW-yr	0	Anthropic-Lambda term HOLE (model keeps an AJ 6-yr term); landlord lease 15 yrs	Nvidia holds the datacenter lease; Hut 8's "high-investment-grade tenant" identified as Nvidia per FT relay	L-137 T2 (2026-08-31); L-167 T1 (2026-07-20)
9	Volta (L2)	10,000 over 6 yrs; 133 MW, Tydal, Norway	HOLE	"Reportedly": single reporting chain	Hydropower site (Bitdeer)	L-048, T3, 2026-08-04
10	Riot Platforms (L2)	9,100 over 20 yrs to Jun-2048; 191 MW powered shell	HOLE	Two 5-yr extension options to 16,100: 7,000 optional	None	L-136, T2, 2026-08-10

#	Counterparty (layer)	Contracted (documented)	Cash paid to date	Cancellable, contingent or unpriced	Vendor backstop	Rows, tier, as-of
11	AMD (L1)	Up to 2 GW MI450 in Helios racks, first GW from H1-2027; $ undisclosed	0	AMD equity up to 5,000 "subject to certain contingencies", funded through FY2028 (inflow); not a priced round	AMD leases datacenter sites itself (9,500 over up to 16 yrs)	L-043 T1 (2026-07-22); L-146 T1
12	Copyright (L5)	Bartz settlement 1,500 (final approval 2026-07-20)	Payable 2026	Music publishers >3,000 demanded; Sony / Warner Chappell suit (Aug 29) unquantified	n/a	L-086, T2
13	Debt facilities	2,500 revolver (May-2025)	Facility, not spend	15,000 expansion "set to finalize" (Sep 3); no close reported as of Sep 9; PB deal still "Upcoming": NOT CLOSED, [VERIFY]	n/a	L-014 T2; L-042 T3; L-160 could-not-verify
	Totals	Documented-$ (items 1, 4-10): 324,100. Plus the SPV lease (item 3): 358,600. Reported-$ incl. the Google leg at ~200,000: 524,100 [VERIFY, T3]. Unpriced GW: Google 5 GW contracted plus 5-10 GW line of sight; AMD 2 GW	Only bounds: SpaceX ≤ 2,560 (Q2), SPV ≤ ~1,700, copyright 1,500	Cancellable inside documented-$: SpaceX ~41,000 + Riot options 7,000 ≈ 48,000 (15%). Contingent inflows: Amazon 20,000 + Google 30,000 + AMD 5,000 = 55,000		

Exhibit 7. Obligation stack, OpenAI ($M; Workbook A / 09_Obligations / A25:J40)

#	Counterparty (layer)	Contracted (documented)	Cash paid to date	Cancellable, contingent or unpriced	Vendor backstop	Rows, tier, as-of
1	Microsoft Azure (L2)	250,000 incremental Azure commitment: could-not-verify (10-K silent; primaries 403). T1 facts: revenue share through 2030 at the same percentage subject to a total cap; non-exclusive IP licence to 2032; any cloud permitted	24,100 FY2026 revenue at Microsoft from OpenAI arrangements incl. revenue share (T1); A/R 6,000 at 2026-06-30; 17,200 CY2025 (recalled, T4, different basis)	Cap 38,000 through 2030 (T3, single outlet); percentage unconfirmed at T1/T2; stake ~27% (~135,000) at the Oct-2025 recap, Foundation 26%, investors and employees 47% (T2); post-dilution "decreased", undisclosed	Microsoft holds equity (recap gain 6,500 FY26); funding commitments 13,000, of which 11,900 funded at 2026-06-30 (T1)	L-121 T4; L-054 T1; L-051 T1; L-052 T3; L-159 T2; L-157 T1; L-158
2	Oracle (L2)	300,000 over 5 yrs from 2027, ~4.5 GW (T2 press). T1 bound: Oracle total RPO 638,000 at 2026-05-31; no customer >10% of FY26 revenue	HOLE (Oracle does not name OpenAI)	Oracle books 20,000-25,000 of customer prepayments in FY27 capex (some may be OpenAI); cancellability unknown	Oracle raises ~40,000 of debt and equity in FY27 to fund it	L-064 T2; L-056 T1; L-057 T1
3	AWS (L2)	~138,000 = 38,000 + "up to" 100,000 over 8 yrs; ~2 GW Trainium (T2/T4)	HOLE	"Up to": the 100,000 expansion may be a ceiling	None disclosed	L-064, T2
4	CoreWeave (L2)	6,500 Sep-2025 order form through 2031-05-31 (T1); ~22,400 cumulative (press)	HOLE	"Committed to pay up to"	CoreWeave funds with 9.0-9.75% notes and OEM financing	L-058 T1; L-059 T1
5	Cerebras (L1)	>20,000 multi-year; 750 MW in tranches 2026-2028, 3-4 yr terms (T1)	1,000 working-capital loan advanced BY OpenAI (a use of cash)	Option for +1.25 GW by end-2030 (2 GW total): optional. Warrant to OpenAI for 33.4M shares vests on milestones	n/a	L-149, T1, 2026-05-11
6	SB Energy PORTS-Pike, Ohio (L2)	17 leases, ~8.0 GW-IT, 20 yrs, OpenAI affiliate tenant, signed 2026-08-17 (T1); rent undisclosed; "OpenAI is not an investment-grade tenant"	0 (ready-for-service from 2028)	Scale proxy only: Nvidia guaranteed value 105,000 on 4.25 GW-IT = 24,700 per GW-IT; × 8.0 ≈ 198,000 (derived, not a disclosed rent)	Nvidia residual-value guaranties 105,000 on the initial ~4.25 GW-IT, option on the remaining 3.78 GW-IT; nothing payable until ready-for-service	L-061, T1, 2026-09-01

#	Counterparty (layer)	Contracted (documented)	Cash paid to date	Cancellable, contingent or unpriced	Vendor backstop	Rows, tier, as-of
7	SB Energy Milam County, TX (L2)	~753 MW across two buildings, 15-yr initial terms (T1); OpenAI warrants 3,991,809 at $0.01	0	Rent undisclosed; SB Energy booked a 2,573 warrant fair-value charge in H1-2026	SoftBank-affiliated landlord	L-061 T1; L-147 T1
8	AMD (L1)	Up to 6 GW; first GW on MI450; warrant for 160M shares at $0.01 vesting on purchase milestones and stock-price targets (T1)	HOLE	90,000 is a press estimate, could-not-verify; purchases are milestone-based, effectively optional	AMD credit support to neoclouds (Core Scientific)	L-060 T1; L-123 T4
9	Broadcom (L1)	Jalapeno 1.3 GW in 2027; >5 GW in 2028 incl. successor; "half the cost of a GPU" (vendor call, T1)	HOLE	350,000 / 10 GW could-not-verify; vendor "on track" language, not disclosed take-or-pay	Broadcom supply secured through FY2028 for six XPU customers	L-050 T1; L-122 T4; L-144 T1
10	Nvidia (L1)	100,000 / 10 GW LOI (Sep-2025) retired: 30,000 equity invested in the Mar-2026 round; the remaining ~70,000 "probably not in the cards" (Huang)	n/a (inflow)	The 105,000 RVG is Nvidia's contingent liability, not OpenAI's	See item 6	L-075 T2 (2026-03-04); L-061 T1
11	Stargate US sites (L2)	>9 GW planned across seven sites; Abilene 0.3 GW of 1.2 GW operational; five sites complete Q4-2028 (T3)	Abilene operating	Overlaps items 2 and 6-7: not added. UAE and other international sites: not verified this run (L-130)	Oracle / Crusoe project finance off-balance-sheet	L-079, T3, 2026-04-17
12	Debt	4,000 revolver (Oct-2024) + 700 (Mar-2026) + 520 term loan (Jul-2026) = 5,220 (T2)	Facilities	No Stargate project debt at the OpenAI level found	n/a	L-023, T2

#	Counterparty (layer)	Contracted (documented)	Cash paid to date	Cancellable, contingent or unpriced	Vendor backstop	Rows, tier, as-of
	Totals	Documented-$ (items 2-5): 480,400. Recalled-$ (items 1, 8, 9): 690,000. Sum 1,170,400 ≈ the "$1.15T obligations through 2035" tally: 59% of it rests on could-not-verify rows. Unpriced T1 leases: 8.75 GW-IT	24,100 FY26 to Microsoft (T1) + 50,000 2026 compute plan (T2) + 1,000 Cerebras loan (T1) + Q1-26 burn 3,700 (T3)	Optional inside documented-$: Cerebras option, AMD milestone purchases, AWS "up to" 100,000		

Exhibit 8. Tally reconciliation: the numbers are scopes ($M; Workbook A / 09_Obligations / A45:H58)

Company	Tally	Value	Scope or components	Vintage, row, tier	Reconciles?
OpenAI	Spend plan (Feb)	665,000	"To be spent training and operating models through 2030": P&L consumption 2026-2030	The Information 2026-02-21; L-082; T3	Same object as the $600B / $750B lines, earlier vintage, broader definition
OpenAI	Spend plan (May)	600,000	"Total compute spending through 2030"	Brockman testimony 2026-05-05; L-078; T2	Company-stated; narrower ("computing power")
OpenAI	Spend plan (Jul)	750,000	Projected computing-power spend through 2030, up from ~600,000	WSJ 2026-07-22; L-062; T2	+25% vs May in eleven weeks; consumption, not contract
OpenAI	Obligation stack	1,150,000	Contractual and announced commitments through ~2035	Register (Feb 27); L-062 notes; T4	Reconciles as 480,400 documented-$ + 690,000 recalled-$ (Exhibit 7); not a spend figure
OpenAI	Headline commitments	1,400,000	Altman's "infrastructure commitments" incl. optional and aspirational capacity	2026; L-062; T2 relay	Not reconcilable to contracts
OpenAI	2026 spend	50,000	Compute spend this calendar year	2026-05-05; L-078; T2	First-year slice of the spend plan

Company	Tally	Value	Scope or components	Vintage, row, tier	Reconciles?
Anthropic	Bloomberg (Aug 31)	"at least 175,000"	Lambda 35,000 + Nscale 45,000 + Fluidstack 50,000 + SpaceX 45,000 (36-month maximum)	L-138; T2	Yes, exactly. Scope: 2026 deals with four counterparties
Anthropic	TechCrunch (Aug 26)	"at least 61,000"	Nscale 45,000 + Volta 10,000 + AMD 5,000 (equity, not compute) + SpaceX 1,250 (one month) = 61,250	L-138; T2	Arithmetically yes, but it mixes an equity inflow with compute: not cited as a figure
Anthropic	Four-neocloud tally	">99,000"	Volta 10,000 + Riot 9,100 + Nscale 45,000 + Lambda 35,000 = 99,100	L-138; T4	Yes. Scope: four neoclouds
Anthropic	"$517B"	517,000	Not reproducible from any ledger row	L-138; T4	No: cut
Anthropic	Register, Jun 12	"$80B+ across 6 partners"	Pre-dates Riot, Volta, Nscale, Lambda	L-138; T4	Superseded: cut
Anthropic	The Information (May 5), Google leg	~200,000 over five years; ">40%" of Google's backlog	Google Cloud commitment only (5 GW from 2027); in no other tally above	L-155; T3 [VERIFY]; Alphabet 10-Q and Google release silent (L-153, L-154, T1)	Page not opened. If true, it alone roughly doubles every published 2026 tally (C-14 addendum). Own scope row; never inside documented-$
Anthropic	v5 ledger stack	324,100 documented-$ (8 contracts) + 34,500 TPU SPV lease = 358,600; reported-$ incl. the Google leg = 524,100 [VERIFY]; plus unpriced GW	Exhibit 6	Mixed T1-T3	Print the components; never a single total without this scope line

The largest single item in Anthropic's stack is the one that can be switched off. SpaceX's S-1 says the customer "has agreed to pay us $1.25 billion per month through May 2029, with capacity ramping in May and June 2026 at a reduced fee", and that "the agreements may be terminated by either party upon 90" days' notice (filed May 20; confirmed, T1; L-046). Thirty-six months at $1.25B is about $45B, which is how every tally counts it; the contracted amount at any moment is three months, about $3.75B, and roughly $41B is cancellable. That cuts both ways (§13): an option Anthropic does not exercise is capacity Anthropic does not have. SpaceX's second-quarter print shows the ramp: AI-segment revenue of $2.56B, up 247%, on a $1.26B operating loss, with $18.4B of consolidated capex (August 4; confirmed, T2; L-047).

The $34.5B that disappeared from PitchBook's Total Raised did not disappear from the stack; it moved. Between July 16 and September 8 Anthropic's PitchBook Total Raised fell by exactly $34,500M, to $126,754M (L-006). That June financing is lessor debt: Apollo-led funds, Blackstone and banks committed $34.5B to an SPV that buys more than 1 GW of Google TPU systems and leases them to Anthropic for five years, $30B of it with Broadcom's support, at Treasuries plus one point, 5.75% and 8.5% across the tranches (Epoch AI, August 12; estimated, T3; L-044). PitchBook removed it because it is not Anthropic's debt, which is right; the equity-only denominator in §8 never contained it. The lease payments, about $6.9B a year of principal, sit in Layer 1, and Broadcom and Google absorb part of the loss if Anthropic or Fluidstack stops paying. A public investor will find this in the commitments note, if anywhere.

OpenAI's vendor list is nine, not seven. Two primary filings opened in this run added counterparties the register did not carry. Cerebras's amended S-1 describes a multi-year deal "valued at more than $20 billion" for 750 MW of dedicated inference capacity in tranches through 2028, an option on 1.25 GW more by 2030, a warrant to OpenAI on 33.4M shares, and a $1.0B working-capital loan advanced by OpenAI to its supplier (filed May 11; confirmed, T1; L-149). SB Energy's S-1 documents seventeen twenty-year leases for about 8.0 GW-IT in Ohio, signed August 17 with an OpenAI affiliate as tenant, plus fifteen-year leases on about 753 MW in Milam County, Texas, rent undisclosed, with Nvidia guaranteeing residual value on the first 4.25 GW-IT up to $105B (filed September 1; confirmed, T1; L-061). Oracle's $300B over five years from 2027 and AWS's "up to" $138B remain press-valued (WSJ, July 22; estimated, T2; L-064); Oracle's 10-K discloses only remaining performance obligations of $638B, up from $138B, and no customer above 10% of revenue (T1; L-056). The Nvidia "$100B for 10 GW" letter of intent is gone: $30B of equity went into the March round, Jensen Huang said the rest was "probably not in the cards", and the guarantee took its place (TechCrunch, March 4; confirmed, T2; L-075; C-20).

The $1.15T is a tally, and it decomposes. Documented contracts of $480.4B plus three register values that could not be verified at T1 or T2, Azure $250B, Broadcom $350B and AMD $90B, sum to $1,170.4B, the "$1.15T of obligations through 2035" (L-121, L-122, L-123; could-not-verify). Fifty-nine percent of the tally rests on those three rows. The checks were real: Microsoft's 10-K contains no $250B and no revenue-share percentage (L-157), Broadcom's 10-Q names no customer commitment (L-156), and AMD's 10-Q confirms 6 GW and the warrant but no dollar value (L-060). Each of the three may be true; none is a fact yet, and the prose of this report uses gigawatts and T1 terms for them.

Cash paid to date is the column that is almost empty, and it is printed that way. The anchors that exist are Microsoft's $24.1B (L-054), OpenAI's $50B 2026 plan (L-078), the $1.0B loan to Cerebras (L-149), SpaceX's $2.56B of second-quarter AI-segment revenue as an upper bound on what Anthropic paid it (L-047), Anthropic's $1.5B settlement (L-086) and OpenAI's $3.7B first-quarter burn (The Information headline; estimated, T3; L-142). Nothing in the ledger supports an estimate of cumulative cash against the $324B and $480B stacks, so none is made.

For a report about what the S-1s will reveal, one finding is the silence itself. Alphabet's 10-Q for the June quarter mentions Anthropic zero times and reports $811.0B of total purchase commitments without itemizing them (confirmed, T1; L-153). Google Cloud's April 6 release promises "multiple gigawatts" and no dollar figure (L-154). Broadcom's 10-Q names no customer and no commitment (L-156). Microsoft's 10-K gives $24.1B of revenue and $13.0B of funding commitments, no percentage and no $250B (L-157). Four counterparties that will supply most of both labs' compute disclose no dollar value for these contracts; the only Google number in existence is a single-outlet scoop whose page returned an error. The S-1 commitments notes will be the first primary documents to put a dollar figure on the largest contracts in the industry.

§4 Training and inference: budgets, runs, and the Q2 prints

Two kinds of training number circulate and they are different objects. The first is an annual training-compute budget: Anthropic $7B in 2026, $14B in 2027, $22B in 2028; OpenAI $25B, $60B, $112B and $120B for 2026 through 2029 (leaked documents in the NEXUS tracker; recalled, T4; L-125); nobody in this run could open the documents. The second is the amortized compute cost of a single final training run: Grok 4 about $500M, GPT-4.5 about $200M of pre-training, GPT-4 about $78M, with the largest runs growing 2.4x a year and passing $1B by 2027 (Epoch AI; estimated, T3; L-115). The budgets are ten to fifty times the runs because they aggregate every run, every failed experiment and all post-training. Logged as a conflict in July (C-11), this is a basis label, and the two columns of Exhibit 9 are printed side by side and never summed.

The one training-side number with a company behind it does not reconcile to the leaked split. Greg Brockman told a court on May 5 that OpenAI would spend $50B on computing power in 2026 (Reuters; confirmed, T2; L-078). Take the $25B training budget from the leak and Sacra's $14.1B of projected 2026 inference cost (estimated, T3; L-083) and $10.9B of the $50B is unexplained: capacity ahead of use, prepayments, research compute outside the training line, or a split that was never right. The report treats the total as the fact and the split as colour.

Anthropic's second quarter is the datapoint on which the thesis turned, so its arithmetic is shown rather than summarized. Revenue was $11.5-11.6B (preliminary; L-033, L-151), 2.45x the first quarter's $4,730M (L-151). The operating result was a positive adjusted operating income (Bloomberg, August 14; confirmed, T2; L-151). The amount the press attached to it, $559M, was a company projection on $10.9B of revenue, about 5.1% (Reuters; L-034, L-152). The measure is non-GAAP; the company has not published its definition; the relays say it includes training cost and excludes stock-based compensation; the WSJ, which confirmed the sign, wrote that "improvements in the efficiency of its computing resources helped it generate an adjusted profit" and that "the methodology behind the adjusted profitability measure is unclear" (L-151, L-152). Reuters' full text confirms that no margin percentage was disclosed and that margins are "still being pressured by enormous spending on computing power, model training and hiring" (L-150). No T2 gross-margin figure exists for Anthropic as of September 9.

What the print implies for gross margin depends on two things the company has not said: quarterly operating expenses outside cost of revenue, and the revenue basis. Assume $3-5B of quarterly opex (analyst judgment; a 5,000-person payroll and the January plan's loss profile make less than $3B implausible). Then gross profit was $3.6-5.6B and gross margin 31-48% on the gross revenue basis, 51-80% over net revenue at 39.75%, and 42-66% at 27% (Exhibit 9c). A second cross-check comes from a relay of CNBC's August 15 report: compute cost per revenue dollar of $0.71 in the first quarter, projected at $0.56 in the second (T4; L-152). If compute is the whole of cost of revenue, that is a gross-basis margin of about 29% rising to about 44% in one quarter, consistent with the low end of the opex grid and with the register's 40% for 2025 (L-119). The point is not the number. The margin a public investor will see in the S-1 depends on the recognition basis as much as on the business, and the profit the same investor has already read about is a non-GAAP measure whose scope is unpublished.

Exhibit 9. Training vs inference: annual budgets vs per-run cost; the $50B split; the Q2 gross-margin cross-check ($M; Workbook A / 03_Costs / A45:K60 and 04_PL / A30:H40)

(a) Annual training-compute budgets (C-11 basis: all runs, experiments, post-training)	2026E	2027E	2028E	2029E	Row, status, tier
Anthropic	7,000	14,000	22,000	HOLE (scenario input; no row)	L-125, recalled, T4
OpenAI	25,000	60,000	112,000	120,000 (peak; 2030 HOLE)	L-125, recalled, T4
Per-run reference, final run only (never summed with the rows above)	Grok 4 ~500; GPT-4.5 ~200 pre-training + ~2 post; GPT-4 ~78; Gemini Ultra ~191; Llama 3.1 405B ~170	Cost growth 2.4x per year since 2016	Largest runs >1,000 by 2027	Cost shares: hardware 47-67%, R&D staff 29-49%, energy 2-6%	L-115, estimated, T3

(b) OpenAI 2026 compute: sworn total vs leaked split	Value	Row, tier
Total 2026 compute spend (Brockman, under oath, May 5)	50,000	L-078, confirmed, T2
Less: training budget 2026E (leak)	25,000	L-125, recalled, T4
Less: inference cost 2026E (Sacra)	14,100	L-083, estimated, T3
Unexplained	10,900	derived
Inference 2025A → 2026E; share of FY2026E revenue (32,400-45,500)	8,400 → 14,100; 31-44%	L-083 (T3), L-066, L-065
Ads run-rate (Aug 31) as a share of 2026E inference	1,000; ~7%	L-081, confirmed, T2

(c) Anthropic Q2-2026 implied gross margin = (adjusted operating income 559 + quarterly opex) ÷ Q2 revenue on each basis	Gross basis (11,600)	Net at 39.75% (6,989)	Net at ~27% (8,468)
Opex 3,000 (AJ)	31%	51%	42%
Opex 4,000 (AJ)	39%	65%	54%
Opex 5,000 (AJ)	48%	80%	66%
Cross-check: compute-cost ratio 0.71 (Q1) → 0.56 (Q2 projected), gross basis, if compute is all of cost of revenue (T4, L-152)	29% → 44%		
Consistency flag: 1 − 0.56 − haircut, if partner payouts sat in cost of revenue beside the ratio		4% (impossible with a positive adjusted result)	17%

Note: the 559 is a non-GAAP adjusted operating income (excludes SBC per relays; definition unpublished), so every cell in (c) is an adjusted gross margin, not a GAAP one (L-034, L-151, L-152, L-126). Opex rows are analyst judgment with no ledger anchor.

Inference is where OpenAI's cost base already lives. Sacra puts inference cost at $8.4B in 2025, 64% of $13.1B of revenue, rising to $14.1B in 2026 (estimated, T3; L-083): 31-44% of the derived FY2026E revenue range of $32.4-45.5B, an improvement whose size depends entirely on the top line. The advertising business, at a $1.0B run-rate by August 31 (Digiday; confirmed, T2; L-081), covers about 7% of the year's inference bill. The v4 estimate that a third of inference cost served users who paid nothing has no external source and is retired (footnote 1); the datapoints behind it, more than 1B active users and 2M businesses (company, July 31; L-080), remain, and they are the reason inference cost cannot fall with revenue.

What this section does not print are breakeven years. The April investor-document figures v4 carried, a breakeven excluding training and one including it for each company, have no fresh row and are cut (argument-map AM-32). The only forward cash-flow datapoint in the ledger is The Information's February report that OpenAI's plan showed positive cash flow of $39B in 2030 after burn of $25B in 2026 and $57B in 2027 (estimated, T3; L-082); a later vintage relayed by Sacra has $27B and $63B. Those are the figures §11 tests.

Footnote 1: the v4 "free-tier subsidy" derivation (~34% of inference cost, ~$2.86B, ~910M non-payers) was internal and could not be traced to any external source (L-083 notes); its surviving datapoints fold into this section.

§5 The margin-commitment incompatibility: the next mispricing

Anthropic's plan for 2028, as shown to the investors who will price its IPO, contains three numbers. Revenue of $190-200B (Reuters exclusive, August 14, two sources familiar with the forecast; estimated, T2; L-036). A gross margin of 77% (The Information's January reporting as carried in the register; recalled, T4; L-119). A training budget of $22B (leaked document; recalled, T4; L-125). Take them at face value and the company has $43.7-46.0B of cost of revenue and $22B of training to spend on compute in 2028: an envelope of about $65.7-68.0B a year (derived), which is what the compute can cost if the margin is to be what the plan says it is.

Against that envelope, lay the contracts that already carry a price. On the terms in Exhibit 6 and two stated assumptions (Azure amortized over five years, Lambda over six), the priced commitments run about $53.4B in 2028: AWS $10.0B, Azure $6.0B, SpaceX $15.0B, Nscale $7.5B, Volta $1.7B, Riot $0.5B, Lambda $5.8B and the TPU lease $6.9B (derived from L-072, L-074, L-046, L-135, L-048, L-136, L-137, L-044). That leaves $12-15B of the envelope for everything not yet priced, which is the largest part of the stack.

The largest part is Google. Anthropic has contracted 5 GW of TPU capacity from Google Cloud over five years from 2027 (TechCrunch, April 24; confirmed, T2; L-073), and Broadcom's chief executive told analysts on September 2 that after 1 GW of Ironwood in 2026 he expects Anthropic "to deploy another five gigawatts of TPU V version 8i in 2027" with "clear line of sight" to 10 GW more in 2028 (earnings call; confirmed, T1; L-049). AMD will supply up to 2 GW of MI450 GPUs from the first half of 2027 (AMD release and 10-Q; confirmed, T1; L-043). No filing prices any of it (L-153, L-154, L-156, L-043), so the Google leg is carried three ways, and Exhibit 10 carries all three.

The first branch is the one reported figure: about $200B over five years, or $40B a year (The Information via Reuters, May 5; estimated, T3; the page could not be opened and the figure is [VERIFY] wherever it appears; L-155). Add it to the priced stack and 2028 commitments are about $93.4B against a $65.7-68.0B envelope: an excess of $25-28B a year, before AMD and before Broadcom's line of sight. The second branch prices the contracted 5 GW on a proxy: the TPU lease SPV implies about $6.9M per MW-year for chips alone (L-044) and Riot's powered shell about $2.4M (L-136), $9.3M all in; Volta and Nscale's full-stack leases imply $12.5-16.3M (derived, T3; L-134). At $9.3-12.5M, 5 GW costs $46.5-62.5B a year and the excess over the $12-15B remainder is $30-50B. The third branch adds Broadcom's 10 GW line of sight: 15 GW at the same rates is $140-188B a year, an excess of $125-175B. The three are the same contract at three levels of disclosure, and even the mildest breaks the plan.

One consistency check says the reported figure is about the contracted tranche and not the line of sight: $200B over five years over 5 GW is $8.0M per MW-year, close to the $9.3M proxy; the same dollars over 16 GW would be $2.5M, below the SPV's chips-only rate. The reported figure fits about 5 GW, and Broadcom's additional 10 GW is additional capacity.

Exhibit 10. Anthropic: priced commitments per year vs the COGS plus training envelope, with the Google leg under each branch ($M; chart data by year; Workbook A / 09_Obligations / A80:K95; band from 07_Sensitivity under SW_GOOGLE_VALUE, SW_TPU_RATE, SW_GW_2028)

Line ($M)	2026E	2027E	2028E	2029E	2030E
Envelope: COGS + training (derived)	30,000-42,000	65,800-69,500	65,700-68,000	HOLE (GM and training are scenario inputs)	HOLE
Priced commitments, L1 + L2 (derived, two AJ terms)	≈ 28,158 (incl. SPV half-year)	≈ 47,730	≈ 53,355	≈ 44,605 (SpaceX ends May-2029)	≈ 38,355
Google leg, reported branch ($40,000/yr; T3 [VERIFY])	0 (the 1 GW of Ironwood sits inside the SPV lease)	40,000	40,000	40,000	40,000
Google leg, proxy: contracted 5 GW at $9.3-12.5M per MW-yr	0	46,500-62,500	46,500-62,500	HOLE	HOLE
Google leg, proxy: 15 GW incl. the 2028 line of sight	n/a	n/a (line of sight is 2028)	139,500-187,500	HOLE	HOLE
Gap, reported branch (envelope − priced − Google leg)	+1,842 to +13,842	−18,230 to −21,930	−25,355 to −27,655	HOLE	HOLE
Gap, proxy 5 GW	+1,842 to +13,842	≈ −25,000 to −45,000	≈ −33,000 to −49,000	HOLE	HOLE
Gap, proxy 15 GW	n/a	n/a	≈ −126,000 to −174,000	HOLE	HOLE

Note: negative = the plan does not cover the commitments. AMD's 2 GW is unpriced and sits on top of every branch. Envelope = revenue × (1 − GM) + training: 2027E revenue $140-150B (AJ) at 63% GM and $14B training; 2028E $190-200B (L-036) at 77% (L-119) and $22B (L-125). Consistency: $200B / 5 yr / 5 GW = $8.0M per MW-yr; over 16 GW $2.5M, below the SPV's chips-only $6.9M (L-044), so the reported figure fits ~5 GW. Rows L-036, L-119, L-125, L-072, L-074, L-046, L-135, L-048, L-136, L-137, L-044, L-049, L-043, L-134, L-155, L-153, L-154.

So one of three things is untrue, and the report says "cannot all be true" and stops there: it does not know which of the three gives, and each exit has a different consequence for the multiple.

The first exit is the margin. Every point of gross margin below 77% on $200B adds $2B a year of room, and at 60% and $200B the reported-branch gap closes with $8.6B to spare (Exhibit 17b). The plan then works, and the 77% that bankers are capitalizing at two-years-forward multiples (Reuters cites Palantir at 53x 2026 revenue as the comparable investors are shown; L-036) does not exist. An S-1 carries no projection, but it carries the historical cost-of-revenue line, and a 2026 gross margin in the 40s on a gross basis (§4) is what the trend would show.

The second exit is training. If the $22B is really $30B or $40B the envelope grows by the same amount, but training is operating expense, and every dollar moved into it comes out of the operating margin the same investors are pricing; the S-1 shows this only through the R&D line, and only for closed periods.

The third exit is the contracts. If the Google and AMD gigawatts are consumption-based or optional rather than take-or-pay, the stack shrinks to what the plan can afford. This is the exit the S-1 reveals directly: the purchase-obligations table must list minimum, fixed and non-cancellable amounts by year, and a 5 GW contract that appears there at $200B looks nothing like one that appears as "up to". It also has the strangest consequence: capacity a lab is not obliged to pay for is capacity it is not assured of having, and the $190-200B of 2028 revenue the multiple rests on needs the capacity.

Every leg of this section is a plan, a derivation or an unverified scoop: the margin and training rows are T4, the reported Google value a T3 single-outlet figure whose page returned an error, the lease rates T3 derivations. The conclusion is conditional and is presented as one. It is also the only place in either company's disclosure where three numbers the market is already pricing can be tested against each other with arithmetic, and the test fails on every branch. The public S-1 settles it.

§6 Who holds the risk: vendor-backstopped financing

The labs' cost of capacity is lower than any cost-stack model that charges them neocloud rates would say, and the reason is not on their balance sheets. Nvidia, Broadcom, Google and AMD carry the tail. The subsidy is real, it has a price, it has an accounting cost at the counterparties, and the one datapoint on whether it persists says it may not. Exhibit 11 is the ladder.

Exhibit 11. Who holds the risk: financing ladder and backstop register ($M; Workbook A / 10_Financing / A5:H25)

Structure	Lab	Rate or terms	Risk holder	Row, tier, as-of
TPU lease SPV, tranches A1 / A2 / B	Anthropic	Treasuries +1pt / 5.75% / 8.5%; 5-yr lease; 34,500 (30,000 Broadcom-supported, 4,500 unsupported)	Apollo-led lenders; Broadcom and Google absorb losses if Anthropic or Fluidstack stops paying	L-044 T3 (2026-08-12); L-106 T1
Project-company debt, five TPU site companies	Anthropic	15,200 for 1.43 GW; Google conditional support	Project lenders; Google	L-045 T3

Structure	Lab	Rate or terms	Risk holder	Row, tier, as-of
Nvidia holds the Lambda datacenter lease (Beacon Point)	Anthropic via Lambda	Hut 8 second 352 MW IT lease, 15 yrs; tenant capacity 704 MW; campus base-term value 19,600, 50,200 with renewals ≈ $1.86M per MW-yr (≈ $155 per kW-month); Anthropic-Lambda term undisclosed	Nvidia as tenant (Bloomberg; FT relay identifies the "high-investment-grade tenant" as Nvidia); Hut 8 as landlord	L-137 T2 (2026-08-31); L-167 T1 (2026-07-20)
Nvidia residual-value guaranties, PORTS-Pike	OpenAI	Guaranteed value 105,000 on the initial ~4.25 GW-IT; option on the remaining 3.78 GW-IT; nothing payable until ready-for-service	Nvidia	L-061 T1 (2026-09-01)
Warrants for capacity: AMD (160M shares at $0.01), SB Energy (3.99M at $0.01), Cerebras (33.4M Class N at $0.00001)	OpenAI	Vest on purchase, milestone or market-cap targets	Vendors dilute their own holders to win the customer; SB Energy booked a 2,573 warrant fair-value charge in H1-2026	L-060, L-061, L-149 T1; L-147 T1
AMD as tenant and guarantor	OpenAI, Anthropic, Meta (via neoclouds)	New long-term DC leases 9,500 over up to 16 yrs commencing 2027-2028; leases not yet commenced 4,500; DC lease guarantees max 4,100; unconditional purchase commitments 30,300	AMD shareholders	L-146 T1 (2026-06-27)
Microsoft equity accounting	OpenAI	FY2026 net gains from OpenAI investments 6,500 (recap dilution gain), after net losses of 4,800 (FY2025) and 1,500 (FY2024); funding commitments 13,000, 11,900 funded	Microsoft shareholders	L-054, L-157 T1
Neocloud unsecured notes (reference)	Market	9.00% (2031), 9.75% (2031), 9.625% (2032); CoreWeave accumulated deficit 4,000; OEM equipment financing 4,800	Bondholders	L-059 T1 (2026-06-30)

Structure	Lab	Rate or terms	Risk holder	Row, tier, as-of
Powered-shell leases (reference)	Market	Core Scientific-AMD >14,000 / 530 MW / 15 yrs, 2.5% escalators ≈ $1.76M per MW-yr; Riot ≈ $2.4M per MW-yr	Landlord	L-148 T1; L-136 T2
Full-stack GPU leases	Anthropic	Nscale ≈ $16.3M per MW-yr; Volta ≈ $12.5M per MW-yr	Neocloud and its lenders	L-134 T3 (derived)
Hyperscaler own capex (reference)	Market	Microsoft CY2026 ~175,000, two-thirds short-lived GPUs and CPUs, DC leases reclassified to operating, useful lives 15 → 25 yrs; Meta 125,000-145,000	Hyperscaler shareholders	L-145 T1; L-095 T2
Persistence of the equity leg	Both	Huang: the 30,000 OpenAI investment "might be the last time"; the 100,000 "probably not in the cards"	n/a	L-075 T2 (2026-03-04)

Price the guarantee first. The backstopped SPV tranche prices at 5.75% (L-044); CoreWeave's unbackstopped senior notes, the reference for a neocloud borrowing against GPUs on its own name, carry coupons of 9.00%, 9.75% and 9.625% for 2031 and 2032 maturities (10-Q; confirmed, T1; L-059). The spread is 300-400 basis points, and on $34.5B it is $1.0-1.4B a year of interest that Anthropic's capacity does not cost because Broadcom and Google stand behind the lease. That is the price of the guarantee, paid by the guarantors' shareholders.

It shows up in their accounts. SB Energy booked a $2,573M change in the fair value of the warrant liability in the first half of 2026 on the warrants it issued to OpenAI with the Milam County leases (S-1; confirmed, T1; L-147). AMD carries $4.1B of maximum datacenter lease guarantees and has signed $9.5B of new leases of up to sixteen years, commencing 2027-2028, to host the capacity its customers will consume (10-Q; confirmed, T1; L-146). Microsoft's $6.5B of FY2026 gains on its OpenAI investments were primarily a dilution gain from the recapitalization, not cash (L-054). None of this is an obligation of either lab: the residual-value guarantee is Nvidia's contingent liability, the SPV's debt is the lessor's, the warrants are the vendors' equity. Which is exactly why none of it is in the labs' prices today, and why the S-1s will not show it: it is in someone else's filing.

The most important sentence in any of those filings is the landlord's. SB Energy's S-1 states that "OpenAI is not an investment-grade tenant" and, in the same document, that it has leased 8.0 GW-IT of datacenter capacity to that tenant for twenty years (L-061). Twenty-year leases signed by such a tenant are quasi-debt whatever the balance sheet calls them, and the AIBQ capital-structure and infrastructure scores in §9 carry them as such. The guarantee is Nvidia's answer to the rating: nothing is payable under it until buildings reach ready-for-service, expected from 2028, and it covers residual value, not rent.

Anthropic's version of the same structure is Beacon Point. Bloomberg reported that Nvidia "would hold the lease on the data center" behind the $35B Lambda deal in Nueces County, Texas (L-137). The landlord's own release gives the terms: a second 352 MW IT lease of fifteen years that "doubles the existing high-investment-grade tenant's contracted capacity to 704 MW", a campus base-term contract value of $19.6B rising to $50.2B with renewals (Hut 8, July 20; confirmed, T1; L-167): about $1.86M per MW-year of powered shell, in line with Riot's $2.4M and Core Scientific's $1.76M (L-136, L-148). What Anthropic pays Lambda, and for how long, is not disclosed by anyone.

Whether the subsidy persists is a question the ledger answers with one quotation. Asked about the $30B Nvidia had just put into OpenAI's March round, Jensen Huang said it "might be the last time", and that the $100B envisaged the previous September was "probably not in the cards" (TechCrunch, March 4; confirmed, T2; L-075). The guarantees and warrants that followed are cheaper for Nvidia than equity, and also finite. A public investor should look for these structures in three places in an S-1: the commitments and contingencies note, the related-party note, and any variable-interest-entity disclosure. The disclosure will be thin, because most of the risk is held by parties that are not the registrant.

§7 Microsoft and OpenAI: the $24.1B anchor

One filed number anchors the whole relationship. Microsoft's FY2026 10-K states: "For fiscal year 2026, we recorded revenue from commercial arrangements with OpenAI, inclusive of revenue-sharing payments, of $24.1 billion, and accounts receivable from OpenAI as of June 30, 2026 was $6.0 billion" (filed July 29; confirmed, T1; L-054). It is the only T1 cash figure in the OpenAI stack: Azure consumption plus revenue share received, for the twelve months to June 2026, about half of the $50B OpenAI says it will spend on compute in calendar 2026. Everything else in Exhibit 12 is either an official term without a number, a single-outlet number, or a documented silence.

Exhibit 12. The Microsoft anchor ($M; Workbook A / 01_Data, Microsoft block, and 09_Obligations / A25:J26)

Item	Value	Source, as-of	Status, tier, row
FY2026 revenue from OpenAI arrangements, incl. revenue sharing	24,100; A/R 6,000 at Jun 30, 2026	Microsoft 10-K, FY ended 2026-06-30	confirmed, T1, L-054
Funding commitments to OpenAI	13,000 committed; 11,900 funded at Jun 30, 2026 (= PB rounds 1,000 + 2,000 + 10,000; 1,100 of the 2023 tranche unfunded)	Microsoft 10-K, equity-method note	confirmed, T1, L-157
Revenue-share terms (official)	OpenAI pays Microsoft through 2030 "at the same percentage but subject to a total cap"; Microsoft no longer pays revenue share to OpenAI; non-exclusive IP licence through 2032; OpenAI may serve any cloud	OpenAI post, 2026-04-27	confirmed, T1, L-051

Item	Value	Source, as-of	Status, tier, row
Revenue-share cap	38,000 cumulative through 2030	The Information via Reuters, 2026-05-12	estimated, T3 (single outlet), L-052
Revenue-share percentage	Not stated in the 10-K, the official post, or Bloomberg's full text on the >40,000 run-rate	10-K full-text search; Bloomberg 2026-08-13	silence, L-157, L-158, L-051
Ownership at the Oct-2025 recapitalization	Microsoft ~27% (~135,000); OpenAI Foundation 26% plus a warrant; investors and employees ~47%	TechCrunch, 2025-10-28	confirmed, T2, L-159
Ownership after the Mar-2026 round	"Our proportionate ownership of OpenAI decreased"; no percentage	Microsoft 10-K	confirmed (silence), T1, L-054; C-19
Equity accounting	FY2026 net gains 6,500 (primarily recap dilution gain); FY2025 net losses 4,800; FY2024 1,500	Microsoft 10-K	confirmed, T1, L-054
Commercial RPO	678,000 (+84% YoY); +25% excluding OpenAI; ~30% recognized in the next 12 months	Microsoft Q4 FY26 call, 2026-07-29	confirmed, T1, L-055
Incremental Azure commitment	250,000 (register value); not in the 10-K, the call, or the reachable recap coverage	Primaries 403; archive unreachable	could-not-verify, T4, L-121
CY2025 payments	17,200 (calendar-year basis, FT leak): not comparable with the 24,100	Register	recalled, T4, L-054 notes

The second T1 number is new to this run and reconciles PitchBook's records exactly. Microsoft has "made total funding commitments of $13.0 billion related to our investment, of which $11.9 billion has been funded as of June 30, 2026" (10-K; L-157). PitchBook's three Microsoft rounds, $1.0B in 2019, $2.0B in 2021 and $10.0B in 2023, sum to $13.0B, and PitchBook's synopsis of the 2023 tranche says it was to be paid over multiple years (L-022); so $1.1B was still unfunded at Microsoft's year-end. It is 0.6% of the $181.2B equity denominator, and stripping it would move OpenAI's capital efficiency from 0.221x to 0.222x, within rounding; the report keeps PitchBook's committed-capital basis (footnote 2).

The terms are official; the numbers attached to them are not. OpenAI's April 27 post says revenue-share payments to Microsoft "continue through 2030, independent of OpenAI's technology progress, at the same percentage but subject to a total cap", that Microsoft's IP licence runs through 2032 and is now non-exclusive, and that OpenAI can serve customers on any cloud (confirmed, T1; L-051). The cap is $38B, from The Information by way of Reuters, which "could not immediately verify the report" (estimated, T3; L-052). The percentage appears in no primary document reachable in this run, not the 10-K and not Bloomberg's full text on the run-rate (L-157, L-158); the figure that circulates is T4 and is not printed. What the cap does can be said without it: it bounds a cost line through 2030. The v4 arithmetic that turned the cap into a shift in OpenAI's breakeven year has no row and is cut.

The stake is a documented silence with a T2 starting point. At the October 2025 recapitalization Microsoft held "a roughly 27% stake, valued at about $135 billion, with the remaining 47% held by investors and employees", the OpenAI Foundation holding 26% (TechCrunch, October 28, 2025; confirmed, T2; L-159). The 10-K says the stake "decreased due to the OpenAI Recapitalization and other funding activity" and gives no figure (L-054). The right sentence is "about 27% at the recap; below 27% now; undisclosed" (C-19). Bloomberg's report on the >$40B run-rate, opened in full, does not say whether the figure is before or after Microsoft's share (L-158), so the basis switch in §2 stays open on the company's own silence.

Two other figures are bounds rather than facts. Microsoft's commercial remaining performance obligation reached $678B at June 30, up 84%, and "increased 25% when excluding OpenAI", which makes the OpenAI component large and undisclosed (Q4 FY26 call; confirmed, T1; L-055). The "$250B" of incremental Azure commitment the register carried could not be verified: not in the 10-K, not on the call, primaries returning errors, archive unreachable (L-121). The "$17.2B paid to Microsoft in FY2025" that circulated in July is a calendar-year figure from a leak on a different basis from the $24.1B, and the two are not trended (L-054 notes). Microsoft, Alphabet, Google and Broadcom file no dollar value for their contracts with these two labs. The $24.1B is the exception, and it is a revenue line at the supplier, not a commitment at the customer.

Footnote 2: capital efficiency in §8 uses PitchBook's committed-capital sum ($181,216.5M) as the OpenAI denominator; the $1.1B unfunded at June 30, 2026 (L-157) is inside that sum and is not stripped.

§8 Capital efficiency on an equity-only denominator

The denominators reproduce PitchBook exactly, which is the test they have to pass before anything is divided by them. Anthropic's ten completed equity rounds sum to $124,254M; add the $2,500M revolver of May 2025 (L-014) and the result is PitchBook's Total Raised of $126,754M to the dollar (L-007, L-006). OpenAI's equity rounds sum to $181,216.5M; add the $4,000M revolver, the $700M inside the March round and the $520M July term loan, $5,220M in all (L-023), and the result is PitchBook's $186,436.5M (L-022, L-021). Debt is not equity and does not enter a capital-efficiency denominator (Ruling 1); neither does the $34.5B lease SPV, which was never Anthropic's debt (§3), nor AMD's "up to $5 billion", which is a contingent, milestone-based commitment funded through fiscal 2028 and not a priced round (AMD 10-Q; confirmed, T1; L-043).

Exhibit 13. Capital efficiency walk, July 16 to September 9, equity-only denominators, both bases ($M; Workbook A / 05_Cash_Fund / A40:H55)

Date and basis	Anthropic gross run-rate	Anthropic net run-rate (haircut)	Anthropic equity-only	Anthropic CE	OpenAI run-rate (basis)	OpenAI equity-only	OpenAI CE	CE ratio
Jul 16, 2026 (register; recalled, T4, L-133)	47,000	28,318 (39.75%)	124,254	0.228x	~25,000 (net, estimate)	181,216.5	0.138x	1.65x
Sep 9, 2026, base case (Ruling 5)	65,000	39,162 (39.75%)	124,254	0.315x	40,000 (NET, SW_OAI_BASIS)	181,216.5	0.221x	1.43x

Date and basis	Anthropic gross run-rate	Anthropic net run-rate (haircut)	Anthropic equity-only	Anthropic CE	OpenAI run-rate (basis)	OpenAI equity-only	OpenAI CE	CE ratio
Sep 9, 2026, external branch (C-13)	65,000	47,450 (27%)	124,254	0.382x	40,000 (NET)	181,216.5	0.221x	1.73x
Sep 9, 2026, raw gross (display only; forbidden as a comparison)	65,000	65,000 (none)	124,254	0.523x	40,000	181,216.5	0.221x	2.37x
Sep 9, 2026, OpenAI gross case (20% share, T4)	65,000	39,162 (39.75%)	124,254	0.315x	32,000 (GROSS less 20%)	181,216.5	0.177x	1.78x

Note: rows L-032, L-065, L-007, L-022, L-126, L-133. CE = run-rate ÷ equity-only capital raised; the debt facilities and the lease SPV never enter the denominator (Ruling 1).

On the data, both companies got more efficient and the gap narrowed. Anthropic's equalized ratio rose from 0.228x on July 16 to 0.315x on September 9 as the run-rate rose 38% on an unchanged denominator; OpenAI's rose from 0.138x to 0.221x as its run-rate rose 60% from a lower base. The equalized advantage fell from 1.65x to 1.43x; on the external 27% branch it is 1.73x; raw and unequalized it would be 2.37x, printed once, in red in the workbook, so that nobody uses it. OpenAI grew faster from a lower base, and that is the same fact that flips the relative multiple in §2.

One correction to the Series H. The company's own announcement says the $65B "includes $15 billion of previously committed investments from hyperscalers, including $5 billion from Amazon" (May 28; confirmed, T1; L-038); Google's $10B, committed in April at a $350B valuation with $30B more contingent on performance targets, is the other part (TechCrunch, April 24; confirmed, T2; L-073). New cash in the round was therefore at most $50B, and $10B of the $65B was priced 64% below the round's own valuation. That does not change the denominator, which counts what was raised; it changes what the $965B post-money means: the price at which $50B of new money cleared, with $15B of older money riding at the new mark.

Read this section with the next one, or the two will seem to contradict each other: capital efficiency improved on the data, and the AIBQ capital-efficiency score in §9 falls, because the prior sub-score was not reproducible from the rubric, not because the business got worse.

§9 AIBQ: Capital Efficiency and Compute Independence re-run

Start with the caveat, because it governs every number below. The AIBQ v3.0 changelog states that the May 27 sub-scores were "decomposed from" the earlier dimension scores and "not independently validated yet", so this is the first rubric-based validation of the CE and CI sub-scores for both companies. Moves are labelled data-driven, where a ledger row changed an input, or method-driven, where the prior sub-score could not be reproduced from the rubric on any input. The moves exceed the rubric's daily caps and are logged as a re-baseline with the 24-hour cooling period. RQ, GO and MD were not re-scored; the events that touch them are listed at the end for the NEXUS scorer.

Anthropic's capital-efficiency dimension falls from 8.0 to 7.4 while its capital efficiency on the data improved (§8). The fall is a method correction. CE-1 was 10.0; the rubric's efficiency index on the

September inputs is 0.55, the 7-8 band, and the 0.315x ratio sits in the same band. The 9-10 band requires positive trailing free cash flow, and Ruling 3 keeps Anthropic's unswept: a projected, non-GAAP adjusted operating income that excludes stock-based compensation and whose definition is unpublished is further from free cash flow than a GAAP operating profit would be (L-151, L-152). CE-1 is scored 8.0, a −2.0 move flagged as method-driven; only the raw gross ratio of 0.523x reaches 9-10, and Ruling 5 forbids that basis. CE-2 rises to 5.0 on the Q2 arithmetic in §4 (T4-weighted, flagged); CE-3 holds at 9.0 on direction, the measure flagged as adjusted; CE-4 falls to 7.5, or 7.0 if the $15B facility closes, because the lease stack keeps both branches below the prior 8.0 and the facility is carried as not closed: no closing has been reported and PitchBook still shows the deal as upcoming (L-160). Net, the data-driven moves add +0.175 and the method correction takes −0.80.

OpenAI's moves are data-driven and mostly T1. CE-1 rises from 3.0 to 5.0 because 0.221x sits in the 0.2-0.3x band that a $25B run-rate did not reach (4.0 on the gross case). CE-4 moves from 2.5 to 3.0: the rubric's letter would give 7-8 on $5.2B of debt and a $122B round, but the prior score treated the obligation stack as debt-like, and a $480.4B documented stack at twelve times the run-rate plus twenty-year leases signed by a tenant the landlord calls non-investment-grade justify holding that reading and crediting the round half a point. CI-2 rises from 3.0 to 6.0 on three T1 events, Cerebras's dedicated capacity, the SB Energy leases and Broadcom's dated custom-silicon schedule (L-149, L-061, L-050); CI-4 from 6.0 to 7.0; CI-1 and CI-5 half a point each as exclusivity ends and Microsoft's share of the 2026 compute plan, about 48% on $24.1B against $50B, just clears the 50% line. Anthropic's CI rises from 5.05 to 5.55: nine providers, structures between colocation and ownership, four chip families, and no power contract of its own anywhere in the ledger (L-131, could-not-verify), which takes CI-3 down a point.

Exhibit 14. AIBQ CE and CI sub-scores, both companies: old, new, rubric basis, rows, type, dimension delta (Workbook A / 11_AIBQ / A5:L40)

Company	Sub-score (weight)	Old	New	Rubric basis for the new score	Rows	Type	Δ dimension	Flag
Anthropic	CE-1 Primary efficiency (40%)	10.0	8.0	Efficiency index 0.55 → 7-8 band; ratio 0.315x → 7-8; 9-10 needs TTM FCF+, which Ruling 3 keeps unswept	L-032, L-007, L-126, L-034, L-151, L-152	Method	−0.80	Yes (≥0.5)
Anthropic	CE-2 Gross margin quality (25%)	4.0	5.0	2025 ~40%; 2026 T4 range 44-60%; Q2 arithmetic 31-48% gross / 51-80% net; compute ratio implies ~44% gross; lower edge of the 50-65% band on the net basis	L-119, L-034, L-151, L-152, L-033	Data, T4 input	+0.25	Yes

Company	Sub-score (weight)	Old	New	Rubric basis for the new score	Rows	Type	Δ dimension	Flag
Anthropic	CE-3 Burn trajectory (20%)	9.0	9.0	Q1 4,730 → Q2 11,600 with the adjusted result positive and the compute ratio falling: "improving >50%" on direction; measure non-GAAP, amount a projection	L-034, L-151, L-152	Confirmed direction	0	Yes (adjusted measure)
Anthropic	CE-4 Capital structure (15%)	8.0	7.5 default / 7.0 if the $15B facility closes	Runway >24 months; dilution 17.33%; debt 2,500 on balance sheet; the lease stack (34,500 SPV, 20-yr Riot, Nscale, Lambda, Volta) keeps both branches below 8.0; facility not closed as of Sep 9	L-042, L-160, L-044, L-136, L-135, L-137, L-015	Data, T3	−0.075 / −0.15	Yes
Anthropic	CE dimension	8.00	7.375 ≈ 7.4 default / 7.30	0.4×8 + 0.25×5 + 0.2×9 + 0.15×7.5			−0.625 / −0.70	Composite −0.125 / −0.14
Anthropic	CI-1 Provider diversification (30%)	6.0	7.0	Nine providers under contract; balance unverified (AWS >100,000 / 10 yr is the largest)	L-072, L-073, L-074, L-046, L-135, L-136, L-137, L-048, L-076	Data, T1/T2	+0.30	Yes
Anthropic	CI-2 Infrastructure ownership (25%)	4.0	5.0	Fluidstack build with Anthropic's 50,000, Riot powered shell, TPU systems leased from an SPV: between colocation and owned racks; no own ASIC	L-076, L-136, L-044, L-132	Data, T1/T3	+0.25	Yes

Company	Sub-score (weight)	Old	New	Rubric basis for the new score	Rows	Type	Δ dimension	Flag
Anthropic	CI-3 Energy independence (20%)	5.0	4.0	No PPA attributable to Anthropic found; power procured by partners	L-131 (could-not-verify)	Absence	−0.20	Yes
Anthropic	CI-4 Supply-chain resilience (15%)	5.0	6.0	Four chip families (Nvidia, TPU, Trainium, AMD from 2027); none its own; all TSMC-fabbed	L-043, L-049, L-072, L-135	Data, T1/T2	+0.15	Yes
Anthropic	CI-5 Contractual lock-in (10%)	5.0	5.0	SpaceX terminable on 90 days vs a 10-yr AWS minimum and a 20-yr Riot lease: net unchanged	L-046, L-072, L-136	Confirmed	0	No
Anthropic	CI dimension	5.05 (sub-score base; 5.8 canonical)	5.55	0.3×7 + 0.25×5 + 0.2×4 + 0.15×6 + 0.1×5			+0.50 vs 5.05 / −0.25 vs 5.8	Composite +0.075 / −0.038
OpenAI	CE-1 Primary efficiency (40%)	3.0	5.0 (4.0 gross case)	40,000 / 181,216.5 = 0.221x → 5-6 band; index 0.50 borderline	L-065, L-022	Data, T2	+0.80	Yes
OpenAI	CE-2 Gross margin quality (25%)	3.5	3.5	33% (2025) = 3-4 band; no 2026 GM datapoint	L-083, L-066	Unchanged	0	No
OpenAI	CE-3 Burn trajectory (20%)	3.0	3.0	Operating loss worsened 32% QoQ (9,300 → 12,300): "worsening 10-50%"	L-066	Confirmed	0	No

Company	Sub-score (weight)	Old	New	Rubric basis for the new score	Rows	Type	Δ dimension	Flag
OpenAI	CE-4 Capital structure (15%)	2.5	3.0	Debt 5,220 minimal and runway >24 months would score 7-8 by the letter; the 480,400 documented stack (12x run-rate) and 20-yr leases as a non-investment-grade tenant held as debt-like	L-020, L-023, L-082, L-061	Data, interpretive	+0.075	Yes
OpenAI	CE dimension	3.05	3.925 ≈ 3.9	0.4×5 + 0.25×3.5 + 0.2×3 + 0.15×3			+0.875	Composite +0.175
OpenAI	CI-1 Provider diversification (30%)	5.5	6.0	Azure, Oracle, AWS, CoreWeave, Cerebras, SB Energy sites, any cloud; Microsoft ≈ 48% of the 2026 compute plan	L-051, L-054, L-078, L-058, L-064, L-149	Data, T1/T2	+0.15	No
OpenAI	CI-2 Infrastructure ownership (25%)	3.0	6.0	20-yr leases on 8.0 GW-IT and 15-yr on 753 MW as tenant; custom ASIC with a dated 2027 deployment; dedicated Cerebras capacity; not "owned" because SB Energy and Oracle own the sites and Nvidia guarantees residual value	L-061, L-050, L-149, L-132	Data, T1 (three events)	+0.75	Yes

Company	Sub-score (weight)	Old	New	Rubric basis for the new score	Rows	Type	Δ dimension	Flag
OpenAI	CI-3 Energy independence (20%)	4.0	4.0	No OpenAI PPA in the ledger; power with Oracle campuses and SB Energy	L-079, L-061	Unchanged	0	No
OpenAI	CI-4 Supply-chain resilience (15%)	6.0	7.0	Nvidia + AMD 6 GW + Broadcom custom XPU + Cerebras: "multiple + custom", low end	L-050, L-060, L-149	Data, T1	+0.15	Yes
OpenAI	CI-5 Contractual lock-in (10%)	4.0	4.5	Exclusivity gone; minimums (300,000 Oracle, ~138,000 AWS, 20-yr leases) pull down	L-051, L-064, L-061	Data, offsetting	+0.05	No
OpenAI	CI dimension	4.50	5.60	0.3×6 + 0.25×6 + 0.2×4 + 0.15×7 + 0.1×4.5			+1.10	Composite +0.165

Note: weights per Ruling 6 (CE 20 / RQ 25 / CI 15 / GO 20 / MD 20; CE sub-weights 40/25/20/15; CI 30/25/20/15/10). Old sub-scores are the May 27 vintage (recalled, T4; L-133). Flag = a move of ≥0.5 at sub-score level or ≥0.1 at composite level.

Composites move less than the sub-scores because the dimension weights are 20% and 15%. Anthropic goes from 8.20 to 8.15 by default, 8.13 if the facility closes and 8.04 if the delta is applied to the 5.8 canonical CI base rather than the 5.05 sub-score base; it stays Strong. OpenAI goes from 4.53 to 4.87, or 4.79 on the gross-basis case; it stays Developing, below 5.0. Neither tier changes. The only valuation expression the rubric permits is dollars of mark per point of quality, and Exhibit 15 prints it. The correlation coefficient between quality and valuation across the coverage set is embargoed (Ruling 2) and does not appear in this report in any form.

Exhibit 15. The $/AIBQ-point ladder, old and new ($B of post-money mark per composite point; Workbook A / 06_Valuation / A25:F32)

Company and branch	Old composite (May 27)	New composite (Sep 9)	Tier	Mark ($M)	$/point, old	$/point, new
Anthropic, default (facility does not close; sub-score CI base)	8.20	8.15	Strong (unchanged)	965,000 (L-005)	118	118
Anthropic, facility closes	8.20	8.13	Strong	965,000	118	119
Anthropic, canonical CI base (5.8)	8.20	8.04	Strong	965,000	118	120

Company and branch	Old composite (May 27)	New composite (Sep 9)	Tier	Mark ($M)	$/point, old	$/point, new
OpenAI, net case	4.53	4.87	Developing (unchanged)	852,000 (L-019)	188	175
OpenAI, gross case (C-16)	4.53	4.79	Developing	852,000	188	178
Ladder spread (OpenAI ÷ Anthropic)	1.60x	1.48x default (1.47x if the facility closes)				The market pays $57B more per point for OpenAI, down from $70B

The ladder narrows from 1.60x to 1.48x because OpenAI's scores improved on ledger evidence while its mark stood still. Nothing more is claimed for it.

Events outside this section's scope, listed for the NEXUS scorer and not scored here: both confidential S-1s (GO-3; L-039, L-003, L-001, L-002); OpenAI's C-suite departures from April to August (GO-5, down; L-089, L-070); the Musk verdict and the pending NYT summary judgment (GO-4; L-087, L-088); the completed recapitalization (GO-2, up; L-071, L-159, L-157); Anthropic's Bartz approval and new music-publisher suits (GO-4, mixed; L-086); its more than 1,000 customers above $1M and held price ladder (RQ-3, RQ-5, up; L-139, L-084); OpenAI's ads run-rate and user base (RQ-3, mixed; L-081, L-080); the retention program and executive exodus (MD-3, down; L-093, L-166, L-089).

§10 Outside view: reference classes and base rates for every forecast

Before any forecast in this report is used, it should be asked what usually happens to forecasts like it. The ledger contains its own reference class, and the base rate it yields is not subtle: in every observed case a revenue plan was beaten, a cost plan was revised up, a margin plan was missed and a dated milestone slipped. Where a base rate has an external source, Exhibit 16 cites it; where none exists, it says "analyst judgment".

Revenue plans first. Anthropic's January 2026 plan called for $18B of revenue in 2026, $55B in 2027, $102B in 2028 and $148B in 2029 (The Information, as relayed; recalled, T4; L-120). By the end of July the run-rate was $65B (L-032) and by August the 2028 forecast shown to IPO investors was $190-200B (L-036): the 2028 number roughly doubled in seven months. OpenAI's February plan had $30B of 2026 revenue (L-082); the July run-rate was more than $40B (L-065). PitchBook's own projections moved 29% and 38% in eight weeks (L-009, L-024). Cost plans went the other way: OpenAI's compute plan through 2030 rose from about $600B in May to about $750B in July (L-078, L-062), and its cumulative burn forecast by about $111B in one February revision (L-082). Margin plans were missed: OpenAI's 2025 gross margin was 33% against a 46% plan (Sacra; estimated, T3; L-083), and Anthropic lowered its projection in January (L-119). Milestones slipped: no public S-1 from either company by September 9 against PitchBook windows of September and October (L-001, L-002, L-003, L-004), Anthropic's prospectus now reported for mid-October (T3; L-040), ChatGPT's billion users seven months late (L-080), Abilene capped at 1.2 GW from 2.1 GW with five Stargate sites clustered at Q4 2028 (L-079). Datacenter build cost rose 21% per megawatt to $17.6M (Cushman & Wakefield; confirmed, T2; L-109).

Exhibit 16. Outside view: reference class and base rate for every forecast this report makes (Workbook A / 13_OutsideView / A5:H30)

Forecast	Reference class	Base rate	Source row or analyst judgment	What is different here	By how much
Anthropic revenue 2027-2030	Frontier-lab revenue plans in this ledger; software growth decay	Plans beaten in every observed case; growth endurance ~70% year to year for private cloud companies, ~80% public (SaaS dataset)	L-120, L-032, L-036 (plan vs actual); L-161 (Bessemer; estimated, T3; SaaS caveat)	The SaaS rule applied to FY2026E growth (+476% on PitchBook's FY2025 $10B) gives +333% in 2027 and ~$249B, above the company's own $190-200B 2028 plan; the plan already embeds faster decay than the base rate	Base rate shapes only the 2029-2030 tails; the bear risk is a regime break (consumption revenue, price compression), not decay
OpenAI revenue 2027-2030	Same	Plans beaten ($30B 2026 plan vs >$40B run-rate); PB projection +38% in eight weeks	L-082, L-065, L-024; CAGR 40 / 70 / 100% is analyst judgment; L-069 (+35% QTD) as colour	Growth from a $40B base at 100% a year has no software precedent	No sourced decay rate for a consumer-plus-API mix
Anthropic gross-margin path (40% → 77%)	Capex-heavy compute platform	AWS operating margin 9.9% (2014) → 25.4% (2016) → 27-30% (2018-2023) → 37.0% / 35.4% (2024-2025): 25% in three years, a plateau near 30% for six	L-162 (four Amazon 10-Ks; confirmed, T1); L-119 (T4 plan)	Anthropic's 77% is a GROSS-margin target on a gross-revenue basis; AWS's is an operating margin: different line, different basis; margin plans were missed at both labs	The 2028 GM is the input §5 and §11 flex first (60 / 70 / 77%)
OpenAI gross-margin path (33% → ?)	Same	Margin plan missed by 13 points (33% vs 46%)	L-083 (T3)	No 2026 GM datapoint exists; inference 31-44% of FY2026E revenue	Scenario input 50 / 40 / 30% of revenue 2027E+ (analyst judgment)
Both training ramps	Cost plans in this ledger	Revised up in every observed case (+25% in eleven weeks; +$111B burn)	L-078, L-062, L-082; training rows L-125 (T4)	Budgets 10-50x per-run cost because they aggregate runs and experiments (L-115)	2029-2030 training is a scenario input (analyst judgment); Bear adds 25% to plan

Forecast	Reference class	Base rate	Source row or analyst judgment	What is different here	By how much
Both burn paths	Same	Two vintages in six weeks for OpenAI ($25B / $57B → $27B / $63B); Anthropic's January "lose ~$11B" vs a positive Q2 adjusted result (C-12)	L-082, L-142, L-120, L-151, L-152	The adjusted measure is non-GAAP and unpublished; GAAP trajectory unknown	Burn bridge from operating loss is a HOLE (SBC, prepayments, vendor financing)
Both IPO windows	Dated milestones in this ledger; IPO filed-to-priced timing	Slipped in every observed case (S-1 windows, ChatGPT 1B users, Abilene, Stargate sites); no filed-to-priced or withdrawal statistic exists (PitchBook, Renaissance, SEC)	L-001, L-002, L-004, L-040, L-080, L-079; L-163 (could-not-verify): analyst judgment	The one documented frontier-adjacent case is Cerebras: confidential file → cancel → refile → price about nine months later; PitchBook Q2 colour: median step-up 1.3x at listing	September is closed by arithmetic (15-day rule); Q4-2026 vs 2027 for OpenAI; October pricing needs a public flip by late September for Anthropic
Obligation consumption (contracts to cash)	Datacenter build and GPU economics	Build cost +21% per MW to $17.6M; H100 new $25-40K vs used $8.2-25K; hyperscalers extended DC useful lives 15 → 25 years while GPUs stay short-lived (two-thirds of capex)	L-109 (T2), L-103 (T4), L-145 (T1)	Both labs consume via leases with vendor backstops (§6), so the residual-value risk sits with Nvidia, Broadcom, Google, AMD	Lease terms (Azure, Lambda) and PORTS rent are HOLEs; annualized run uses two AJ terms
Greenfield capital to milestones (§12)	2025-2026 entrants and xAI	$2-8B buys a credible model and little revenue; $40B+ of losses buys vertical integration and ~$3B of revenue; $124B and $181B of equity bought $65B and $40B of run-rate	L-097, L-098, L-099, L-100, L-096, L-007, L-022	Custom-silicon supply booked through 2028 by six incumbents (L-144)	Each scenario in §12 is checked against the band or flagged "outside reference class"

The growth base rate cuts against the bear case, and the arithmetic should be shown. Bessemer's benchmark says a private cloud company should expect next year's growth rate to be about 70% of this year's (September 2021; estimated, T3; a SaaS dataset of 200-plus portfolio companies, not consumption-billed labs; L-161). Anthropic's derived FY2026E gross revenue of $57.6B is +476% on PitchBook's FY2025 $10.0B (L-010). Apply 70% endurance and 2027 growth is +333%, which puts

2027 revenue near $249B, above the company's own $190-200B plan for a year later. The plan already embeds faster decay than the base rate, so the base rate cannot generate a bear case on its own; the bear risk for Anthropic is a break in regime, consumption revenue that reprices or a customer base that concentrates. The report uses the base rate only to shape the 2029-2030 tails in §11, where no company figure exists.

The margin base rate has to be read with its basis. AWS is the one audited history of a capex-heavy compute platform moving from loss to profit at scale: operating margin of 9.9% in 2014, 25.4% by 2016, a plateau of 27-30% from 2018 to 2023, and 37.0% and 35.4% in 2024 and 2025 after a useful-life change and scale (four Amazon 10-Ks; confirmed, T1; L-162). That says a platform can reach 25% in three years and sit near 30% for six. It does not say anything about 77%. Anthropic's target is a gross margin on gross revenue; AWS's figures are operating margins after everything. The two are not the same line and the report does not conflate them, but the comparison still carries information: no capex-heavy platform in the reference class has shown the margin expansion Anthropic's plan requires between 2025 and 2028, and both labs have already missed the margin plans they set.

The one base rate with no source is the one the IPO calendar rests on. No statistic for time from confidential filing to pricing, or for withdrawal rates, could be found: PitchBook has none, Renaissance's is paywalled, the SEC publishes counts only (L-163, could-not-verify). The slippage base rate in Exhibit 16 is analyst judgment and says so. The one documented case in the window is Cerebras, which filed confidentially, cancelled a 2025 listing, refiled and priced about nine months later at $34.3B (PitchBook Q2 2026 US VC Valuations report; confirmed, T2; L-163); the same report's 1.3x median step-up at listing is colour, not a forecast.

§11 Scenarios and sensitivities

Three inputs move the answer and everything else is noise at this level of rigor: the gross-to-net haircut, Anthropic's 2028 gross margin, and OpenAI's revenue growth against a compute plan that averages about $175B a year from 2027 to 2030 on the $750B figure (L-062). Exhibit 17 flexes each across the range the ledger supports.

Bear, Base and Bull are defined in the workbook and stated here once. Anthropic: year-end 2026 run-rate flat at $65B, $110B or $120B (L-032, L-035); 2027 revenue $120B, $145B or $160B (analyst judgment between the year-end range and the 2028 forecast); 2028 revenue $150B, $190B or $220B (L-036 at Base; the Bear shortfall and the Bull beat are analyst judgment); 2029 and 2030 from Bessemer endurance on each scenario's 2028 growth rate, 70% in Bear and Base and 80% in Bull as a public company (L-161): Bear +25.0% → +17.5% → +12.2%, $176.3B and $197.8B; Base +31.0% → +21.7% → +15.2%, $231.3B and $266.4B; Bull +37.5% → +30.0% → +24.0%, $286.0B and $354.6B. Gross margin 2026 / 2027 / 2028 of 44 / 55 / 60% in Bear, 52 / 63 / 70% in Base and 60 / 63 / 77% in Bull; the Base 2028 margin sits below the 77% plan on purpose, because both labs have missed every margin plan in the ledger and the §5 envelope does not close at 77%. Training at plan plus 25% in Bear, plan in Base, plan less 15% in Bull (L-125). The Google leg is the reported $40B a year at Base, flagged [VERIFY] on every cell that reads it, because it is the only dollar figure in existence for that contract; the proxy at $16.3M per MW-year in Bear and the chips-only $6.9M in Bull. OpenAI: second-half 2026 monthly growth 0 / 10 / 20% (L-065); 2027-2030 revenue CAGR 40 / 70 / 100% (analyst judgment; the CFO's +35% quarter-to-date is colour, L-069); compute plan $900B / $750B / $600B, Bear being the higher spend (L-062, L-078); other opex 40 / 30 / 20% of revenue. The haircut is a switch, not a scenario: 39.75% throughout, because Ruling 5 is gate-verified on the May 2026 mix and the external 27% is adversarial and undocumented; the 27% branch is printed beside it wherever a net figure appears.

Exhibit 17a. Basis grid: haircut × OpenAI basis → relative multiple and CE ratio (Workbook A / 07_Sensitivity / A5:H20)

Haircut on Anthropic's $65,000M gross	OpenAI basis	Anthropic multiple	OpenAI multiple	Relative multiple (Anthropic vs OpenAI)	CE ratio (Anthropic ÷ OpenAI)
39.75% (Ruling 5)	NET (40,000)	24.6x	21.3x	+15%	1.43x
33%	NET	22.2x	21.3x	+4%	1.59x
27% (C-13 external)	NET	20.3x	21.3x	−5%	1.73x
39.75%	GROSS less 20% (T4) = 32,000	24.6x	26.6x	−7%	1.78x
33%	GROSS less 20%	22.2x	26.6x	−17%	1.98x
27%	GROSS less 20%	20.3x	26.6x	−24%	2.16x

Exhibit 17b. Anthropic 2028E envelope grid: revenue × gross margin → envelope minus priced commitments ($53,355M) minus the Google leg ($M; negative = the plan does not cover the commitments; Workbook A / 07_Sensitivity / A25:H40)

2028E revenue	GM 60%	GM 70%	GM 77% (plan)
Panel 1: Google leg reported 40,000/yr [VERIFY]; training 22,000 (plan)			
150,000	−11,355	−26,355	−36,855
175,000	−1,355	−18,855	−31,105
200,000	+8,645	−11,355	−25,355
Panel 2: Google leg reported 40,000/yr; training 30,000 (the training exit)			
150,000	−3,355	−18,355	−28,855
175,000	+6,645	−10,855	−23,105
200,000	+16,645	−3,355	−17,355
Panel 3: Google leg proxy, 5 GW at $9.3M per MW-yr = 46,500; training 22,000			
150,000	−17,855	−32,855	−43,355
175,000	−7,855	−25,355	−37,605
200,000	+2,145	−17,855	−31,855

Note: cell = revenue × (1 − GM) + training − 53,355 − Google leg. AMD's 2 GW and Broadcom's 10 GW line of sight are outside every panel; adding the line of sight at the proxy rate subtracts a further ~93,000-125,000 from each cell.

Exhibit 17c. OpenAI 2030E cash-flow sign: 2027-2030 revenue CAGR × compute plan ($B; Workbook A / 07_Sensitivity / A45:H60)

2027-2030 revenue CAGR	Compute plan $600B (2030 compute 192.5)	Compute plan $750B (245.0)	Compute plan $900B (297.5)
40% (2030E revenue 146.4)	−90	−142	−195
70% (2030E revenue 318.4)	+30	−22	−75
100% (2030E revenue 609.9)	+234	+182	+129

Note: derived on analyst-judgment drivers: 2026E revenue base $38,119M (Base, 10% monthly growth from the July run-rate); 2027-2030 compute = plan less the 2026 $50B, phased 15 / 22 / 28 / 35%; other opex 30% of revenue. The Information's own figure for the company's plan is positive cash flow of $39B in 2030 (Feb-2026 vintage; estimated, T3; L-082). Sign, not magnitude, is the output.

Three readings follow from the grids and none is a point estimate. The relative value of the two names swings about $190B on the haircut alone, from a 15% premium to a 5% discount, and further if OpenAI's run-rate turns out to be gross. Anthropic's 2028 plan covers its priced commitments and the reported Google leg only if the gross margin is around 60% with revenue at the top of the forecast, or if training is well above the leaked budget; at the plan's own 77% no cell in any panel is positive. OpenAI's 2030 cash flow is positive only if revenue compounds at about 70% a year or more against the $750B plan, a rate with no reference class at a $40B base; at 40% the single-year shortfall is $90-195B. The Information's $39B for 2030 sits between the 70% and 100% rows of the middle column, which says what the company's own plan assumes.

§12 The greenfield entry fee: starting from zero in 2026

The two incumbents' stacks say what a seat at the frontier has cost so far: $124B and $181B of equity for $65B and $40B of run-rate, $324B and $480B of contracts to keep the seat, and, at one of them, a first positive quarter on an adjusted basis. That is the price paid by two companies that bought their compute and their people when both were cheaper. It is not the price of a seat bought today. This section prices one from zero, three ways, line by line over five years, and checks each answer against what the entrants of 2025 and 2026 actually paid. Three numbers come out of it, with ranges, each fenced by a reference class.

The line items come with their anchors, and the anchors are mixed in quality. Chips: a GB200 NVL72 rack at $2-3M ($27,800-41,700 per GPU), a GB300 rack at $3-4M, H100s at $25-40K new and $8,200-25,000 used, from a price index rather than a list (estimated, T4; L-103). Renting: CoreWeave's on-demand rates on September 9 were $6.16 per H100-hour, $8.60 per B200 and $10.50 per GB200 GPU-hour, Lambda's B200 $6.69-6.99, Nebius's $7.15 (confirmed, T1; L-104); contracted large clusters price at $2.40 per B200-hour at the neocloud 25th percentile, $3.10 at a hyperscaler median and $4.00 in a 5,184-GPU GB300 example (SemiAnalysis; estimated, T3; L-105). Debt: 9.00-9.75% unbackstopped, 5.75% with a vendor behind the lease (L-059, L-044). Power: about 2.0-2.2 MW of facility per 1,000 GB200-class GPUs (T3; L-107) at 6.58 cents per kWh in Texas and 9.17 for the US industrial average (EIA, June 2026; T1; L-108). Building: $17.6M per megawatt all-in, up 21% since late 2024 (Cushman & Wakefield; T2; L-109); renting the shell instead, $1.76-2.4M per MW-year (L-148, L-167, L-136); renting the whole stack, $12.5-16.3M (derived, T3; L-134). Network 3-15% of GPU cost (T3; L-111); storage 2-5% of capex, from aggregators only (could-not-verify; L-112). A frontier final run about $500M today and more than $1B by 2027 (L-115), against incumbent annual budgets of $7B and $25B as the ceiling (T4; L-125). People: Anthropic's H-1B filings show Member of Technical Staff base salaries of $1.12M and $1.38M inside a $133,952-1,380,000 range across about 80 roles, base only (Business Insider via Business Today; estimated, T3; L-165); Meta has called a $100M four-year package for a very senior leader "not inconceivable" (T2; L-094); OpenAI's 2025 retention program paid about $1.5M per person to about 1,000 staff over two years, engineers $200-600K and key researchers "mid-single-digit millions" (relayed; T3, verbatim not reachable; L-093, L-166). Data: $50-70M a year per marquee text source (Reddit S-1; T1; L-113) and $1-3B a year of expert data at scale (T4; L-114). Safety $55-115M by one estimate (could-not-verify; L-117); legal $1.5B paid and more than $3B demanded at Anthropic (L-086), a pending NYT case with no figure at OpenAI (T4; L-088); sales and G&A only a workforce-mix proxy (T4; L-118).

Two of those lines deserve arithmetic in the text. The first is serving cost, which the ledger can now anchor at T2. SGLang on a GB200 NVL72 running DeepSeek V3/R1 delivers 13,386 output tokens per second per GPU, about 48M output tokens per GPU-hour (LMSYS, September 2025; confirmed, T2; L-164). At CoreWeave's on-demand $10.50 per GB200 GPU-hour that is about $0.22 per million output tokens, and at contracted $2.40-4.00 it is $0.05-0.08, both at 100% utilization; at 60% multiply by 1.67 (an estimate on those three inputs). A dense frontier-class model is a different object: Llama 3.1 405B runs at 138-224 tokens per second per GPU on the same hardware, 0.5-0.8M tokens per GPU-hour, about $13-21 per million on-demand and $3.0-8.1 contracted. Anthropic's list prices are $10 and $50 per million input and output tokens for Fable 5.1, $5 and $25 for Opus 5, $2 and $10 for Sonnet 5 (pricing page, September 9; confirmed, T1; L-084). The gap between the efficient-MoE benchmark and list is enormous; the gap between the dense proxy and list is not, and no lab discloses which its frontier models resemble. The $0.22 is a benchmark, not a frontier-model COGS.

The second is ownership. Owning 1 GW-IT of Blackwell-class capacity costs about $36.4B up front at Base: $17.6B of building, $16.7B of chips for roughly 476,000 GPUs, $1.7B of network and $0.5B of storage (derived; L-109, L-103, L-107, L-111, L-112), plus $0.4-0.6B a year of power at 70-85% utilization. Renting the same gigawatt as full-stack capacity costs $12.5-16.3B a year; renting only the powered shell, $1.8-2.4B. The lease-versus-own crossover is about 2.5-3 years of full-stack rent, before residual-value risk (used H100s at 33-63% of new; L-103) and before the 24-36 months a build takes (analyst judgment). Which is why the incumbents rent, and why the vendors stand behind the leases (§6).

Exhibit 18a. Greenfield: three scenarios and their line items (Workbook B / 00_Assumptions and 02-06; $M per year unless stated)

Line item	1. Lean fast-follower	2. Full frontier	3. Vertically integrated with custom silicon	Anchor rows
Compute model	100% leased, contracted large-cluster rates $2.40-4.00 per GPU-hr; no debt	Leased at scale plus purchased clusters from Y2, half of them GB300 at rack price ÷ 72; 60% debt-financed at 9.0-9.75%, or 5.75% if a vendor backstop is obtained (switch)	Own 1 GW-IT by Y4 (~476,000 GB200-class GPUs at 2.1 MW per 1,000); custom ASIC program from Y1; merchant GPUs until ASIC availability (no custom supply before 2029, AJ)	L-105, L-104, L-103, L-059, L-044, L-107, L-144
GPU-equivalents (AJ)	25,000 (Y1) → 50,000 (Y3)	100,000 (Y1) → 250,000 (Y3) → 400,000 (Y5)	50,000 leased Y1; 240,000 purchased Y2 (500 MW); 476,000 by Y4	L-103
Power, cooling, datacenter	Inside the lease	Colocation powered shell at $1.76-2.4M per MW-yr for owned clusters	Greenfield build at $17.6M per MW (AI-dense sensitivity $20M): $17,600-20,000 over Y1-Y4; power 403-619 per year at Texas rates	L-148, L-167, L-136, L-109, L-108
Network and storage	Inside the lease	8% of purchased GPU BOM; storage 2-5% of cluster capex (flagged)	5.8-8% of GPU BOM ≈ 1,000-1,700 for 1 GW; storage ≈ 500 (flagged)	L-111, L-112
Training runs	One frontier-class final run per year at ~500 rising 2.4x/yr, capped by cluster size; development multiple 3x; failed-run allowance 30% (AJ)	2-3 generations in five years; annual training budget 1,500 (Y1) → 7,000 (Y3) → 14,000 (Y5), the Anthropic 2026-2027 budgets as ceiling; multiple 4x; allowance 40%	As scenario 2 for budget; compute-hours priced at owned cost (3-yr depreciation plus power, AJ)	L-115, L-125
Inference serving	Efficient-MoE class: $0.05-0.08 per M output tokens contracted at 100% utilization, $0.08-0.20 at 40-70%; inference cost 64% of revenue (Y1) → 35% (Y5), AJ	Dense-frontier class: $3.0-8.1 contracted at 100%, $4-20 at 40-70%; 64% → 30%	Same, with the vendor's "half the cost of a GPU" ASIC claim from Y5 (flagged)	L-164, L-104, L-105, L-083, L-050
Research and engineering headcount	100 → 300; 10% leaders at $5-25M, 40% senior IC at $0.5-1.5M, 50% engineers at $0.3-0.76M ≈ 180-450 per year	500 → 3,000; same bands; retention pool from Y3 ≈ 750 per year	800 → 4,000 incl. datacenter operations and a silicon team; ASIC program 1,000-3,000 per year for 3-4 years (AJ)	L-165, L-094, L-093, L-166
Data licensing and labeling	50-200 and 100-300	300-1,000 and 1,000-3,000 by Y3	As scenario 2	L-113, L-114

Line item	1. Lean fast-follower	2. Full frontier	3. Vertically integrated with custom silicon	Anchor rows
Safety, legal, GTM and G&A	Safety 3% of research payroll (flagged); legal reserve 100 + compliance 20; GTM and G&A 20% of headcount cost	5%; reserve 300 + compliance 50; 50%	5%; as scenario 2; 50%	L-117, L-086, L-088, L-118

Exhibit 18b. Cumulative capital to milestones at Base, with ranges ($M; Workbook B / 07_Cash / A5:H60 and 09_Output / A1:H25)

Milestone	1. Lean fast-follower	2. Full frontier	3. Vertically integrated	Reference band and check
Year-1 cash out	1,000-1,500	4,000-6,000	5,000-8,000 (Y2: 15,000-25,000 as the 500 MW build and 240,000 GPUs land)	Entrant seeds and Series A-C of 2025-2026: $300M to $3B (L-097-L-102)
Cumulative to first frontier-class model	2,500-4,000 (Y2)	12,000-18,000 (Y2)	25,000-40,000 (Y3; the owned cluster lands Y2-Y3)	SSI $7.0B raised with zero revenue; TML $2.0B seed; Reflection $2.2B; Mistral $7.5B: Lean must land in the $2-8B band (PASS)
Cumulative to first $1B of net revenue	5,000-8,000 (Y4 if reached; probability 30-50%, AJ)	20,000-30,000 (Y3; 60-80%)	40,000-60,000 (Y4; 60-80%)	Mistral: $7.5B raised for a $1.2B FY2026 projection; TML "a few hundred million" annualized on $2B; xAI $41.3B of deficit for $3.2B of 2025 revenue
Cumulative to cash-flow breakeven	8,000-15,000 (Y5), or not reached in Bear	60,000-100,000+ (beyond Y5)	80,000-150,000 (beyond Y5)	Full frontier must sit below Anthropic's $124.3B of equity and the ~$12-20B it burned to a first positive adjusted result (derived, flagged), and near xAI's $41.3B deficit by Y4-Y5; Vertically integrated must sit at or above xAI's deficit plus its 2025 AI-segment debt (~$57B) by Y5

Milestone	1. Lean fast-follower	2. Full frontier	3. Vertically integrated	Reference band and check
Owning 1 GW-IT, up front	n/a	n/a	≈ 36,400 (build 17,600 + chips 16,700 + network 1,700 + storage 500) plus 400-620 per year of power; lease equivalent 12,500-16,300 per year full-stack, 1,760-2,400 shell only; crossover ≈ 2.5-3 years	L-109, L-103, L-111, L-112, L-108, L-134, L-148, L-136

Exhibit 18c. Reference class: what entry actually cost ($M; Workbook B / 08_Reference / A5:J20)

Entrant	Capital raised or spent	Revenue	Headcount	Valuation	Resembles scenario	Rows, tier
xAI (inside SpaceX)	Accumulated deficit 41,311 at Mar 31, 2026 (consolidated); AI-segment 2025 debt proceeds 16,055; pre-merger raise not reproducible	AI segment 2025 revenue 3,201, loss from operations (6,355); Q2-2026 revenue 2,560 (+247%), operating loss (1,260); consolidated Q2 capex 18,400	n/a	Acquired at 250,000 (Feb 2026)	3 (owned Colossus)	L-096 T1; L-047 T2
Safe Superintelligence	7,000 raised (PB) or 8,000 (press) incl. Nvidia 5,000 in Jul 2026 (C-08 frozen)	Zero	40	32,000	1 (research only)	L-097 T2
Thinking Machines Lab	2,000 seed (Jun 2025) at 8,000 pre / 10,000 post (press 12,000; C-09); seeking 1,000 at 40,000 pre, in talks [VERIFY]	"At least a few hundred million" annualized; PB 100 FY2026 projection	~200	10,000-12,000 (seed)	1	L-098 T2
Reflection AI	2,155 raised; Series C 2,500 at ~25,000 pre in progress	n/a	~200	8,000 (Series B)	1	L-099 T2
Mistral AI	7,492 raised incl. 830 debt; Series D EUR 3,000 (Sep 8, 2026)	PB 1,167 FY2026 projection	~900	~24,374 (EUR ~21B post)	1 → 2 (first French datacenter)	L-100 T2
Periodic Labs	300 seed at 1,000 pre; talks at ~7,000 [VERIFY]	n/a	48	1,300 (seed)	1	L-101 T2

Entrant	Capital raised or spent	Revenue	Headcount	Valuation	Resembles scenario	Rows, tier
Humans&	480 seed at 4,000 pre / 4,480 post	n/a	20-30	4,480	1	L-102 T2
Meta Superintelligence Labs	2026 capex 125,000-145,000 (2025: 72,200); Scale AI 14,300 for 49%; packages to 100M over 4 yrs; 6 GW AMD deal	n/a (inside Meta)	n/a	n/a	3 (hyperscaler-funded)	L-095, L-094 T2; L-060 T1
Anthropic (incumbent, for scale)	Equity 124,254; documented contracts 324,100 (reported-$ 524,100 incl. the unverified Google leg)	Run-rate 65,000 gross; Q1 4,730, Q2 11,600; first positive adjusted operating income (non-GAAP)	5,000 (PB, stale; ~4,020 per Revelio, C-18)	965,000	Beyond 2	L-007, L-032, L-151, L-033, L-034, L-152, L-155, L-008, L-091, L-005
OpenAI (incumbent, for scale)	Equity 181,216.5; documented contracts 480,400	Run-rate >40,000; H1 12,400; H1 operating loss 21,600	4,500 → 8,000 plan	852,000	Beyond 3 (leases, custom silicon)	L-022, L-065, L-066, L-027, L-090, L-019

The reference class says three things at once, and the greenfield numbers have to be consistent with all three. Two to eight billion dollars buys a credible model and almost no revenue: Safe Superintelligence has raised $7.0B by PitchBook's count or $8B by the press's, 40 people, no product (C-08; L-097); Thinking Machines turned a $2B seed into "at least a few hundred million dollars" of annualized revenue and is in talks, not closed, at $40B pre-money (L-098); Reflection has raised $2.2B and is raising $2.5B more at about $25B (L-099); Mistral has raised $7.5B including debt and closed a EUR 3B Series D on September 8 against a $1.2B PitchBook projection for 2026 (L-100). Forty billion dollars of accumulated losses buys vertical integration and about $3B of revenue: SpaceX's S-1 shows a $41.3B accumulated deficit at March 31, 2026, with the AI segment earning $3.2B in 2025 on a $6.4B operating loss (L-096, L-047). And the incumbents raised $124B and $181B to reach $65B and $40B of run-rate. The Lean scenario lands in the first band and the workbook's check prints PASS; Full frontier must land below what Anthropic burned to its first positive adjusted quarter and near xAI's deficit by year four or five; Vertically integrated at or above xAI's deficit plus its 2025 AI-segment debt, about $57B, by year five; anything outside its band prints "outside reference class: explain".

One constraint applies to the third scenario before any of its arithmetic. Broadcom expects AI semiconductor revenue of about $58B in fiscal 2026, $115B in 2027 with supply "secured", and $230B in 2028 in its line of sight, across six custom-silicon customers including Google, Anthropic, OpenAI and Meta (earnings call; confirmed, T1; L-144). Custom-silicon supply for a new entrant is booked through 2028 by the incumbents: an entrant buys merchant GPUs at merchant prices, or waits, and the vendor claim that an XPU costs "half the cost of a GPU" (L-050) describes capacity the entrant cannot get.

The greenfield verdict, one sentence per scenario, each fenced by its reference class. A lean fast-follower can buy a frontier-class model for $2.5-4B and reach $1B of net revenue for $5-8B with perhaps one chance in three, which is what SSI, Thinking Machines and Reflection are pricing at $8-40B valuations before the revenue exists; up-rounds are prices, not returns. A full-frontier entrant needs $12-18B to its first frontier model and $60-100B or more to breakeven, roughly what Anthropic spent to its first positive adjusted quarter and less than the $124B it raised, on terms an entrant in 2026 will not get from vendors who have already chosen their customers. A vertically integrated entrant needs $80-150B, cannot buy custom silicon before 2029, and has one precedent, xAI, which reached $3B of revenue on $41B of deficit and was absorbed into SpaceX at $250B.

§13 Where the thesis got weaker, where it got stronger, the verdict, and the falsifiers

The bull case deserves to be stated at full strength before it is answered, because most of it is true. Run-rate growth is without precedent: Anthropic from about $9B to $65B in seven months (L-139, L-032), OpenAI from more than $20B to more than $40B with July up more than 20% in a month (L-065) and enterprise run-rate up 50% quarter to date (L-069), every revenue plan in the ledger beaten within months (L-120, L-082, L-009, L-024). The unit economics have turned: positive adjusted operating income at Anthropic with compute cost per revenue dollar falling from $0.71 to $0.56 in one quarter (L-034, L-151, L-152), Fable 5.1 at 25-45% lower cost per task (L-085), serving stacks at about $0.22 per million output tokens on-demand (L-164), custom XPUs at half the cost of a GPU (L-050), cache reads down 75% (L-084), $230B of Broadcom AI silicon in sight for fiscal 2028 (L-144). Capacity is financed by vendors at near-investment-grade rates: 5.75% on SPV debt (L-044), a $105B Nvidia guarantee (L-061), Nvidia on the Beacon Point lease (L-137, L-167), AMD credit support (L-148), warrants instead of cash (L-060, L-149). The commitments are options: SpaceX cancellable on 90 days (L-046), AMD purchases vesting by milestone (L-060), AWS "up to" (L-064), Cerebras's second gigawatt an option (L-149), Broadcom's gigawatts "line of sight" (L-049); the stack is a menu, not a bill. Equalized, the multiples are ordinary for the growth, 24.6x and 21.3x against bankers citing Palantir at 53x 2026 revenue (L-036). And no frontier lab with more than $1B of revenue has failed, while the $2-8B entrants are all up-rounding (L-097-L-100).

The answers, point by point. Growth is real and §2 says so; the same ledger says cost plans rose faster than revenue plans (L-062, L-078, L-082) and margin plans were missed at both labs (L-083, L-119). The Q2 result is one quarter, a projection for the amount, a non-GAAP measure with an unpublished definition, about 5.1% of revenue on a basis whose gross margin is 31-48% at plausible opex (§4); Reuters says margins are "still being pressured by enormous spending on computing power, model training and hiring" (L-150); the vendor cost claims are single-source (L-085, L-050); the $0.22 is an efficient-MoE benchmark at 100% utilization, not a frontier-model cost of goods (L-164). The subsidy is real and priced; it is also finite on the one datapoint that exists (Huang, L-075) and lives on the vendors' books (L-147, L-146), which is why it is not in the labs' prices and why the S-1s will not show it (§6). Optionality cuts both ways: an option the lab does not exercise is capacity it does not have, and the growth case needs the capacity; §5 asks which of margin, training or take-or-pay gives. The equalized multiples depend on a haircut disputed by thirteen points (C-13), and the one new datapoint, a T4 compute ratio, points to a payout share below 39.75%, which would move Anthropic to a discount (§2); the report does not claim the marks are expensive, it claims the basis risk is unpriced. And the reference class also contains xAI, $41.3B of deficit for $3.2B of revenue (L-096), and the incumbents themselves; up-rounds are prices, not returns (§12).

Where the thesis got weaker:

- Anthropic's positive adjusted operating income in Q2 2026 at $11.5-11.6B of quarterly revenue, projected at $559M or about 5.1%, on a non-GAAP measure that excludes stock-based compensation and whose definition is unpublished (L-034, L-151, L-152): real counter-evidence, less than "first operating profit" implied, and not "no profit".
- Revenue beating every plan at both labs, Anthropic's quarterly ladder now T2 at every step, $787M to $4,730M to more than $11.5B (L-151), and OpenAI's run-rate doubled since the end of 2025 (L-065).
- Vendor-backstopped financing and unit-cost deflation: the SPV at 5.75% (L-044), Nvidia's $105B guarantee (L-061) and its Beacon Point lease (L-137, L-167), XPUs at "half the cost of a GPU" and Fable 5.1 at 25-45% lower cost per task (L-050, L-085; both vendor claims), and serving stacks at about $0.22 per million output tokens on efficient models (L-164, T2; a 100%-utilization benchmark, not a frontier COGS).

Where it got stronger:

- OpenAI's first-half operating loss of $21.6B on $12.4B of revenue, 1.74x revenue and widening from $9.3B to $12.3B quarter on quarter (L-066), with the compute plan revised up twice, from $600B to $750B (L-078, L-062), and cumulative burn raised by $111B in one revision (L-082).
- The incompatibility inside Anthropic's own plan (§5): at the reported Google value the 2028 stack exceeds the plan's envelope by $25-28B a year before AMD and before Broadcom's line of sight; on the proxy, by $30-50B for the contracted gigawatts and $125-175B with the line of sight.
- Margin plans missed at both labs while revenue plans were beaten (L-083, L-119): the pattern the base rate in §10 predicts, and the one the 77% target has to break.

The verdict. The thesis survives as a statement about commitments and margins, not about profits. "Structurally incompatible with investable returns" was falsified at the level of an adjusted operating quarter and was never tested at the level of the contracts, because no filing has yet put a price on them. The mispricing this report names is two things, neither of them "too expensive": the basis risk in §2, twenty points of relative multiple the S-1 revenue-recognition note will resolve, and the triad in §5, three numbers the market is pricing together that cannot all be true. The market moved since April, and the report does not say it was right. It says what the prices still do not carry, before the S-1s do.

Four falsifiers, each dated:

1. Anthropic's public S-1 commitments note listing the Google and AMD gigawatt deals as consumption-based or optional, with no minimum purchase amounts by year. That kills §5's strongest exit. Date: the public flip, which an October pricing requires by roughly late September (15-day rule; L-001).
2. An Anthropic FY2026 gross margin at or above 60% on a gross basis, or a GAAP operating profit for the second quarter, in the S-1's historical financials. That kills the margin exit and restores the Q2 counter-evidence to full strength. Date: the same filing.
3. OpenAI's third-quarter operating loss narrowing quarter on quarter. That weakens §4 and the first "stronger" bullet. Date: on the pattern of the Q2 leaks, six to seven weeks after quarter-end, so mid-November.

4. Either lab pricing an IPO above its standing mark with the obligation table in the prospectus. That tests §0's claim that the cost architecture is not in the prices. Date: October for Anthropic if the reported timetable holds; Q4 2026 or 2027 for OpenAI.

The IPO windows as they stand on September 9. Neither company has a public S-1 on EDGAR (confirmed, T1; L-001, L-002). Anthropic submitted its draft confidentially on June 1 (L-039); PitchBook expects October (L-003); the prospectus reported due after Labor Day has not appeared and the latest report puts the listing at mid-October (T3; L-040). OpenAI filed confidentially on June 8 (L-004); its chief financial officer told employees on August 19 that it "will be a public company in 2027", or sooner if "our business continues to inflect" (L-068); PitchBook's September window is closed by arithmetic, and the live question is Q4 2026 against 2027 (C-02, residual). Prediction markets gave Anthropic a 72% chance of listing first (T4; colour only; L-141). The first prospectus to flip replaces the first six of the ten figures this report turns on with audited numbers.

§14 Frozen conflicts, methodology rulings, provenance

Nine conflicts stay frozen, both branches in the text and a named switch on Workbook A's 12_Conflicts tab; two carry addenda from the second pass, and one house-rule item is added.

Frozen conflicts and switches (Workbook A / 12_Conflicts, column C)

Conflict	Branch A	Branch B	How the report renders it	Switch
C-02 (residual)	OpenAI IPO in Q4-2026 "if our business continues to inflect" (L-068)	2027 (L-068); PitchBook's "September 2026" (L-004) retired by L-002 and the 15-day rule	The September window is closed by arithmetic; the live question is Q4-2026 vs 2027	SW_IPO_OAI
C-03	Anthropic FY2025 recognized revenue $10,000M (PB, L-010)	~$4,500-6,000M implied by the Jan-2026 guidance (L-120, T4); $9,000M is the exit run-rate (L-139)	PitchBook's FY2025 field equals the December exit run-rate; recognized revenue is lower and undisclosed	SW_A_FY25
C-05	FY2024 net loss −$5,300M (Jul-16 PitchBook vintage)	−$8,300M (PB deal records, L-012)	Both printed with vintages	SW_A_FY24_NI
C-08	SSI Apr-2025 round $6,000M, total raised $7,000M (PB)	$2,000M round, total $8,000M (press)	Reference class only; both printed	SW_SSI (Workbook B)
C-12, with addendum	Anthropic "expects to lose ~$11B in 2026 and 2027" (Jan-2026, L-120, T4; a GAAP-style plan figure)	Q2-2026 positive ADJUSTED operating income (L-151, T2; non-GAAP, definition unpublished; includes training cost and excludes SBC per T3 relays, L-152); "2026 planned burn ~$3B" (T3 recalled)	Definitions and vintages differ (GAAP loss incl. SBC and D&A; operating cash burn; a non-GAAP adjusted income of unpublished scope); H2 training load could reconcile them; unresolved	SW_A_2026_LOSS
C-13	39.75% equalization (Ruling 5, L-126)	~27% (OpenAI CRO memo relay, adversarial, T4); the T4 compute ratio $0.56 (L-152) cannot coexist with a 39.75% payout in cost of revenue (§2)	Base case 39.75% with both branches printed wherever a net figure appears; the L-152 corollary as a flag	SW_HAIRCUT

Conflict	Branch A	Branch B	How the report renders it	Switch
C-15	Secondary at $925 a share, $1.50T implied (T4)	"Low-to-mid $800B" (T4)	Secondary prints disagree; the primary mark is $965B	display-only cell
C-16	OpenAI's >$40B is net of the Microsoft share (register convention)	Basis unstated; could be gross (Bloomberg caveat, L-065, L-158)	Base case net; gross case shown with a T4 20% haircut, flagged	SW_OAI_BASIS
C-18	Anthropic headcount 5,000 (PB, Apr 21, L-008)	~4,020 (Revelio, March, L-091, T4)	Both printed; neither company-stated	SW_A_HC
C-14 addendum (Google leg)	Google leg unpriced in every filing (L-153, L-154, L-156, T1)	~$200B over five years, ">40%" of Google's backlog (The Information via Reuters, L-155, T3, page not opened)	Own scope row in Exhibit 8 with [VERIFY]; never inside documented-$; §5 shows the reported branch beside the proxy branches	SW_GOOGLE_VALUE
$15B facility (house rule)	"Set to finalize" (Bloomberg Sep 3, L-042, T3)	No close reported as of Sep 9; PB still "Upcoming" (L-160)	[VERIFY] until closed; default "does not close"; CE-4 shown at 7.5 and 7.0	SW_15B_FACILITY

Resolved by the ledger and stated as resolved: C-01 (the $34.5B is lessor debt; equity-only unaffected), C-04 (PitchBook's −$42B net income on both companies is an artifact; neither used), C-06 (PitchBook's $20B is run-rate vintage; FY2025 is $13.1B), C-07 (Fast Mode is 2x standard pricing; "3x cheaper" is retired), C-09 and C-10 (minor; a PB date artifact), C-11 (basis label), C-14 (scope table), C-17 (company statement governs), C-19 (about 27% at the recap, T2; below 27% now, undisclosed; L-157, L-159), C-20 (the Nvidia LOI is superseded), O-15 (the "Co-CEO" title was a PitchBook artifact; no OpenAI co-CEO structure; L-028, L-070, L-089).

The six rulings that bind every sentence. Ruling 1: the capital-efficiency denominator is equity raised only; no debt, no lease. Ruling 2: the quality-valuation correlation coefficient is embargoed; the $/AIBQ-point ladder is the only permitted expression. Ruling 3: Anthropic's free cash flow stays unswept; an adjusted operating income is not an operating profit and neither is free cash flow. Ruling 4: PitchBook TTM revenue fields are forward projections, never current figures. Ruling 5: Anthropic's gross-to-net equalization is 39.75%, gate-verified on the May 2026 mix, with the ~27% branch carried beside it. Ruling 6: AIBQ v3.0 weights are CE 20 / RQ 25 / CI 15 / GO 20 / MD 20.

The ten load-bearing figures, with as-of dates (Workbook A / 01_Data, LB01-LB10)

#	Figure	Value	As-of	Basis	Row	Status	Tier
1	Anthropic annualized run-rate	$65.0B	end-Jul 2026 (reported Aug 17)	GROSS run-rate incl. cloud-partner resale; company investor update via CNBC, Bloomberg, TechCrunch	L-032	confirmed	T2
2	OpenAI annualized run-rate	>$40.0B	Jul 2026 performance (reported Aug 13-14)	Basis not stated (register convention: net of the Microsoft share); +35% QTD on Aug 19	L-065, L-069	confirmed	T2

#	Figure	Value	As-of	Basis	Row	Status	Tier
3	Anthropic equity-only capital raised	$124,254M	2026-05-28	Sum of ten PB equity rounds incl. converted convertibles; excludes the $2.5B revolver and the $34.5B lease SPV	L-007	estimated (derived)	T2
4	OpenAI equity-only capital raised	$181,216.5M	2026-03-31	Sum of PB equity rounds; excludes $5,220M of debt and $1,030M of grants	L-022	estimated (derived)	T2
5	Anthropic Q2-2026 revenue and first positive ADJUSTED operating income	$11.5-11.6B revenue (Q1 $4,730M); adjusted operating income projected $559M ≈ 5.1% on $10.9B, sign confirmed	Q2 2026 (reported Aug 14-19)	Preliminary recognized revenue, gross presumed; the profit measure is non-GAAP, definition unpublished; per T3 relays includes training cost and excludes SBC	L-033, L-034, L-151, L-152	confirmed (revenue, sign) / estimated (amount, definition)	T2 (T3 for the definition)
6	OpenAI Q2-2026 revenue and operating loss	$6.7B revenue (Q1 $5.7B); operating loss $12.3B incl. SBC (Q1 $9.3B)	Q2 2026 (WSJ Aug 19)	Recognized quarterly, net presumed	L-066	confirmed	T2
7	OpenAI compute spend, 2026 and through 2030	$50B (2026); ~$750B (2026-2030 plan, up from ~$600B)	2026-05-05 (sworn) / 2026-07-22	Company statement in court; WSJ on the plan	L-078, L-062	confirmed / estimated	T2
8	Microsoft revenue from OpenAI arrangements	$24.1B FY2026; A/R $6.0B	2026-06-30 (10-K filed Jul 29)	Azure consumption plus revenue sharing received, at Microsoft	L-054	confirmed	T1
9	SB Energy PORTS-Pike leases and the Nvidia guaranty	17 leases, ~8.0 GW-IT, 20 yr, OpenAI tenant; Nvidia RVG $105B on ~4.25 GW-IT; "OpenAI is not an investment-grade tenant"	2026-08-17 (S-1 filed Sep 1)	Issuer filing	L-061	confirmed	T1

#	Figure	Value	As-of	Basis	Row	Status	Tier
10	Anthropic 2028 revenue forecast to IPO investors	$190-200B	FY2028E (Aug-2026 vintage)	Company forecast per two sources familiar (Reuters exclusive); presumably gross	L-036	estimated	T2

Frame constants every figure is divided by: Anthropic's $965B post-money (L-005, L-038; confirmed; May 28, 2026); OpenAI's $852B (L-019, L-020; confirmed; March 31, 2026; re-affirmed by the August 10 tender, L-067); the 39.75% equalization (L-126; recalled, T4; Ruling 5) with its ~27% branch (C-13). Only the eighth and ninth of the ten are T1; the run-rates are company communications relayed by T2 outlets, the Q2 prints were leaked, and the equity sums are derived from PitchBook deal records. The S-1s will replace figures 1 through 6 with audited numbers. This is the last report written before they do, and it says so here, once.

Provenance. Tiers: T1 primary or audited; T2 reputable, official, or PitchBook fields; T3 single-outlet scoop or analyst estimate; T4 aggregator, secondary, or internal prior; T5 rumor. Status: confirmed (T1/T2 with a verbatim sentence read on the page); estimated (modeled, derived, or forward); recalled (prior internal register, T4 at most); could-not-verify. Decay: VOLATILE, QUARTERLY, STABLE, PERMANENT. Ledger: 168 rows (L-150 to L-168 appended in a second bounded pass on September 9, 2026; existing rows untouched); Gate 1 passed twice; confirmed 97, estimated 51, recalled 6, could-not-verify 14; L-161 and L-165 tagged estimated/T3 by rule. Rigor dial 5: a figure resting on a T3 or T4 row ships only with the flag in the sentence. Tool limits this run: NEXUS and Notion not available (prior extracts used as T4 only); Bigdata.com not exposed; Aiera not connected; cnbc.com, forbes.com, axios.com, theinformation.com and blogs.microsoft.com blocked to the fetcher, so syndications and PitchBook news chunks were used. Analyst judgment is labelled "AJ" or "analyst judgment" wherever it appears; no figure was filled from memory. Disclosure. Every input is public or licensed: SEC filings, company statements, PitchBook subscription data, and press reporting; no non-public information from either company, its bankers or its investors was used, and nothing from either confidential draft registration statement appears here. The PitchBook model-portfolio allocation referenced in prior notes carried Anthropic at 6% overweight and OpenAI at 2% underweight as of the July 16, 2026 register; personal holdings, if any, are disclosed under PitchBook research policy before external distribution. Author: Harrison Rolfes, Senior Research Analyst, PitchBook. Version 5.0, September 9, 2026.