

LATE-STAGE COMPANY RESEARCH **Ranking the AI Giants: A New Framework for the Frontier Five**

Where the models end
and the business begins



Contents

Key takeaways	3
How to read this report	4
Note on methodology	4
AI Business Quality (AIBQ) scorecard	5
Part I: The AIBQ framework	6
Part II: Market context	15
Part IIIA: Company profiles: The Frontier Five	34
Databricks: The enterprise foundation	35
Anthropic: The institutional bet	41
xAI: The infrastructure gambit	51
OpenAI: The first-mover paradox	60
Safe Superintelligence (SSI): The pure research bet	69
Part IIIB: Comparative analysis	78
Part IIIC: Exit path analysis	86
Part IV: Investment implications	94
Five conclusions	98
Appendix	100

Institutional Research Group



Harrison Rolfes
Senior Research Analyst,
Late-Stage Research
harrison.rolfes@pitchbook.com

pbinstitutionalresearch@pitchbook.com

Published on March 4, 2026

Key takeaways

- **The valuation-quality paradox is real and intensifying:** Our AIBQ framework scores Databricks (with a \$134 billion valuation)¹ an 8.7 out of 10 and Anthropic (with a \$380 billion valuation)² a 7.4 out of 10, yet Databricks commands a fraction of xAI's \$250 billion valuation (with a score of 5.4)³ and OpenAI's \$840 billion valuation (with a score of 4.8),⁴ and Anthropic's completed \$380 billion round now exceeds xAI's acquisition price despite Anthropic having a materially higher AIBQ score. OpenAI's \$110 billion funding round widened the paradox dramatically: Its valuation increased 68%, its capital efficiency ratio collapsed from 0.31x to about 0.11x, and its overall AIBQ score dropped 0.4 points. The market now pays \$175 billion per AIBQ point for OpenAI versus \$15.4 billion for Databricks, an 11.4x premium per unit of business quality, up from 6.7x before the round. The market currently prices narrative and optionality, whereas AIBQ measures business fundamentals. For investors seeking sustainable returns, this gap creates a significant opportunity.
- **Revenue quality matters more than revenue quantity:** Based on estimated quality metrics, one dollar of Anthropic revenue (Anthropic has an NRR over 140%, an 80% enterprise revenue concentration, and multiyear contracts)⁵ is worth roughly \$2.20 of OpenAI revenue (OpenAI has an estimated NRR around 115%, a roughly 60% consumer revenue concentration, and shorter contract terms).⁶ On our quality-adjusted framework, Anthropic trades at 27.1x

current ARR (\$380 billion/\$14 billion) versus OpenAI's 42x (\$840 billion/\$20 billion)—a premium for OpenAI that has inverted since its \$110 billion round, with OpenAI now trading at a 55% multiple premium to Anthropic despite a materially lower revenue quality score.⁷ Claude Code's acceleration to \$2.5 billion in ARR, with enterprise users representing more than half of its revenue and business subscriptions quadrupling since January 1, 2026, further validates Anthropic's revenue quality advantage. Recent enterprise customer disclosures across the cohort further support the divergence in revenue-quality metrics.

- **Databricks is systematically misclassified:** Consensus places Databricks alongside frontier AI labs; however, Databricks is primarily a data infrastructure company with AI capabilities. This distinction makes it the most traditionally investable company in the Frontier Five, having already achieved positive FCF with gross margins over 80%, an NRR over 140%, and IPO-ready governance.⁸ Databricks is also the only Frontier Five company completely insulated from the Anthropic-Pentagon standoff, with no government contract dependency, no political risk exposure, and no involvement in the broader debate over military AI use restrictions. Its capital efficiency ratio of 0.16x now leads the cohort following OpenAI's ratio compression, and its AIBQ lead over the second-ranked company has widened to 1.3 points (8.7 versus Anthropic's 7.4), up from 1.1 points before the crisis.

1: Databricks' \$134 billion valuation is tied to its Series L, a \$7 billion raise including \$5 billion in equity and \$2 billion in debt, on February 9, 2026. Its AIBQ score was revised from 8.5 to 8.7 to reflect its updated capital efficiency ratio (0.16x) following \$33.1 billion in cumulative capital raised. "Databricks Grows >55% YoY, Surpasses \$4.8B Revenue Run-Rate, and Is Raising >\$4B Series L at \$134B Valuation," Databricks, December 16, 2025.

2: Anthropic achieved a \$380 billion post-money valuation with its \$30 billion Series G funding round completed on February 12, 2026. "Anthropic Raises \$30 Billion in Series G Funding at \$380 Billion Post-Money Valuation," Anthropic, February 12, 2026.

3: xAI achieved a \$230 billion valuation with its Series E in January 2026; it was acquired by SpaceX at \$250 billion on February 2, 2026. "Elon Musk's xAI Raises \$20 Billion From Investors Including Nvidia, Cisco, Fidelity," CNBC, Lora Kolodny, January 6, 2026.

4: OpenAI's \$840 billion valuation reflects its \$110 billion round in February 2026.

5: "Anthropic Raises \$13B Series F at \$183B Post-Money Valuation," Anthropic, September 2, 2025. See also: "Anthropic Lowers Gross Margin Projection as Revenue Skyrockets," The Information, Sri Muppidi, January 22, 2026. Reuters confirmed Anthropic's enterprise revenue concentration: "Exclusive: Anthropic Aims to Nearly Triple Annualized Revenue in 2026, Sources Say," Reuters, Krystal Hu and Deepa Seetharaman, October 15, 2025.

6: OpenAI's net retention rate (NRR) and consumer/enterprise mix are estimated; the company does not publicly disclose its NRR. Enterprise customers account for roughly 40% of OpenAI's business, according to CFO Sarah Friar ("OpenAI, Anthropic Set Sights on Enterprise Customers at Davos," CNBC, Ashley Capoot, January 21, 2026), implying consumer subscriptions represent approximately 60% of revenue. Deutsche Bank Research flagged stalling European ChatGPT subscription growth in October 2025 ("OpenAI's 'Flatlining' Subs in Europe: The AI Boom's Poster Child May Be Struggling to Recruit New Subscribers, Deutsche Bank Warns," Fortune, Jim Edwards, October 14, 2025), noting only about 5% of weekly active users convert to paid subscriptions. Revenue estimates are from Sacra. "OpenAI," Sacra, n.d., accessed February 18, 2026.

7: The revenue multiple comparison is based on the most recent valuations and disclosed annual recurring revenues (ARRs).

8: "Databricks Grows >55% YoY, Surpasses \$4.8B Revenue Run-Rate, and Is Raising >\$4B Series L at \$134B Valuation," Databricks, December 16, 2025.

How to read this report

- Part I introduces the AIBQ framework—the analytical lens applied throughout the report.
- Part II provides market context: sizing, capital formation, and competitive dynamics.
- Part III profiles each company in order of AIBQ score (highest to lowest), applying consistent analysis to enable direct comparison.
- Part IV synthesizes investment implications with actionable triggers.

Note on methodology

- PitchBook is the primary source for the data in this report.
- All financial data was accessed between January 15 and February 28, 2026, unless otherwise noted. All closed-round data is as of February 28, 2026. Secondary market pricing is as of February 19 or 20, 2026.
- Material developments that occurred on February 27, 2026, including OpenAI's \$110 billion funding round, the presidential executive order banning federal use of Anthropic's technology, and xAI's classified-systems agreement with the Pentagon, have been incorporated into the analysis and are noted throughout.
- Post-money valuations as of February 2026:
 - OpenAI's \$840 billion valuation reflects its \$110 billion funding round closed on February 27, 2026, with Amazon (\$50 billion), NVIDIA (\$30 billion), and SoftBank (\$30 billion) as lead investors. The round remains open for additional participants. OpenAI's prior valuation of \$500 billion, based on its October 2025 secondary transaction, is used as the baseline for comparative purposes where noted. Additionally, data on OpenAI's classified-systems agreement with the Pentagon is based on CEO Sam Altman's February 27, 2026, statement. This agreement includes stated prohibitions

on domestic mass surveillance and human responsibility requirements for autonomous weapons, with forward-deployed OpenAI engineers embedded at the Pentagon.

- Anthropic's \$380 billion valuation is from its Series G completed on February 12, 2026. Its pre-money valuation was \$350 billion.
- xAI's \$250 billion valuation is from its acquisition by SpaceX on February 2, 2026. Prior to its acquisition, xAI achieved a \$230 billion valuation in its Series E.
- Databricks' \$134 billion valuation is from its Series L on February 9, 2026.
- SSI's \$32 billion valuation is from its \$2 billion round in April 2025.
- Data on the SpaceX-xAI acquisition is based on February 2, 2026, reporting.⁹
- Anthropic's NRR and enterprise mix are estimated from The Information reporting;¹⁰ the company does not publicly disclose this information. Its enterprise customer count and ARR of \$14 billion are referenced in its Series G announcement on February 12, 2026.¹¹ Claude Code's \$2.5 billion ARR and the quadrupling of business subscriptions since January 1, 2026, are based on company disclosures during the Series G period.
- OpenAI metrics are estimated; the company does not publicly disclose its NRR. OpenAI's cumulative capital raised on roughly \$174 billion includes the \$110 billion round closed on February 27, 2026. Revenue projections of \$280 billion by 2030 and computing spending guidance of \$600 billion by 2030 are based on investor materials disclosed in connection with the round. Codex usage (1.6 million weekly users and 9 million paying business users) is based on company disclosures.
- Databricks' gross margins, NRR, and free cash flow trajectory are based on company investor presentations and press releases.¹²

9: SpaceX acquired xAI on February 2, 2026, in an all-stock merger valuing the combined entity at \$1.25 trillion. "Elon Musk's SpaceX Officially Acquires Elon Musk's xAI, With Plan to Build Data Centers in Space," TechCrunch, Sean O'Kane, February 2, 2026.

10: "Anthropic Lowers Gross Margin Projection as Revenue Skyrockets," The Information, Sri Muppidi, January 22, 2026.

11: "Anthropic Raises \$30 Billion in Series G Funding at \$380 Billion Post-Money Valuation," Anthropic, February 12, 2026.

12: "Databricks Grows >55% YoY, Surpasses \$4.8B Revenue Run-Rate, and Is Raising >\$4B Series L at \$134B Valuation," Databricks, December 16, 2025.

AI Business Quality (AIBQ) scorecard



CE = Capital efficiency (20%) RQ = Revenue quality (25%) CI = Computing independence (15%) GO = Governance optionality (20%) MD = Moat durability (20%)

Company	Post-money valuation (\$B)	ARR (\$B)	Multiple	AIBQ score	Capital efficiency	Revenue quality	Computing independence	Governance optionality	Moat durability
Databricks	\$134.0	\$5.4	24.8x	8.7	9	9	8	9	8
Anthropic	\$380.0	\$14.0	27.1x	7.4	7	9	5	7	8
xAI	\$250.0	\$3.2	78.1x	5.4	4	4	9	5	6
OpenAI	\$840.0	\$20.0	42.0x	4.8	4	5	5	4	6
SSI	\$32.0	\$0.0	∞	2.3	1	1	3	5	2

The valuation-quality paradox: AIBQ scores correlate inversely with valuation.

OpenAI (\$840 billion, 4.8) and xAI (\$250 billion, 5.4) are overvalued on fundamentals.
Databricks (\$134 billion, 8.7) and Anthropic (\$380 billion, 7.4) offer superior quality-adjusted returns.

The market prices narrative; the AIBQ framework measures the business.

Dimension leaders

Dimension	Company	AIBQ score
Capital efficiency	Databricks	9
Revenue quality	Databricks	9
Computing independence	xAI	9
Governance optionality	Databricks	9
Moat durability	Databricks	8

Portfolio allocation

Company	Portfolio allocation	Risk
Databricks	8%	Low-medium
Anthropic	6%	Medium
xAI	3%	Medium
OpenAI	2%	Medium-high
SSI	<1%	Very high
Total AI allocation	~20%	

Note: Companies are ordered by AIBQ score. Anthropic's \$380 billion post-money valuation is from its Series G completed on February 12, 2026 (its pre-money valuation was \$350 billion).¹³ ARR is used for valuation multiples; ARR figures are comparable for enterprise companies.¹⁴ xAI's \$250 billion valuation is from its acquisition by SpaceX on February 2, 2026; prior to its acquisition, it achieved a \$230 billion valuation in its Series E in January 2026.¹⁵ xAI's standalone Grok API revenue is about \$500 million; ecosystem revenue including X Premium+ attribution approaches \$3.2 billion. xAI acquired X on March 28, 2025.¹⁶ OpenAI closed a \$110 billion funding round on February 27, 2026, bringing its post-money valuation to \$840 billion.

13: "Anthropic Reportedly Raising \$10B at \$350B Valuation," TechCrunch, Rebecca Bellan, January 7, 2026.

14: "Anthropic's Revenue Run Rate Tops \$9 Billion as VCs Pile In," Yahoo! Finance, Bloomberg, Natasha Mascarenhas, Shirin Ghaffary, and Dawn Lim, January 21, 2026.

15: "xAI Raises \$20B Series E," xAI, January 6, 2026.

16: "Musk's Social Media Firm X Bought by His AI Company, Valued at \$33 Billion," Yahoo! Finance, Reuters, Greg Bensinger, March 28, 2025.

Part I: The AIBQ framework

We introduce the AIBQ framework as a standardized methodology for evaluating frontier AI companies across five dimensions that determine long-term enterprise value creation. The AIBQ framework is a subset of the PBQ: a granular methodology to value late-stage private companies across multiple industries. Unlike traditional valuation metrics, the AIBQ framework accounts for the unique characteristics of AI businesses: capital efficiency, revenue quality, computing independence, governance optionality, and moat durability.

Scope note: Why not Google, Meta, or Amazon?

Google (Gemini), Meta (Llama), and Amazon (via Amazon Web Services) are significant frontier AI players that appear throughout this report’s competitive benchmarking. They are excluded from the Frontier Five for a specific analytical reason: They are publicly traded conglomerates whose AI operations cannot be independently valued, invested in, or exited. The AIBQ framework is designed for private market investors evaluating standalone AI companies where capital-allocation decisions—entry price, position sizing, and exit timing—are direct. Google’s \$2 trillion+ market cap reflects search, advertising, cloud, and hardware alongside Gemini; applying the AIBQ framework to a segment that may represent 5% to 10% of enterprise value would produce misleading scores. The Frontier Five are selected as the most significant privately held, venture-backed frontier AI model developers as of February 2026, offering direct private market investment exposure to AI.

Changes since last edition

This is the inaugural edition of this report. Future editions will include a “Changes since last edition” summary tracking AIBQ score movements, valuation changes, capability shifts, and material corporate events. Readers should expect quarterly updates as new funding rounds, benchmark releases, and financial disclosures become available. The AIBQ framework may be refined as the frontier AI market matures and new evaluation dimensions emerge.

AIBQ overview

Company	Post-money valuation (\$B)	ARR (\$B)	Multiple	AIBQ	Position	Maximum size	Risk
Databricks	\$134.0	\$5.4	24.8x	8.7/10	Core	8%	Low-medium
Anthropic	\$380.0 ²¹	\$14.0 ²²	27.1x	7.6/10	Overweight	7%	Low-medium
xAI	\$250.0	\$3.2 ²³	78.1x	5.4/10	SpaceX proxy	3%	Medium
OpenAI	\$500.0	\$20.0	25.0x	5.2/10	Underweight	2%	High
SSI	\$32.0	\$0.0	Infinity	2.3/10	Optionality	<1%	Very high

Why traditional metrics fail

Standard software-as-a-service (SaaS) metrics—such as annual recurring revenue (ARR) growth, NRR, and the Rule of 40—were designed for businesses with gross margins over 70%, minimal infrastructure requirements, and straightforward corporate structures. Frontier AI companies violate all three assumptions:

OpenAI is projected to burn approximately \$14 billion in 2026, up from around \$8 billion to \$9 billion in 2025.^{17, 18} Its \$110 billion funding round brings its cumulative capital raised to approximately \$174 billion against \$20 billion in ARR, a 0.11x efficiency ratio that no traditional SaaS framework is equipped to evaluate.

Anthropic’s gross margins reportedly swung from -94% in 2024 to approximately 40% in 2025 (revised downward from an earlier 50% internal projection in January 2026, after inference costs came in about 23% above expectations).¹⁹

xAI’s 78.1x revenue multiple on ecosystem revenue—or 500x on standalone Grok API revenue—would be absurd for any traditional technology company.²⁰

The AIBQ framework provides evaluation criteria designed for this new reality.

Sources: PitchBook, [Anthropic](#), [TechCrunch](#), [CNBC](#), [Yahoo! Finance](#), [Reuters](#), [Sacra](#) • As of February 28, 2026
Note: ARR is used for valuation multiples; ARR figures are comparable for enterprise companies. Post-money valuations are used throughout. Companies are ordered by AIBQ score, from highest to lowest.

17: “OpenAI’s First Half Results: \$4.3 Billion in Sales, \$2.5 Billion Cash Burn,” [The Information](#), Stephanie Palazzolo, Amir Efrati, and Cory Weinberg, September 30, 2025.
18: “OpenAI Says It Plans to Report Stunning Annual Losses Through 2028—and Then Turn Wildly Profitable in 2030,” [Fortune](#), Dave Smith, November 12, 2025.
19: “Anthropic Lowers Gross Margin Projection as Revenue Skyrockets,” [The Information](#), Sri Muppidi, January 22, 2026.
20: xAI’s \$250 billion valuation from its acquisition by SpaceX, divided by its \$3.2 billion ecosystem ARR (including X ecosystem contributions), would equal a 78.1x revenue multiple. “[Elon Musk’s xAI Raises \\$20 Billion From Investors Including Nvidia, Cisco, Fidelity](#),” [CNBC](#), Lora Kolodny, January 6, 2026. xAI’s ARR estimate is from [Sacra](#). “[xAI](#),” [Sacra](#), n.d., accessed February 18, 2026.
21: “[Anthropic Raises \\$30 Billion in Series G Funding at \\$380 Billion Post-Money Valuation](#),” [Anthropic](#), February 12, 2026. See also: “[Anthropic Reportedly Raising \\$10B at \\$350B Valuation](#),” [TechCrunch](#), Rebecca Bellan, January 7, 2026.
22: “[Anthropic’s Revenue Run Rate Tops \\$9 Billion as VCs Pile In](#),” [Yahoo! Finance](#), Bloomberg, Natasha Mascarenhas, Shirin Ghaffary, and Dawn Lim, January 21, 2026. See also: “[Exclusive: Anthropic Aims to Nearly Triple Annualized Revenue in 2026](#),” [Sources Say](#), Reuters, Krystal Hu and Deepa Seetharaman, October 15, 2025.
23: xAI’s ARR includes contributions from the X ecosystem following xAI’s acquisition of X on March 28, 2025. xAI’s standalone ARR is estimated to be around \$500 million. “[Musk’s Social Media Firm X Bought by His AI Company, Valued at \\$33 Billion](#),” [Yahoo! Finance](#), Reuters, Greg Bensinger, March 28, 2025. xAI’s ARR estimate is from [Sacra](#). “[xAI](#),” [Sacra](#), n.d., accessed February 18, 2026.

The five dimensions of the AIBQ framework

1. Capital efficiency (20% weight)

Capital efficiency measures how effectively a company converts capital into sustainable revenue. For this analysis, capital efficiency is calculated by dividing ARR by cumulative capital raised, adjusted for capital deployment timing, gross margin trajectory, and path to profitability.

- Score of 10: Exceptional efficiency (over 0.5x with positive free cash flow).
- Score of 1: Capital destruction with no profitability path.

This metric penalizes companies that raise faster than they monetize, regardless of growth rate. OpenAI's \$110 billion round illustrates the dynamic: Its ARR remained at \$20 billion while its cumulative capital raised increased from \$63.9 billion to approximately \$174 billion, compressing its efficiency ratio from 0.31x to 0.11x—a 65% decline that moved it from cohort-leading to trailing both Databricks (0.16x) and Anthropic (0.23x). Companies that convert capital efficiently will survive market corrections, while those that do not become cautionary tales. Regardless of the interest rate environment, capital efficiency separates companies that can self-fund from those dependent on continuous external capital—a distinction that becomes existential when private market liquidity tightens.

2. Revenue quality (25% weight)

Revenue quality is the highest-weighted dimension because it determines business durability. We evaluate NRR (expansion versus churn), the enterprise percentage of revenue, contract duration, customer concentration, and gross margin sustainability.

- Score of 10: NRR over 140%, 80% enterprise revenue concentration, multiyear contracts.
- Score of 1: Pre-revenue or entirely transactional business.

For example, a dollar of Anthropic revenue (NRR over 140%, 80% enterprise revenue concentration, multiyear contracts) is worth more than a dollar of OpenAI revenue (estimated 115% NRR, around 60% consumer revenue concentration, shorter contract duration). Anthropic's Claude Code—now at \$2.5 billion in ARR with enterprise users generating more than half of revenue—demonstrates how product-level quality metrics can reinforce company-level scores. Revenue quality determines whether current ARR translates to future value. We note that NRR figures are estimates with confidence ranges of +/-15 percentage points; the quality differential persists across plausible ranges, but its magnitude may vary.

3. Computing independence (15% weight)

Computing independence evaluates a company's ability to control its most critical input: computing. We assess cloud-provider concentration, owned infrastructure, custom silicon development, and negotiating leverage.

- Score of 10: Fully owned infrastructure with no hyperscaler dependency.
- Score of 1: Single-provider dependency with no optionality.

Computing is the oxygen of AI. Companies dependent on a single hyperscaler face margin compression, capacity constraints, and strategic vulnerability. xAI's Colossus (more than 555,000 physical GPUs,

comparable to 1 million H100 computing units)²⁴,

²⁵ represents genuine independence. Anthropic's historically heavy Amazon Web Services (AWS) dependency is now diversifying across Google Cloud (up to 1 million tensor processing units), Microsoft Azure (up to \$30 billion in computing purchases), and planned owned infrastructure, reducing single-provider risk but not yet eliminating it. OpenAI's prior \$250 billion Azure commitment represents structural constraint,²⁶ though the Amazon partnership announced alongside its \$110 billion round—a \$100 billion expansion of the existing \$38 billion AWS relationship, bringing total AWS commitments to about \$138 billion—improves multicloud optionality while deepening aggregate hyperscaler dependency. Databricks benefits from multicloud neutrality.

4. Governance optionality (20% weight)

Governance optionality assesses a company's ability to execute strategic transactions: IPO readiness, M&A optionality, and capital-raising flexibility. We evaluate corporate-structure complexity, board composition, key-person risk, and regulatory positioning.

- Score of 10: Clean structure with immediate IPO optionality.
- Score of 1: Structural impediments or single-person control.

Governance optionality determines whether a company captures value or destroys it through structural complexity. OpenAI's conversion from

capped-profit company to public benefit corporation (PBC) is a highly complex corporate restructuring. Databricks could file an S-1 tomorrow. xAI's governance optionality was effectively exercised—and eliminated—when SpaceX acquired xAI on February 2, 2026; investors now access xAI only through SpaceX equity, with no independent governance levers remaining.

The Anthropic-Pentagon standoff introduces a new viewpoint: political risk from principled governance. Anthropic's refusal to permit Claude's use for mass surveillance and fully autonomous weapons—decisions rooted in its PBC charter and responsible scaling commitments—triggered a presidential

²⁴: "xAI Raises \$20B Series E," xAI, January 6, 2026.

²⁵: Colossus GPU counts require independent verification.

²⁶: The OpenAI-Microsoft Azure partnership was restructured on October 28, 2025. Microsoft holds a roughly 27% stake (down from 32.5% pre-dilution), valued at around \$135 billion. OpenAI committed to \$250 billion in incremental Azure purchases.

"OpenAI Completes Restructure, Solidifying Microsoft as a Major Shareholder," CNBC, Ashley Capoot, October 28, 2025.

executive order banning all federal use of its technology. Hours later, OpenAI signed a classified-systems deal with the Pentagon incorporating functionally identical restrictions, and the Pentagon accepted those terms. For Anthropic, this does not reflect a decline in governance quality but a reduction in strategic flexibility: A supply chain risk designation (which Anthropic is legally challenging), Defense Production Act threats, and IPO timing uncertainty

are new constraints on optionality. The OpenAI deal demonstrates that the terms themselves were achievable—the penalty was political, not substantive—but the constraints on Anthropic’s strategic flexibility are real regardless of their origin. This case demonstrates that governance optionality must account for external political responses to governance decisions, not only internal structural readiness.

5. Moat durability (20% weight)

Moat durability evaluates the sustainability of competitive advantages. We assess brand strength, distribution relationships, capability differentiation, switching costs, and data/network effects.

- Score of 10: Multiple reinforcing moats.
- Score of 1: Purely reputational positioning with no structural advantage.

Capability gaps between frontier models have narrowed on preference benchmarks (within 5% to 15% for Elo points on Arena), though significant gaps persist on reasoning benchmarks (for example, there are roughly 24 percentage points between Claude

Opus 4.6 and the next-best model on ARC-AGI-2). The AI wars will not be won on capability alone. The Pentagon’s classified-systems deals illustrate how political dynamics can override capability advantages: xAI gained access with no restrictions despite Grok trailing Claude by 23.7 percentage points on ARC-AGI-2, while OpenAI secured access with Anthropic-style guardrails that the Pentagon had refused to grant Anthropic. Capability moats remain necessary but insufficient when procurement decisions are politically mediated. Companies with durable moats—as seen with Databricks’ data relationships, Anthropic’s enterprise trust, and OpenAI’s brand recognition—will capture disproportionate value.

AIBQ composite scores

Dimension	Databricks	Anthropic	xAI	OpenAI	SSI
Capital efficiency	9/10	7/10	4/10	4/10	1/10
Revenue quality	9/10	9/10	4/10	5/10	1/10
Computing independence	8/10	5/10	9/10	5/10	3/10
Governance optionality	9/10	7/10	5/10	4/10	5/10
Moat durability	8/10	8/10	6/10	6/10	2/10
Composite score	8.7/10	7.4/10	5.4/10	4.8/10	2.3/10

Source: PitchBook • As of February 28, 2026


Strongest and weakest AIBQ scores by company

	Databricks	Anthropic	xAI	OpenAI	SSI
Strongest	Capital efficiency (9)	Revenue quality (9)	Computing independence (9)	Moat durability (6)	Governance optionality (5)
Weakest	Computing independence (8)	Computing independence (5)	Governance optionality (4)	Capital efficiency (3)	Revenue quality (1)
Spread	1pt - Balanced	4pts - Moderate	5pts - Unbalanced	3pts - Moderate	4pts - Moderate

Source: PitchBook • As of February 28, 2026

Across the five AIBQ dimensions, Databricks shows the most uniform strength (with no dimension scoring below an 8 out of 10), while Anthropic demonstrates a distinctive high-revenue-quality profile with improving capital efficiency at 7 out of 10 (up from 5 out of 10 at year-end 2025, as its \$14 billion in ARR outpaced the capital-base expansion)—a shift from the prior high-revenue-quality/low-computing-independence characterization, driven by its Series G nearly doubling its cumulative capital raised to \$60.55 billion. Anthropic's governance optionality now represents its second-weakest dimension alongside capital efficiency, a profile shift driven not by structural deterioration but by the political constraints imposed by the federal use ban. OpenAI and xAI both show moderate gaps between their strongest and weakest dimensions (five points for xAI, and two points for OpenAI), indicating unbalanced business models. OpenAI's capital efficiency score now ties its governance optionality as its weakest dimension following the \$110 billion round, compressing its relatively flat profile into one with a clearer low-efficiency, low-governance drag. xAI's profile has become more extreme following the SpaceX acquisition: Its computing independence remains best in class (9/10), but its governance optionality is constrained at 5/10, reflecting Elon Musk's single-person control and the elimination of independent transaction

Frontier Five AIBQ score overlay



Source: PitchBook • As of February 28, 2026

optionality. SSI scores below 4 on every dimension except governance optionality, where its score of 5 reflects a clean corporate structure despite its pre-revenue status.

AIBQ and valuation

The AIBQ framework reveals a counterintuitive pattern: The companies closest to the AI frontier score lowest on business fundamentals, while those with diversified, enterprise-focused models score highest. The clearest trend: AIBQ scores correlate inversely with how “pure-play AI lab” a company is. The company with the highest AIBQ score (Databricks) is not primarily an AI model company—it is a data infrastructure platform that added AI capabilities. The companies building

the most advanced frontier models (OpenAI, xAI, and SSI) score lowest on business quality. This suggests that AIBQ implicitly measures the distance from a pure research lab to a sustainable business. The events of February 27 sharpen this pattern: OpenAI’s \$110 billion round compressed capital efficiency, while Anthropic’s political retaliation constrained its governance optionality. Databricks, the least “AI lab” of the cohort, was the only company that remained unaffected. For investors, this creates tension between betting on capability leadership versus business durability.


AI lab purity versus business quality trade-off

Company	AIBQ score	Post-money valuation (\$B)	AI lab purity	Business characteristics and interpretation
Databricks	8.7/10	\$134.0	Lowest	Data platform first, AI second; enterprise SaaS model; multicloud; IPO-ready structure. Highest-quality, most attractive risk-adjusted return. Capital efficiency ratio (0.16x) now leads the cohort.
Anthropic	7.4/10	\$380.0 ²⁷	Medium	AI lab but enterprise obsessed (80%); B2B distribution; capital efficiency improving (0.23x, recovered from a brief 0.15x compression at the moment of Series G close, and up from 0.29x on the pre-Series G capital base). Strong quality; 300,000+ paying customers validates strategy; Claude Opus 4.6 extends capability lead. Governance optionality reduced by federal use ban despite Pentagon accepting identical terms from OpenAI hours later. AIBQ score declines due to political constraint on strategic flexibility, not governance quality deterioration.
xAI	5.4/10	\$250.0	High	Frontier lab with single-person risk; SpaceX acquisition in February 2026 makes xAI an operating subsidiary with constrained independent governance optionality. Postmerger integration risk; valuation reflects SpaceX parent ecosystem optionality. Signed deal with Pentagon on February 23 with no use restrictions, gaining access alongside OpenAI while Claude is phased out.
OpenAI	4.8/10	\$840.0	High	Frontier lab with consumer exposure; hyperscaler locked; complex governance. Highest valuation relative to AIBQ score—now \$175B per AIBQ point. \$110 billion round compressed capital efficiency from 0.31x to 0.11x, driving the capital efficiency score down. Signed Pentagon deal with restrictions mirroring Anthropic’s red lines. Valuation reflects distribution scale, brand premium, and \$174B in cumulative capital.
SSI	2.3/10	\$32.0	Highest	Pure research; pre-revenue; optionality pricing only. Venture-only position; binary outcome.

Sources: PitchBook, [CNBC](#), [TechCrunch](#) • As of February 28, 2026

27: “Anthropic Signs Term Sheet for \$10 Billion Funding Round at \$350 Billion Valuation,” CNBC, Ashley Capoot and Kate Rooney, January 7, 2026.

AIBQ score versus valuation by company



The above visual shows the AIBQ-valuation paradox. Plotting AIBQ scores against valuations reveals the central finding of this report: The relationship is inverse. OpenAI (\$840 billion valuation, AIBQ 4.8) and xAI (\$250 billion valuation, AIBQ 5.4) occupy the left half: high valuation, low quality. Databricks (\$134 billion valuation, AIBQ 8.7) and Anthropic (\$380 billion valuation, AIBQ 7.4) sit in the right half: stronger fundamentals. Against a hypothetical “fair-value” line where valuation scales with business quality, Databricks is the only company trading at an AIBQ-justified valuation. OpenAI has moved further from the fair-value line than any other company in the cohort: Its valuation increased 68% while its AIBQ score declined 8%. Anthropic’s position has shifted since year-end 2025: Its completed \$380 billion round pushes it further up the valuation axis, and the improved capital efficiency to 0.23x moves it further right, widening the quality gap with OpenAI

(4.8) and xAI (5.4). However, the federal use ban introduced a governance constraint that pulled Anthropic modestly leftward, narrowing its quality lead over the lab cohort. Anthropic now sits closer to Databricks with its AIBQ score (1.3 points higher) than to the AI labs it is typically grouped with (2 to 2.6 points higher). SSI (\$32 billion valuation, AIBQ 2.3) sits at the bottom of both axes, priced on optionality rather than fundamentals. Capital raised (bubble size) shows that the most funded companies are not necessarily the highest-quality businesses: OpenAI’s bubble, now reflecting \$173.92 billion in cumulative capital, is by far the largest in the cohort, yet the company has the second-lowest AIBQ score. For investors seeking sustainable returns rather than venture-style binary outcomes, the framework suggests overweighting Databricks and Anthropic while treating OpenAI, xAI, and SSI as position-sized optionality bets.

AIBQ weight sensitivity analysis

The AIBQ composite scores above use our base-case weights (capital efficiency 20%, revenue quality 25%, computing independence 15%, governance optionality 20%, and moat durability 20%). Institutional investors may reasonably assign different weights based on their investment thesis. The table below shows how scores shift under alternative weightings.

The top-two-ranking (Databricks #1, Anthropic #2) and bottom-ranking (SSI #5) companies are stable across all weight scenarios tested. The #3 and #4 ordering has consolidated: xAI now ranks above OpenAI in all four scenarios tested (base case, computing heavy, revenue focused, and equal weight),

a shift from before their raise, when OpenAI ranked #3 under revenue-focused weighting. The change is driven by OpenAI's capital efficiency score declining (now matching xAI's), which eliminated the efficiency advantage that previously offset xAI's computing independence edge under revenue-heavy weights. OpenAI's sole remaining dimensional advantage over xAI is revenue quality (5 versus 4), which is insufficient to overcome xAI's computing independence lead (9 versus 5) under any weighting tested. Under equal weighting, xAI (5.6) separates meaningfully from OpenAI (4.8)—a 0.8-point gap driven by xAI's best-in-class computing score (9/10)—a lead that is no longer offset by OpenAI's compressed efficiency profile. This stability validates the framework's core insight: The valuation-quality paradox persists regardless of weighting assumptions.

AIBQ composite scores by weighting

Company	Base case (20/25/15/20/20)	Computing heavy (15/20/25/20/20)	Revenue focused (15/35/10/20/20)	Equal weight (20/20/20/20/20)
Databricks	8.7/10	8.6/10	8.7/10	8.6/10
Anthropic	7.4/10	7.1/10	7.7/10	7.2/10
xAI	5.4/10	5.9/10	5.1/10	5.6/10
OpenAI	4.8/10	4.9/10	4.9/10	4.8/10
SSI	2.3/10	2.5/10	2.2/10	2.4/10

Source: PitchBook • As of February 28, 2026

Part II: Market context

Market sizing

The AI market defies traditional sizing methodologies due to its horizontal applicability. We segment the market into three layers with distinct competitive dynamics and value-capture potential.

Foundation models show the highest growth but also the highest commoditization risk. Pricing per million tokens has declined over 90% since 2023.²⁸ Sustainable value capture will shift to AI infrastructure (where Databricks dominates) and AI applications (where none of the Frontier Five has a

Foundation models show the highest growth but also the highest commoditization risk.


meaningful presence yet). This suggests Databricks has the most defensible market position among the Frontier Five. The Pentagon crisis adds a further dimension to commoditization risk in foundation models: If procurement decisions are politically mediated rather than capability driven, pricing power erodes even for technically superior products.

AI market sizes

Market layer	2025 size (\$B)	2028E size (\$B)	CAGR	Key players	Frontier Five revenue exposure
Foundation models	\$50.0	\$250.0	71%	OpenAI, Anthropic, Google, xAI	High (4 of 5 companies)
AI infrastructure	\$120.0	\$400.0	49%	Databricks, Snowflake, cloud providers	Moderate (Databricks only)
AI applications	\$30.0	\$150.0	71%	Vertical-specific enterprise tools	Low (none currently)

Sources: PitchBook, Gartner, IDC, Menlo Ventures, Andreessen Horowitz • As of February 28, 2026
Note: AI market sizing estimates are based on internal analysis triangulating Gartner, IDC, Andreessen Horowitz, and PitchBook research.

AI total addressable market (TAM) (\$B) by layer



Source: PitchBook • As of February 28, 2026

28: Large language model (LLM) inference cost has been declining roughly 10x per year for equivalent performance. "Welcome to LLMflation – LLM Inference Cost Is Going Down Fast," Andreessen Horowitz, Guido Appenzeller, November 12, 2024. See also: "LLM Inference Prices Have Fallen Rapidly But Unequally Across Tasks," Epoch AI, Ben Cottier, et al., March 12, 2025.

Capital formation

The Frontier Five have absorbed over \$312 billion in cumulative funding, representing more than close to 2x the entire US venture market in 2009. Capital deployment accelerated dramatically from 2024 to 2026, with significant raises across the five companies in that period. OpenAI's \$110 billion round on February 27, 2026, the largest single private funding event in history, accounts for more than a third of total Frontier Five capital and reshapes the cohort's concentration dynamics.

The dominance of hyperscaler and sovereign capital creates implicit strings. Amazon's \$50 billion commitment to OpenAI's February 27 round—structured as \$15 billion up front and \$35 billion contingent on AGI achievement or an IPO by year-end—

represents the single largest hyperscaler in a frontier lab to date and was paired with a \$100 billion expansion of the existing \$38 billion AWS relationship, bringing total AWS commitments to approximately \$138 billion. The deal includes exclusive third-party cloud distribution for OpenAI's Frontier enterprise platform and custom model development for Amazon products, establishing Amazon as OpenAI's second major hyperscaler dependency alongside Microsoft. Microsoft holds a roughly 27% equity stake in the restructured OpenAI Group PBC (valued at around \$135 billion), down from its previous 32.5% equity stake and 49% profit-sharing rights under the 2019 agreement.²⁹ The October 2025 restructuring freed OpenAI to pursue broader partnerships while Microsoft retained IP rights through 2032 and a \$250 billion Azure consumption commitment.³⁰ In November 2025, Microsoft committed \$5 billion in Anthropic alongside NVIDIA's \$10 billion,³¹

Total capital raised by company

Company	Total raised (\$B)	Share of Frontier Five total	2023-2026 raises	Debt facilities	Key investors
OpenAI	\$173.92	56%	• \$6.6B late-stage VC • \$40B Series F • \$8.4M secondary • \$6.5M late-stage VC • \$110B round	\$4B credit	Amazon, NVIDIA, SoftBank, Microsoft, Thrive Capital, Fidelity
Anthropic	\$60.55	19%	• \$0.98B Series B • \$0.45B Series C • \$2.75B Series D • \$3.5B Series E • \$13B Series F • \$30B Series G	\$2.5B credit	Coatue Management, Dragoneer Investment Group, Founders Fund, ICONIQ Capital, GIC, Google, Amazon
xAI	\$42.16	13%	• \$0.14B seed • \$6B Series B • \$6B Series C • ~\$15B Series D • \$20B Series E	\$5B facility	NVIDIA, Qatar Investment Authority, Fidelity, Valor Equity Partners
	\$33.11	11%	• \$0.5B Series I • \$8.6B Series J • \$1B Series K • \$78 Series I (\$5B equity + \$2B debt)	\$7.5B debt (\$5.5B facility + \$2B term loan)	Insight Partners, J.P. Morgan, Fidelity, Saraswati Ventures, BlackRock
SSI	\$3.0	1%	• \$1B seed • \$2B Series A	None	Greenoaks, Andreessen Horowitz, Sequoia Capital, Google
Total	\$312.74	100%	\$233.1B	\$19B	


Source: PitchBook • As of February 17, 2026

²⁹: "OpenAI Secures Freedom to Dilute Its Investors," Reuters, Karen Kwok, October 29, 2025.

³⁰: "Microsoft, OpenAI Reach Deal Removing Fundraising Constraints for ChatGPT Maker," Reuters, Deepa Seetharaman, Stephen Nellis, and Deborah Mary Sophia, October 28, 2025.

³¹: "Microsoft, NVIDIA and Anthropic Announce Strategic Partnerships," Microsoft, November 18, 2025.

Frontier Five funding timeline



Frontier Five capital source segmentation

Investor type	Share of Frontier Five capital	Primary targets	Investment thesis
Hyperscalers (Microsoft, Google, Amazon)	~40%	OpenAI, Anthropic, Databricks	Distribution lock-in, cloud revenue
Sovereign wealth (GIC, Qatar Investment Authority, Abu Dhabi Investment Authority)	~20%	xAI (via SpaceX), OpenAI, Anthropic, Databricks	National AI capability, diversification
Growth equity (Thrive Capital, Coatue Management, ICONIQ Capital)	~15%	All five	Pre-IPO positioning
Strategic (NVIDIA)	~15%	xAI, OpenAI, Anthropic, SSI	Computing demand visibility, circular financing
Traditional VC (Andreessen Horowitz, Sequoia Capital)	~5%	SSI, Anthropic	Founder relationships, option value

Source: PitchBook • As of February 28, 2026

which marked the first time hyperscalers had significant positions in competing frontier labs, signaling strategic hedging. Microsoft further deepened its Anthropic position by participating in its \$30 billion Series G, alongside NVIDIA, Goldman Sachs, and Temasek, reinforcing the cross-portfolio hedging pattern.³² Additionally, Amazon has committed up to \$8 billion to Anthropic with associated cloud-consumption requirements. Google holds positions in both Anthropic (more than \$3 billion invested and a roughly 14% equity stake) and SSI.³³ NVIDIA has emerged as a cross-portfolio investor in xAI, OpenAI, Anthropic, and SSI.

The “independent” AI lab remains largely fictional, but the constraint mechanism has shifted from exclusive lock-in to circular dependency. Anthropic now depends on AWS, Google tensor processing units (TPUs), and Azure simultaneously. OpenAI gained nominal freedom but committed to an estimated \$250 billion in Azure purchases and now about \$138 billion in AWS commitments—a total hyperscaler obligation approaching \$390 billion, more than double the company’s cumulative capital raised.³⁴ Capital does not just constrain—it creates competitive entanglement where hyperscalers hedge by backing multiple labs.

32: Participants in Anthropic’s Series G included Microsoft, NVIDIA, Goldman Sachs Alternative Energy Investing Group, Temasek Holdings, BlackRock, Fidelity Management & Research Company, and more than 40 additional investors alongside lead investors Coatue Management, Dragoneer Investment Group, and Founders Fund.

33: “Google Has Given Anthropic More Funding Than Previously Known, Show New Filings,” TechCrunch, Connie Loizos, March 11, 2025.

34: The Microsoft-OpenAI partnership terms are according to Securities and Exchange Commission (SEC) filings and the October 2025 PBC restructuring. “OpenAI Completes Restructure, Solidifying Microsoft as a Major Shareholder,” CNBC, Ashley Capoot, October 28, 2025. See also: “OpenAI Completes Major Reorganization With \$135 Billion Microsoft Stake,” TIME, Chantelle Lee, October 28, 2025 and “Microsoft Starts Talking About OpenAI, But Disclosures Still Fall Far Short of Complete,” Deep Quarry, Olga Usvyatsky and Francine McKenna, November 16, 2025.

Capital efficiency analysis

Before analyzing multiples, investors should understand capital efficiency: the revenue generated per dollar raised. This metric differentiates business models from business projects. We calculate capital efficiency using the following formula:

$$\text{Capital efficiency} = \text{ARR} / \text{cumulative capital raised}$$

The Frontier Five now display three distinct efficiency tiers. OpenAI's \$110 billion round on February 27 collapsed its ratio from 0.31x to 0.11x, a 65% compression that moved it from cohort leader to the second-lowest tier, above only SSI. Databricks (0.16x) and Anthropic (0.23x) now occupy the top two positions, with Anthropic maintaining a clear lead on monetization velocity and Databricks maintaining its distinction as the only company generating positive FCF. Anthropic has separated from Databricks at 0.23x versus 0.16x, reversing the brief convergence that appeared immediately following the \$30 billion Series G close when Anthropic's point-in-time ratio dropped to 0.15x on the lagging \$9 billion ARR figure. The rapid recovery from 0.15x to 0.23x in the weeks since the Series G reflects Anthropic's underlying revenue velocity: ARR grew from \$9 billion to \$14 billion between late 2025 and mid-February 2026—a \$5 billion acceleration that was already in motion before the capital raise closed. For investors, the post-round tier structure tells a revised story: Anthropic leads on monetization velocity, Databricks leads on

capital discipline, and OpenAI's scale advantage in absolute revenue (\$20 billion) is now overwhelmed by the denominator effect of \$174 billion in cumulative capital.

OpenAI's 0.11x multiple means it generates \$0.11 in ARR for every dollar raised—a ratio that would have ranked last among the Frontier Five when excluding SSI just weeks ago. The \$110 billion round added more capital in a single transaction than OpenAI had raised in its entire prior history (\$63.9 billion), while ARR remained at \$20 billion. The prior 0.31x ratio reflected both massive revenue and the compounding effect of H2 2025 growth; the new ratio reflects the denominator shock of the largest private funding round ever closed. OpenAI projects \$280 billion in revenue by 2030 and \$600 billion in total computing spending over the same period. If 2030 revenue materializes, efficiency would recover to approximately 1.6x, but the intervening years require investors to accept a ratio below xAI's infrastructure-first profile. Anthropic (0.23x) presents a more nuanced picture than a single ratio suggests. At the time of the \$30 billion Series G close on February 12, 2026, the company's ARR produced a 0.15x efficiency that nearly halved the 0.29x ratio that prevailed on the pre-Series G capital base of \$30.55 billion. However, this snapshot was misleading: Anthropic's ARR had been accelerating rapidly (\$1 billion to \$4 billion, to \$9 billion to \$14 billion over the prior 14 months), and by the time the round settled, the February ARR had already reached \$14 billion—restoring efficiency to 0.23x. The 0.29x-to-0.23x compression (a roughly 21% decline) is a


Frontier Five capital efficiency

Company	ARR (\$B)	Capital raised (\$B)	Efficiency	Interpretation
OpenAI	\$20.0	\$173.92	0.11x	Collapsed from 0.31x (prior cohort leader) following \$110B round. Denominator shock: Capital raised increased 172% while ARR was unchanged. Now trails both Anthropic and Databricks. Efficiency recovery dependent on projected \$280B revenue by 2030.
Anthropic	\$14.0	\$60.55	0.23x	Cohort leader on monetization velocity. Efficiency compressed from 0.29x at time of Series G close (\$9B ARR on \$30.55B capital) but rapidly recovering as ARR scales on the enlarged capital base. FCF negative but margin trajectory improving. Federal use ban introduces ~\$200M government revenue risk (~1.4% of ARR), but enterprise concentration limits direct financial impact.
Databricks	\$5.4	\$33.1	0.16x	Cohort leader on capital discipline. FCF positive; proven unit economics despite ratio compression from \$7B Series L and \$5.45B debt facility.
xAI	\$3.2	\$42.1	0.08x	Infrastructure heavy; pre-monetization.
SSI	\$0.0	\$3.0	0.00x	Pre-revenue by design.

Sources: PitchBook, [Reuters](#), [OpenAI](#) • As of February 28, 2026

Note: ARR-based efficiency reflects forward momentum; trailing 12-month (TTM)-based efficiency (shown in earlier sections) reflects realized revenue. Both metrics are valid for different analytical purposes. See the ["ARR versus actual revenue"](#) section below.

Frontier Five capital efficiency



Source: PitchBook • As of February 28, 2026

fundamentally different story than the 0.29x-to-0.15x halving (a roughly 48% decline) that the previous \$9 billion figure implied. Critically, Anthropic’s underlying monetization velocity has not deteriorated: The 80% enterprise revenue concentration with minimal consumer infrastructure remains intact, Claude Code alone generates \$2.5 billion in ARR (about 18% of the total), and more than 500 customers now spend over \$1 million annually. The federal use ban introduces revenue risk on the government side (the Pentagon contract represented roughly \$200 million, or about 1.4% of ARR), but Anthropic’s enterprise concentration limits the direct financial impact. Investors should monitor whether Anthropic’s projected \$20 billion to \$30 billion ARR target for 2026 materializes—achieving \$25 billion would restore efficiency to around 0.41x on the current capital base, vindicating the raise. Databricks occupies a distinct tier above OpenAI but remains the only company to achieve positive free cash flow (FCF)—a milestone no frontier lab has reached, and arguably a more meaningful indicator of business sustainability. Databricks’ ratio compressed from 0.18x following its \$7 billion Series L and \$5.5 billion debt facility, which pushed its cumulative capital raised from \$26.1 billion to \$33.1 billion. Unlike Anthropic’s compression, Databricks’ was planned: The capital supports IPO-readiness infrastructure and international expansion, with FCF already proving the model. Databricks’ February 2026 press release confirmed YoY growth of more than 65% and more than \$1.4 billion in AI product revenue, demonstrating that its capital discipline coexists with accelerating growth and is not at its expense. xAI’s 0.08x ratio reveals a deliberate infrastructure-first strategy where Colossus build-out (a cash burn of about \$1 billion/month) significantly outpaces revenue materialization—a bet that computing advantage will compound into future market share. OpenAI’s 0.11x now sits closer to xAI’s infrastructure-first ratio than to Anthropic’s enterprise-monetization ratio—a

convergence that would have been unthinkable before the round and that underscores how a single capital event can reshape cohort positioning.

Capital efficiency now displays a restructured hierarchy following OpenAI’s \$110 billion round. The prior three-tier structure, with OpenAI leading decisively at 0.31x before the round, has inverted. Anthropic (0.23x) now leads on monetization velocity, Databricks (0.16x) leads on capital discipline and remains uniquely FCF positive, and OpenAI (0.11x) has compressed into the lower tier alongside xAI (0.08x). The separation between Anthropic and Databricks masks complementary strengths: Databricks is uniquely FCF positive, indicating its capital generates self-sustaining economics rather than just growth; Anthropic demonstrates the fastest ARR velocity in the cohort, converting capital to revenue at a pace without precedent in B2B software. xAI (0.08x) remains deep in an infrastructure-investment phase, while SSI (0.00x) has no revenue against \$3 billion in capital raised—the widest capital-to-revenue gap in the cohort. Valuation-to-ARR multiples have diverged sharply: OpenAI now trades at 42x (\$840 billion/\$20 billion), close to double the Anthropic (27.1x) and Databricks (24.8x) multiples that have converged. Before the round, OpenAI’s multiple had compressed to 25x—the \$110 billion round reversed that entirely, producing the widest multiple gap in the cohort between the two highest-valued companies. xAI’s 78.1x remains an outlier reflecting SpaceX ecosystem optionality rather than standalone economics. The visual tells the story: Anthropic has separated from the pack on efficiency velocity, Databricks occupies a disciplined second tier with proven unit economics, OpenAI’s denominator shock has pulled it into proximity with xAI, and xAI is in a different phase entirely.

ARR- versus TTM-revenue-based efficiency ratios by company

Company	ARR-based efficiency	TTM-revenue-based efficiency	Delta
Anthropic	0.23x	~0.12x	-0.11x
Databricks	0.16x	~0.14x	-0.02x
OpenAI	0.11x	~0.07x	-0.04x
xAI	0.08x	0.08x	0.00x

Source: PitchBook • As of February 28, 2026

ARR versus actual revenue

Companies in hypergrowth mode often report ARR, which annualizes current monthly revenue for a full-year projection. This creates a critical distinction for investors: ARR captures a company's most recent momentum—what it would earn if today's run rate held for 12 months—while TTM revenue reflects what the company actually collected over the past year. When a company is accelerating rapidly, ARR can far exceed TTM revenue, making the business appear more mature than its cash flows suggest. The gap between the two figures reveals how much of a company's reported scale is forward-looking projection versus demonstrated economics. Among the Frontier Five, these gaps vary dramatically.

Investors should distinguish between momentum (ARR) and realized economics (TTM revenue) when assessing valuation sustainability. This report uses ARR for both capital efficiency calculations and valuation multiples, providing a forward-looking view of each company's monetization trajectory. Investors preferring a more conservative lens should apply TTM revenue, which would produce the following adjusted capital efficiency ratios:

On a TTM basis, the tier structure has compressed and reordered. Databricks (~0.14x) and Anthropic (~0.12x) remain within striking distance of each other at the top, while OpenAI (~0.07x) has fallen below xAI (0.08x)—a reversal that underscores the denominator shock of the \$110 billion round. The Anthropic-Databricks gap narrows from seven basis points (ARR: 0.23x versus 0.16x) to around two basis points (TTM: ~0.12x versus ~0.14x), with Databricks actually leading on TTM revenue—reflecting Databricks' mature revenue profile (~1.2x ARR/TTM ratio) versus Anthropic's steep but unproven growth curve (~2.0x ratio). The ~2.0x ARR/TTM ratio for Anthropic is the highest in the cohort, meaning its ARR-based metrics

are the most forward-loaded and carry the greatest execution risk. xAI's 1.0x ratio—indicating ARR equals TTM—reflects the relative stability of X Corp ecosystem revenue rather than hypergrowth acceleration. OpenAI's TTM-based ratio of ~0.07x places it last among revenue-generating companies in the cohort on realized economics—a position that would have been inconceivable before the round and that highlights how ARR-based metrics (0.11x) mask the gap between projected momentum and collected revenue.

Competitive dynamics

The dispersion in multiples across the Frontier Five has widened dramatically. OpenAI's \$110 billion round pushed its multiple from 25x to 42x, breaking the convergence with Databricks (24.8x) that had emerged in late 2025. The cohort now splits into three valuation tiers: "valued as businesses" (Anthropic at 27.1x and Databricks at 24.8x), "valued as narrative" (OpenAI at 42x), and "valued as bets" (xAI at 78.1x and SSI at infinity)—placing investors across fundamentally incompatible pricing frameworks.

Anthropic's multiple compression from 42x (at \$9 billion in ARR) to 27.1x (at \$14 billion in ARR) is a significant shift. At 27.1x, Anthropic's valuation premium over Databricks (24.8x) narrows to just two turns—a far cry from the 14-turn gap that prevailed under the \$9 billion figure. This 27.1x multiple, while still premium to enterprise SaaS comparables, is within the range that growth-stage software companies with over 100% revenue growth have historically commanded. The 107% valuation jump in five months (\$183 billion in September 2025 to \$380 billion in February 2026) reflects both revenue acceleration and FOMO-driven multiple expansion ahead of a potential 2026 IPO. The completed round resolves the uncertainty that surrounded the January term sheet—but at \$60.55 billion in cumulative capital raised, the 27.1x revenue

multiple coexists with a 0.23x capital efficiency ratio. The federal use ban introduces IPO timing uncertainty but has not affected the private market valuation: Anthropic's Series G closed at \$380 billion on February 12, and secondary market pricing has held steady through the crisis, consistent with Dario Amodei's statement that revenue and valuation have grown during the standoff. Investors should still monitor whether the \$20 billion to \$30 billion 2026 ARR target materializes to further compress the multiple toward Databricks' range.

The convergence of OpenAI and Databricks at 25x that prevailed before the \$110 billion round has been decisively broken. OpenAI now trades at 42x—a 69% premium to Databricks' 24.8x—after its valuation increased from \$500 billion to \$840 billion while ARR remained at \$20 billion. The prior parity implied the market assigned roughly equivalent forward growth expectations to both companies despite fundamentally different business models and risk profiles. The \$110 billion round repriced OpenAI on a different framework entirely: At 42x on \$20 billion in ARR, investors are underwriting revenue reaching approximately \$100 billion within three to four years to justify the valuation on a growth-adjusted basis—a target consistent with OpenAI's own projection of \$280 billion by 2030 but requiring sustained 80%+ annual growth from an already massive base. For Databricks, 25x on \$5.4 billion in ARR requires over \$13 billion on the same timeline—a more achievable target given 65% current growth. The multiple divergence reinforces the AIBQ paradox: The market now pays a 69% multiple premium for OpenAI (AIBQ 4.8) over Databricks (AIBQ 8.7).

xAI's 78.1x multiple mirrors that of early SpaceX, which traded at 50x-100x revenue during its infrastructure build-out before Starlink monetization compressed multiples to around 15x. The comparison is no longer hypothetical. SpaceX acquired xAI at \$250 billion, making xAI a SpaceX operating subsidiary. The critical difference: SpaceX's moat was physical (reusable rockets competitors could not replicate), while xAI's is computational (GPUs that hyperscalers can match given sufficient capital). At a \$250 billion acquisition valuation with around \$500 million in standalone revenue (separate from X) and a \$1 billion/month cash burn, xAI must reach around \$5 billion in revenue by 2028 to justify current pricing—a 10x increase from standalone guidance. The Pentagon classified-systems deal signed on February 23 adds a government revenue stream but on terms ("all lawful purposes" with no restrictions) that tie revenue to political alignment rather than product-market fit—a dependency the AIBQ framework penalizes under revenue quality. The embedded bet: Colossus is Starlink (transformative) rather than The Boring Company (aspirational). For investors, xAI exposure is now accessible only through SpaceX equity, where xAI represents around 20% of combined value—diluting the purity of the AI thesis but anchoring it to SpaceX's proven \$7.5 billion to \$8 billion EBITDA.

SSI is compared to DeepMind because both are pure research labs valued on team quality. In 2014, DeepMind was acquired for an estimated \$650 million with zero revenue. SSI's \$32 billion valuation implies the market believes Ilya Sutskever is worth 49x what Demis Hassabis was or that AI talent has appreciated 49x in a decade. Both interpretations have merit.

Frontier Five valuation/ARR multiples

Company	Post-money valuation (\$B)	ARR (\$B)	Valuation/ARR	Implied category	Comparable range
OpenAI	\$840.0	\$20.0	42.0x	Consumer/enterprise	Meta (7x-8x), Google (5x-6x)
Anthropic	\$380.0	\$14.0	27.1x	Enterprise AI	ServiceNow (10x), Salesforce (6x-7x)
xAI	\$250.0	\$3.2	78.1x	Infrastructure bet	Early SpaceX
Databricks	\$134.0	\$5.4	24.8x	Enterprise SaaS	Snowflake (16x), MongoDB (12x)
SSI	\$32.0	\$0	∞	Pure research	DeepMind acquisition (around \$500M-\$650M)

Sources: PitchBook, [Reuters](#), [OpenAI](#), [TechCrunch](#) • As of February 28, 2026

Note: Post-money valuations are as of February 28, 2026. xAI's valuation reflects the SpaceX acquisition price (February 2, 2026); Anthropic's valuation reflects its completed Series G (February 12, 2026). OpenAI's valuation reflects its \$110 billion round closed on February 27, 2026.

Capability comparison: Benchmark performance

Benchmark performance has converged dramatically. The top three frontier models are within 5% on most metrics, fundamentally changing competitive dynamics. Anthropic’s Claude Opus 4.6 reopened the capability gap on key benchmarks at launch on February 5, but Google’s Gemini 3.1 Pro (released on February 19) reclaimed the lead on reasoning—most notably on ARC-AGI-2, where it scored 77.1% versus Opus 4.6’s 68.8%. The benchmark race now features

three distinct leaders across different capability domains rather than a single front-runner.

When the top models are within 5% on most benchmarks, capability is no longer a sustainable moat. The playing field has shifted to distribution, pricing, and enterprise trust—favoring Anthropic and OpenAI over xAI and SSI. This dynamic reinforces the AIBQ framework’s moat durability dimension, where Anthropic and Databricks both score 8/10 on durable competitive advantages beyond raw model capability, while xAI

Frontier AI model benchmark scores

Benchmark	OpenAI (GPT-5.2)	Anthropic (Claude Opus 4.6)	Google (Gemini 3.1 Pro)	xAI (Grok 4.1)
MMLU (knowledge)	91.5%	92.0%	92.6%	87.0%
HumanEval (code)	92.0%	94.0%	~92.0%	87.0%
GPQA Diamond (science)	~78.0%	91.3%	94.3%	~68.0%
SWE-bench Verified	75.0%	80.8%	80.6%	65.0%
ARC-AGI-2	54.2%	68.8%	77.1%	40.0%
Terminal-Bench 2.0	60.0%	65.4%	68.5%	50.0%
OSWorld (agentic)	65.0%	72.7%	~65.0%	55.0%
Arena Elo rating	1,440	1,506	1,500	1,475 ³⁵
Context window	400K	1M (beta)	1M	2M

Sources: GitHub, Artificial Analysis, [ARC Prize](#), [Terminal-Bench](#), [OSWorld](#), [Arena](#), [SWE-bench](#), [Anthropic](#) • As of February 28, 2026

Note: Arena Text Leaderboard results are as of late February 2026. Anthropic benchmark data is according to the company’s February 5, 2026, announcement. Google benchmark data is according to the Gemini 3.1 Pro model card released on February 19, 2026. Gemini 3.1 Pro is in preview; scores may shift at general availability. Arena Elo reflects the Claude Opus 4.6 thinking variant (1,506). HumanEval is largely saturated and less differentiating than SWE-bench for frontier models. GPQA Diamond has replaced MATH in this table as a more discriminating measure of scientific reasoning at the frontier. GPT-5.3-Codex has appeared on Chatbot Arena under anonymous code names but has not been officially released; its reported benchmarks (77.3% Terminal-Bench 2.0, 56.8% SWE-bench Pro) are excluded pending official disclosure.

Frontier AI model benchmark score heatmap

Model	MMLU-Pro	GPQA Diamond	HumanEval+	SWE-bench Verified	ARC-AGI-2	Creative writing	Agentic tasks
Gemini 3.1 Pro	91	94	92	81	77	80	90
Claude Opus 4.6	92	91	95	81	69	97	94
GPT-5.2	91	72	93	75	54	82	85
Grok-4.1	88	68	90	65	40	75	78
DeepSeek V3.2	90	72	91	55	42	70	75
Llama 4 Maverick	87	65	91	45	35	72	74

■ Top performer ■ Above average □ Average ■ Below average ■ Lowest

Sources: GitHub, Artificial Analysis, [ARC Prize](#), [Terminal-Bench](#), [OSWorld](#), [Arena](#), [SWE-bench](#), [Anthropic](#) • As of February 28, 2026

35: The cited score is for the Grok 4.1 thinking variant; the nonthinking variant scores around 1,466 Elo.

(6/10) and SSI (2/10) remain disproportionately dependent on computing and talent, respectively. Notably, Anthropic's Claude Opus 4.5 achieved 80.9% on SWE-bench Verified on November 24, 2025,³⁶ which explains the rapid adoption of Claude Code.³⁷ OpenAI's GPT-5.2, released in December 2025, narrowed the gap but remained approximately six percentage points behind (75% versus 80.8%), making real-world factors such as developer experience and integration more decisive than raw benchmark performance.

The reasoning benchmark race has accelerated since early February. Claude Opus 4.6's ARC-AGI-2 score of 68.8%—14.6 percentage points ahead of GPT-5.2 (54.2%)—represented a significant lead at its launch on February 5. However, Google's Gemini 3.1 Pro, released on February 19, reclaimed the ARC-AGI-2 lead at 77.1%, surpassing Opus 4.6 by 8.3 percentage points and more than doubling Gemini 3 Pro's prior score of 31.1%. Gemini 3.1 Pro also matched Opus 4.6 on SWE-bench Verified (80.6% versus 80.8%) and surpassed it on Terminal-Bench 2.0 (68.5% versus 65.4%) and GPQA Diamond (94.3% versus 91.3%). ARC-AGI-2 tests reasoning on problems that are easy for humans but hard for AI, making the rapid lead changes on this benchmark a meaningful signal of how quickly capability advantages erode at the frontier. For investors evaluating Anthropic's \$380 billion post-money valuation at 27x ARR, Claude Opus 4.6 retains benchmark leadership in enterprise-critical domains—SWE-bench Verified (80.8%, highest), OSWorld agentic tasks (72.7%, highest), GDPval-AA expert work tasks (1,606 Elo, highest), and creative writing—while Gemini 3.1 Pro leads on abstract reasoning and scientific knowledge. This domain-specific fragmentation provides a capability justification for the multiple—the 27x multiple at a \$14 billion ARR falls within the range that 100%+ growth companies have historically commanded, making the capability-to-valuation bridge more credible. The question is whether capability translates to revenue conversion fast enough to further compress the multiple toward Databricks' 25x before the next funding cycle. Additional capabilities

of Claude Opus 4.6 include agent teams for multiagent coordination; a 1-million-token context window (beta); 128,000 output tokens (a 2x increase from Opus 4.5); Microsoft PowerPoint integrations; adaptive thinking; and 500+ zero-day vulnerabilities found.

The Arena leaderboard has shifted twice since early February 2026. At launch, Claude Opus 4.6 took the top two positions: the thinking variant at 1,506 Elo and the nonthinking variant at 1,503—both ahead of Gemini 3 Pro (#3 at ~1,487), Grok 4.1 thinking (#4 at 1,475), and Gemini 3 Flash (#5 at ~1,473). Gemini 3.1 Pro, released on February 19, has entered the Arena and is expected to challenge for the top position as votes accumulate, given its benchmark dominance on 13 of 16 evaluated metrics. GPT-5.2 (~1,440) has fallen outside the top five entirely, and GPT-5.3-Codex has appeared under anonymous code names ("vortex" and "zephyr") but has not been officially released. This multimodel churn undercuts the "capability convergence" narrative that dominated early 2026—not because one provider has pulled away, but because leadership now rotates across providers on a weekly cadence, making any single benchmark lead ephemeral. Anthropic retains the strongest position in enterprise-applicable benchmarks (SWE-bench at 80.8%, OSWorld at 72.7%, GDPval-AA at 1,606 Elo) while Google leads on abstract reasoning (ARC-AGI-2 at 77.1%) and scientific knowledge (GPQA Diamond at 94.3%)—a combination that reinforces the AIBQ framework's emphasis on moat durability over capability snapshots. Grok 4.1's #4 Arena Elo placement is notable given xAI's 5.4 AIBQ score: The model is competitive even as the business scores poorly on governance and capital efficiency. The Pentagon's classified-systems deal with xAI—despite Grok trailing the ARC-AGI-2 leader by 37.1 percentage points—further demonstrates that political alignment, not capability, now drives government procurement decisions. Post-acquisition, Grok's development trajectory is now subject to SpaceX's capital-allocation priorities rather than independent AI lab incentives.

³⁶: "Official Leaderboards," SWE-bench, n.d., accessed February 19, 2026.

³⁷: "Anthropic Releases Opus 4.5 With New Chrome and Excel Integrations," TechCrunch, Russell Brandon, November 24, 2025.

³⁸: "ARC-AGI-2 Leaderboard," ARC Prize, n.d., accessed February 5, 2026.

Pricing analysis: Consumer versus enterprise

Pricing strategy has emerged as a critical battleground in frontier AI. The past 24 months have witnessed a decline in per-token costs of over 90%, driven primarily by DeepSeek’s aggressive market entry and subsequent competitive responses from Western providers. Understanding the distinction between consumer and enterprise pricing models is essential for evaluating revenue quality and margin sustainability. This analysis directly informs the AIBQ revenue quality dimension (25% weight), where the consumer/enterprise revenue mix is a primary scoring input.

Consumer subscription models

Consumer pricing follows a tiered subscription model standardized around free (limited access), pro/plus (\$20/month), and premium/max (\$100-\$200/month) subscriptions.

All major providers have converged on around \$20/month for the primary consumer tier, creating intense value competition at this price point. The premium tier shows meaningful variation—with OpenAI at \$200, Anthropic at \$200, xAI at \$300, and Google at \$250—reflecting different positioning strategies. Despite offering free access, xAI charges 50% more than competitors at the SuperGrok tier (\$30 versus \$20), justified by real-time X data integration. Google’s context window of over 1 million tokens at the \$20/month tier represents significant value differentiation versus competitors’ 128,000-to-400,000-token limits.

Consumer subscription models by company

Company	Free tier	Pro/plus tier	Premium/max tier	Key differentiator
OpenAI	GPT-5.2 mini (limited)	\$20/month (Plus)	\$200/month (Pro)	Unlimited o1, Sora video, operator preview
Anthropic	Claude Sonnet (limited)	\$20/month (Pro)	\$100-\$200/month (Max)	5x-20x usage, extended context, priority
xAI	Free at grok.com	\$30/month (SuperGrok)	\$300/month (SuperGrok Heavy)	Real-time X data; X Premium+ (\$40/month) bundles platform
Google	Gemini Flash (limited)	\$20/month (AI Pro)	\$250/month (AI Ultra)	1M+ context window; Workspace integration; Veo 3 video

Sources: OpenAI, Anthropic, xAI, Google • As of February 28, 2026

39: "pricing-details-usd," DeepSeek, n.d., accessed February 19, 2026.

Understanding the distinction between consumer and enterprise pricing models is essential for evaluating revenue quality and margin sustainability.

The over 90% API cost reduction since 2024 was catalyzed by DeepSeek’s January 2025 entry at prices 90% below competitors (\$0.55 input/\$2.19 output per million tokens),³⁹ forcing OpenAI and others to respond with aggressive cuts. Consumer subscription prices have remained relatively stable, with the API price war not yet fully translating to consumer tiers. OpenAI has begun testing advertising in its Free and Go tiers—a monetization lever that could improve consumer-tier economics but risks diluting brand positioning among premium users. For Anthropic, this stability matters: With an 80% enterprise revenue concentration and over 300,000 business accounts, consumer subscription is a secondary revenue driver. Anthropic’s business subscriptions quadrupled between January 1 and mid-February 2026, and Claude Code’s \$2.5 billion ARR—with enterprise users generating more than half of revenue—demonstrates that Anthropic’s pricing power is concentrated in developer and enterprise workflows rather than consumer subscriptions. For OpenAI, where consumer subscriptions represent around 60% of revenue, the lack of consumer price increases constrains a major growth lever—a dynamic reflected in its revenue quality score of 5/10 versus Anthropic’s 9/10. OpenAI’s \$110 billion round, which pushed its ARR multiple from 25x to 42x, increases the pressure on consumer monetization: At 42x, the market is pricing in revenue acceleration that consumer subscription stability alone cannot deliver.

Enterprise pricing models

Enterprise pricing diverges significantly from consumer models, combining per-seat subscriptions with usage-based billing. The key distinction is the emphasis on security, compliance, and administrative controls rather than raw model access. Enterprise deals typically involve annual commitments, volume discounts, and dedicated support. The Pentagon AI procurement landscape shifted dramatically in late February 2026. Anthropic's approximately \$200 million defense contract is being terminated following the presidential executive order banning all federal use of its technology. xAI signed a classified-systems deal on February 23 with no use restrictions, and OpenAI signed its own classified-systems deal on February 27 with guardrails against mass surveillance and autonomous weapons—terms functionally identical to those Anthropic had requested. Google retains its existing Pentagon contracts. The net result: Three of four major providers now hold classified-systems access, while the provider with the highest-scoring model on enterprise-critical benchmarks has been excluded on political rather than capability grounds. Post-acquisition, xAI's enterprise and government sales pipeline is now subject to SpaceX's existing government contracting relationships—a potential advantage given SpaceX's deep ties with the US Department of Defense, but also a potential conflict

given Elon Musk's political profile and the Committee on Foreign Investment in the United States' (CFIUS's) sensitivities around xAI's original investor base (including sovereign wealth funds).

The enterprise pricing table highlights a critical AIBQ insight: Databricks' consumption-based model (DBU) with one-to-three-year commitments is fundamentally different from the per-seat models used by the AI labs. This structure generates an NRR over 140% through natural usage expansion within existing accounts—a pattern validated across enterprise SaaS over two decades—while per-seat models must continually expand head count to grow within accounts. Databricks' pricing model is a primary driver of its 9/10 revenue quality score and a key reason its 0.16x capital efficiency ratio now leads the cohort: Both Databricks and Anthropic reflect rapid capital deployment ahead of full revenue realization, but the revenue Databricks generates is structurally stickier and more predictable. For Anthropic, the path to justifying its \$380 billion valuation involves converting its over 300,000 business accounts from team-tier seats to enterprise-tier annual commitments with usage-based expansion—a transition that would move its revenue closer to Databricks' quality profile and support the 27.1x multiple that the market currently assigns. The federal use ban creates a paradoxical

Enterprise pricing models by company

Company	Product	Team tier	Enterprise	Enterprise features
OpenAI	ChatGPT Business	\$25-\$30/user/month	~\$60/user/month (150+ seats)	Single sign-on (SSO), System for Cross-Domain Identity Management (SCIM), 128K-token context window, System and Organization Controls 2, no training on data. Pentagon classified-systems access with stated prohibitions on mass surveillance and autonomous weapons; forward-deployed engineers embedded at Pentagon.
Anthropic	Claude for Teams	\$25-\$30/user/month	Custom (~\$60/seat, 70+ minimum, annual)	500K-to-1M-token context window, role-based access control, audit logs, SCIM, compliance API, Claude Code (premium seats). Federal use banned by executive order; ~\$200M defense contract being terminated; supply chain risk designation pending legal challenge.
xAI	Grok Business/Enterprise/Enterprise Vault	Custom (self-serve)	Custom	SSO, SCIM, customer-controlled encryption with Enterprise Vault; Pentagon classified-systems access under "all lawful purposes" with no use restrictions; government pricing at \$0.42/agency via the US General Services Administration.
Databricks	Unity Catalog + AI	Consumption (DBU)	One-to-three-year commitments	Data governance, machine learning (ML) lifecycle, multicloud, Unity Catalog included with Premium/Enterprise. No government AI contract exposure; entirely insulated from Pentagon procurement dynamics.

Sources: OpenAI, Anthropic, xAI, Databricks • As of February 28, 2026

dynamic: Commercial enterprise customers who value principled AI governance may accelerate adoption precisely because Anthropic demonstrated willingness to forfeit government revenue rather than compromise on surveillance and autonomous weapons restrictions—strengthening the enterprise trust moat that underpins its 9/10 revenue quality score. Anthropic disclosed that eight of the Fortune 10 are Claude customers, and revenue and valuation have continued to grow through the standoff. The \$200 million contract loss represents approximately 1.4% of ARR—material but not existential for a company with \$14 billion in enterprise-concentrated revenue.

API pricing: The margin battleground

API pricing reveals true foundation model economics, bifurcated into flagship reasoning models and “fast” inference models. xAI’s Grok 4.1 Fast at \$0.20/\$0.50 per million tokens (input/output) represents the most aggressive pricing in the market—a discount of more than 90% versus comparable models—while flagship Grok 4 prices at market parity.

The strategic divergence is clear: OpenAI’s GPT-5.2 at \$1.75/\$14 represents a 40% price increase over GPT-5.1

(\$1.25/\$10), reflecting the computing intensity of its advanced reasoning capabilities. GPT-5.1 remains available as a volume-tier option. Even at the higher GPT-5.2 price point, OpenAI’s models remain significantly cheaper than Anthropic’s Claude Opus 4.6 (\$5/\$25): a 65% discount on input and 44% on output. At its price point, Anthropic’s Claude Opus 4.6 maintains premium positioning while delivering significantly improved capabilities (such as leadership on SWE-bench Verified at 80.8%, OSWorld at 72.7%, and GDPval-AA at 1,606 Elo, and a 1-million-token context window), expanding the value proposition without raising prices. This represents a 67% price reduction from the legacy Claude Opus 4/4.1 (\$15/\$75), democratizing frontier intelligence while preserving premium positioning. The pricing gap between OpenAI and Anthropic’s flagships has narrowed: GPT-5.2’s \$14 output price is now within 1.8x of Claude Opus 4.6’s \$25, compared with the 2.5x gap at the GPT-5.1 price point. Google’s Gemini 3.1 Pro—which now leads on ARC-AGI-2 (77.1%) and 13 of 16 tracked benchmarks—maintains the same \$2/\$12 pricing as Gemini 3 Pro, making it the most competitively priced frontier-class model on a capability-per-dollar basis. At less than half the cost of Opus 4.6 with broadly comparable or superior benchmark performance, Gemini 3.1 Pro intensifies pricing pressure across the entire flagship tier. Anthropic’s

Flagship AI model API pricing

Company	Flagship model	Input/million tokens	Output/million tokens	Estimated margin	Strategy
OpenAI	GPT-5.2 ⁴⁰	\$1.75	\$14	~50%-55%	Flagship reasoning
	GPT-5.1	\$1.25	\$10	~55%-60%	Volume/legacy
Anthropic	Claude Opus 4.6 ⁴¹	\$5	\$25	~40%-47% ⁴²	Premium reasoning
	Claude Sonnet 4.6 ⁴³	\$3	\$15	~45%-55%	Enterprise workhorse
xAI	Grok 4	\$3	\$15	~40%-50%	Market parity
	Grok 4.1 Fast	\$0.20	\$0.50	~15%-25%	Loss leader/land grab
Google	Gemini 3.1 Pro ⁴⁴	\$2	\$12	~60%	Cloud bundle

Sources: [Price Per Token](#), [OpenAI](#), [Anthropic](#), [xAI](#), [Google](#), [The Information](#) • As of February 28, 2026

Note: Margin estimates represent per-model inference margins (revenue minus serving computing costs), not company-wide gross profit margins. Company-level blended gross margins are lower: OpenAI’s internal computing margin reached ~70% on paid products by October 2025, but GAAP gross margin is estimated at ~48%; Anthropic projected a 40% company-wide gross margin for 2025, down from an earlier 50% expectation, reflecting the computing intensity of scaling frontier inference. Google reduced Gemini serving unit costs by 78% over 2025 through TPU optimization, with the Google Cloud segment operating margin reaching 30.1% in Q4 2025. xAI reported a \$63 million gross profit on \$107 million Q3 2025 revenue (~59% blended gross margin), though the company posted a net loss of \$1.46 billion in the same quarter, burning \$7.8 billion in cash over the first nine months of 2025.

40: OpenAI’s GPT-5.2 was released on December 10, 2025, at \$1.75/\$14 per million tokens (input/output)—a 40% increase over GPT-5.1 (\$1.25/\$10), released on November 13, 2025. GPT-5.2 is OpenAI’s current flagship with a 400,000-token context window. “[API Pricing](#),” OpenAI, n.d., accessed February 19, 2026. See also: “[GPT 5.2 API Pricing 2026](#),” [Price Per Token](#), February 19, 2026.

41: Anthropic’s Claude Opus 4.6 API pricing is unchanged from that of Opus 4.5 at \$5/\$25 per million tokens (input/output), according to Anthropic’s February 5, 2026, announcement. “[Introducing Claude Opus 4.6](#),” Anthropic, February 5, 2026. See also: “[Pricing](#),” Anthropic, n.d., accessed February 19, 2026.

42: “[Anthropic Lowers Gross Margin Projection as Revenue Skyrockets](#),” [The Information](#), Sri Muppidi, January 22, 2026.

43: Claude Sonnet 4.6 was released on February 17, 2026, replacing Sonnet 4.5 as the default model across all Claude plans. The pricing was unchanged at \$3/\$15 per million tokens (input/output). Anthropic describes it as its “most capable Sonnet model yet” with improvements across coding, computer use, long-context reasoning, and agent planning. Sonnet 4.6 features a 1-million-token context window in beta. Users preferred Sonnet 4.6 over Opus 4.5 59% of the time in Claude Code testing. “[Introducing Claude Sonnet 4.6](#),” Anthropic, February 17, 2026.

44: Google Gemini 3.1 Pro is priced at \$2/\$12 per million tokens (input/output) for contexts of 200,000 or fewer tokens and \$4/\$18 for contexts over 200,000 tokens. “[Gemini Developer API Pricing](#),” Gemini API, n.d., accessed February 19, 2026.

margin trajectory from -94% to 40%, with projections of 63% by 2027 and 77% by 2028, represents the most dramatic margin expansion path in the Frontier Five. The January 2026 downward revision from 50% to 40% signals that scaling inference costs continue to outpace efficiency gains at the current growth rate.

Anthropic's model lineup received a significant update on February 17, 2026, with the release of Claude Sonnet 4.6, which replaced Sonnet 4.5 as the default model across all Claude plans at the same \$3/\$15 pricing. Anthropic describes Claude Sonnet 4.6 as approaching Opus-level intelligence at the Sonnet price point—a claim supported by early customer evaluations showing users preferred Sonnet 4.6 over Opus 4.5 59% of the time in Claude Code.⁴⁵ Sonnet 4.6 achieved 72.5% on OSWorld (versus Opus 4.6's 72.7%) and surpassed Opus 4.6 on GDPval-AA by 27 Elo points, demonstrating near-parity on enterprise-critical tasks. For the pricing analysis, this is significant: If Sonnet 4.6 genuinely closes the capability gap with Opus, the \$5/\$25 Opus tier faces cannibalization risk from within Anthropic's own lineup. The counterargument—and Anthropic's stated position—is that Opus 4.6 remains superior for tasks demanding the deepest reasoning, multiagent coordination, and maximum reliability. For enterprise customers evaluating cost per quality-adjusted token, Sonnet 4.6 is delivering around 90% of Opus capability at \$3/\$15, representing a compelling value proposition that could accelerate enterprise adoption and support the revenue acceleration Anthropic needs to justify its \$380 billion valuation.

The middle tier is now remarkably competitive. Claude Sonnet 4.6 (\$3/\$15), Grok 4 (\$3/\$15), and Google Gemini 3.1 Pro (\$2/\$12) all occupy a narrow pricing band, with capability

differentiation, rather than price, as the primary battleground. Gemini 3.1 Pro's benchmark dominance at the lowest price point in this tier creates significant pressure: Developers can access the ARC-AGI-2 leader and GPQA Diamond leader at 33% below Sonnet pricing, a value proposition that could erode Anthropic's mid-tier share if enterprise-critical benchmark gaps (SWE-bench, GDPval-AA, OSWorld) continue to narrow. OpenAI's GPT-5.1 (\$1.25/\$10) undercuts all three, while GPT-5.2 (\$1.75/\$14) sits between the mid-tier cluster and Anthropic's premium Opus pricing. This creates a clear market structure: budget tier (Grok 4.1 Fast at \$0.20/\$0.50), middle tier (\$1.25-\$3 input), and premium tier (Claude Opus 4.6 at \$5/\$25).

xAI's dual-tier strategy warrants scrutiny. Grok 4 at \$3/\$15 competes at parity with Claude Sonnet 4.6, but Grok 4.1 Fast at \$0.20/\$0.50 is the true loss leader. Investors should question whether xAI can sustain these prices and determine whether training and inference costs have dropped dramatically or whether the pricing likely represents subsidized usage to build market share. With estimated fully loaded costs of \$0.40 to \$0.60 per million tokens, xAI likely operates at break-even or a slight loss on high-volume inference. This is a deliberate land-grab strategy. The question: Can Colossus efficiency eventually make Grok 4.1 Fast profitable at scale, or is it a perpetual subsidy to capture developer mindshare? Since the acquisition, this question now sits on SpaceX's balance sheet—and SpaceX's \$7.5 to \$8 billion EBITDA provides substantially more runway to sustain below-cost pricing than xAI's standalone roughly \$1 billion/month burn rate allowed. The Pentagon classified-systems deal adds a government revenue stream for xAI, but on "all lawful purposes" terms that prioritize political access over commercial pricing discipline—consistent with the land-grab pattern visible in its API pricing strategy.

Pricing strategy decoded

xAI’s pricing strategy is nuanced. Flagship Grok 4 at \$3/\$15 per million tokens competes at market parity with Claude Sonnet 4.6, while Grok 4.1 Fast at \$0.20/\$0.50 operates over 90% below market as a deliberate loss leader.⁴⁶ This can be interpreted in two ways: 1) Colossus computing efficiency enables sustainable low pricing on the fast tier, or 2) this is a loss-leader strategy to buy market share before unit economics matter. Given xAI’s 0.08x capital efficiency, the second interpretation is more likely. The fast-tier pricing is not sustainable without continued capital injection or significant efficiency gains. With the SpaceX acquisition, xAI’s pricing decisions now flow through SpaceX’s capital-allocation framework rather than standalone fundraising, providing more runway but less pricing autonomy.

Value capture is shifting from raw inference to applications, agents, and enterprise workflows, explaining why Claude Code

(now at \$2.5 billion ARR, up from around \$400 million in mid-2025)⁴⁷ and Databricks’ AI products represent more durable streams than pure API access. Anthropic’s announcement of agent teams and PowerPoint integration in Claude Opus 4.6 signals strategic expansion into enterprise workflow tools. The release of Claude Sonnet 4.6 further accelerates this shift, making frontier-quality AI accessible for high-volume enterprise workflows that were previously cost prohibitive at Opus pricing. Google’s Gemini 3.1 Pro at \$2/\$12—with benchmark leadership on 13 of 16 metrics—represents the most significant pricing pressure on the middle tier, compressing the value proposition for both Claude Sonnet 4.6 and Grok 4 at their \$3/\$15 price points. For Anthropic, the strategic response is clear: Differentiate on enterprise-specific value (compliance, trust, and Claude Code integration) rather than benchmark-per-dollar metrics where Google now holds the advantage.

AI pricing strategy by company

Company	Pricing strategy	Sustainable?	Investment implication
Anthropic	Premium pricing, value capture	Yes	Claude Opus 4.6 delivers more capability at same price; Claude Sonnet 4.6 approaches Opus level at \$3/\$15, accelerating enterprise adoption; margins support research & development reinvestment. Federal use ban removes ~\$200M government revenue but reinforces enterprise trust positioning that supports premium pricing power.
OpenAI	Market pricing, volume play	Uncertain	Amazon and Microsoft partnership provides flexibility; GPT-5.2 at \$1.75/\$14 captures flagship reasoning tier. \$110B round pushes multiple to 42x, requiring aggressive revenue acceleration that stable consumer pricing cannot deliver alone. Ad testing in free/Go tiers signals search for incremental monetization.
xAI	Dual-tier strategy: Flagship parity and aggressive fast-tier discounting	Uncertain	SpaceX acquisition provides capital runway, but Grok development now subject to SpaceX’s capital-allocation priorities. Pentagon classified-systems deal adds government revenue on “all lawful purposes” terms, consistent with market-share-first pricing philosophy.

Source: PitchBook • As of February 28, 2026

46: Grok 4.1 Fast’s margin was estimated as a loss leader based on computing costs relative to pricing. “Welcome to xAI Documentation,” xAI, n.d., accessed February 19, 2026. See also: “Price Per Token,” Price Per Token, February 19, 2026.
47: “Anthropic,” Sacra, n.d., accessed February 19, 2026.

Distribution advantage

With capability near parity, distribution determines winners. Here is how each of the Frontier Five accesses customers:

OpenAI’s distribution is unmatched (over 800 million weekly ChatGPT users plus Microsoft’s enterprise footprint), but distribution does not equal monetization.⁴⁸ OpenAI had over \$20 billion in ARR by year-end 2025, with PitchBook estimating \$35 billion by Q4 2026.⁴⁹ Consumer subscription revenue is estimated at \$5 billion to \$8 billion, with enterprise representing the faster-growing and higher-margin segment. OpenAI’s Codex has reached 1.6 million weekly users—tripling since the start of the year—with 9 million paying business users, demonstrating momentum in developer distribution. The Amazon partnership announced alongside the \$110 billion round adds exclusive third-party cloud distribution for

OpenAI’s Frontier enterprise platform across AWS, significantly expanding enterprise reach beyond the existing Microsoft channel. The conversion gap reveals the opportunity: Whoever cracks enterprise conversion at consumer scale will lead the next decade. xAI has over 600 million X users but generates \$3.2 billion in revenue (including the X ecosystem).⁵⁰ xAI’s Pentagon classified-systems deal and OpenAI’s classified-systems agreement both add government and defense as distribution channels—a domain where xAI’s political alignment and OpenAI’s scale provide access, while Anthropic has been excluded by executive order. Anthropic’s over 300,000 paying business customers demonstrate enterprise conversion at scale without consumer distribution dependency. Anthropic’s federal use ban closes one distribution channel (government/defense) while the enterprise trust generated by Anthropic’s principled stance may strengthen commercial channels.

Frontier Five distribution advantages

Company	Primary distribution	Reach	Stickiness	Advantage score
OpenAI	ChatGPT (~800M-900M+ weekly active users), Codex (1.6M weekly, 9M paying business users), Microsoft365, Azure, AWS (Frontier enterprise platform), Pentagon classified systems	Massive	Medium	5/5
Anthropic	Claude.ai, Claude Code (\$2.5B ARR, 54% coding share), AWS Bedrock, Google Cloud Vertex AI, enterprise direct. Federal government channel closed by executive order	Large	High	4/5
Databricks	Existing 20K+ enterprise customers	Focused	Very high	4/5
xAI	X platform (~600M+ users), Grok integration, Pentagon classified systems	Massive	Low	3/5
SSI	None (no product)	Zero	N/A	1/5

Sources: [TechCrunch](#), X, Databricks, The Information, Reuters • As of February 28, 2026
Note: The X user base is according to company disclosures; the xAI/Grok integration was announced in 2024. The Databricks customer count is according to a company press release in February 2026. OpenAI’s 2025 TTM revenue of ~\$13 billion is according to The Information (November 2025); its \$20 billion ARR run rate is according to Reuters (December 2025); the consumer/enterprise split is estimated. PitchBook estimates a \$35 billion ARR by Q4 2026.

48: “Sam Altman Says ChatGPT Has Hit 800M Weekly Active Users,” [TechCrunch](#), Rebecca Bellan, October 6, 2025.
49: “OpenAI to Focus on ‘Practical Adoption’ in 2026, Says Finance Chief Sarah Friar,” [CNBC](#), Salvador Rodriguez, January 19, 2026. See also: “OpenAI’s Annual Recurring Revenue Tripled to \$20 Billion in 2025,” [PYMNTS](#), January 19, 2026.
50: “X has 600 million monthly active users...” X, [Elon Musk \(@elonmusk\)](#), May 23, 2024.

Moat durability assessment

Databricks has the most durable moat: over 20,000 enterprise customers (60% of the Fortune 500), with deep data integration creating genuine switching costs.⁵¹ Databricks is also the only Frontier Five company entirely insulated from the Pentagon crisis—no government AI contract exposure, no political risk, no involvement in the military AI debate—reinforcing its position as the most traditionally investable company in the cohort. Anthropic’s enterprise trust moat (Constitutional AI, safety-first positioning) is underappreciated—the company holds 54% market share in enterprise AI coding tools versus OpenAI’s 21%.⁵² Anthropic’s moat is strengthening: Claude Opus 4.6’s benchmark leadership on enterprise-critical metrics (SWE-bench Verified at 80.8%, OSWorld at 72.7%, GDPval-AA at 1,606 Elo), over 300,000 paying business customers, and over 500 customers at \$1 million+ ARR make the “enterprise AI” thesis increasingly compelling. The Anthropic-Pentagon standoff creates a paradoxical moat dynamic: The federal use ban closes the government channel, but Anthropic’s willingness to forfeit around \$200 million in defense revenue rather than compromise on surveillance and autonomous weapons

restrictions may deepen enterprise trust among commercial customers who value principled governance—particularly given that the Pentagon accepted functionally identical terms from OpenAI hours later, validating the substance of Anthropic’s position. OpenAI’s moat is real but increasingly complex: Its \$110 billion round provides unmatched capital firepower, the Amazon partnership adds a second major cloud distribution channel, the Pentagon classified-systems deal adds government access, and Codex’s 1.6 million weekly users demonstrate developer traction. But the moat faces structural pressure from OpenAI’s 42x multiple (requiring sustained 80%+ growth to justify); governance complexity spanning Microsoft, Amazon, and nonprofit-to-PBC conversion obligations; and a capital efficiency ratio (0.11x) that trails the cohort. xAI’s only durable advantage is computing independence via Colossus; its Pentagon classified-systems access was secured through political alignment rather than capability superiority—the Pentagon acknowledged replacing Claude would be “very difficult”—making it a politically contingent rather than a structurally durable moat. SSI has no moat until it ships a product.

Frontier Five moat durability assessment

Moat type	OpenAI	Anthropic	xAI	SSI	Databricks
Brand/mindshare	5/5	3/5	2/5	1/5	3/5
Distribution	5/5	3/5	3/5	1/5	4/5
Capability lead	4/5	4/5	2/5	1/5	3/5
Data flywheel	4/5	3/5	4/5	1/5	5/5
Enterprise trust	3/5	5/5	2/5	1/5	5/5
Computing independence	2/5	2/5	5/5	1/5	4/5
Overall moat	Strong	Very strong	Moderate	Weak	Very strong

Source: PitchBook • As of February 28, 2026

Note: Scoring methodology: Brand/mindshare scoring was based on consumer recognition surveys, media volume, and developer community engagement. Distribution scoring was based on multiplying the addressable user base by the conversion potential. Capability lead scoring was based on the domain-weighted benchmark composite (SWE-bench, ARC-AGI-2, and multimodal evaluations). The data flywheel scoring was based on multiplying the proprietary data volume by the feedback loop maturity. The enterprise trust scoring was based on compliance certifications, safety positioning, and enterprise market share according to Menlo Ventures. The computing independence scoring was based on the percentage of self-owned infrastructure versus cloud infrastructure. Government/defense distribution is included in distribution scoring where applicable. Capability lead scoring has been updated to reflect Gemini 3.1 Pro’s February 19 benchmark results, which shifted ARC-AGI-2 leadership from Anthropic to Google.

51: "Databricks Grows >65% YoY, Surpasses \$5.4 Billion Revenue Run-Rate, Doubles Down on Lakebase and Genie," Databricks, February 9, 2026.
52: "2025 Mid-Year LLM Market Update: Foundation Model Landscape + Economics," Menlo Ventures, Tim Tully, et al., July 31, 2025. Note: Menlo Ventures is an Anthropic investor; the report’s data is based on a survey of over 150 technical leaders. PM Insights references this data secondarily. "Anthropic Approaches \$7B Run Rate in 2025, Outpaces OpenAI," PM Insights, November 10, 2025.

Assessment summary: Who is winning?

The release of Claude Opus 4.6 and the subsequent release of Gemini 3.1 Pro have shifted the competitive landscape twice in three weeks. Capability convergence—the dominant narrative of early 2026—is now replaced by domain fragmentation: Google leads on abstract reasoning (ARC-AGI-2 at 77.1%) and scientific knowledge (GPQA Diamond at 94.3%),⁵³ Anthropic leads on enterprise-critical metrics (SWE-bench Verified at 80.8%, OSWorld at 72.7%, GDPval-AA at 1,606 Elo),⁵⁴ and OpenAI retains the broadest distribution. Combined with the 1-million-token context window, agent teams, \$2.5 billion Claude Code ARR, and over 300,000 paying business customers, Anthropic leads on several AIBQ-weighted metrics: enterprise trust, monetization velocity, and pricing power. Benchmark leadership now rotates across providers on a weekly cadence, reinforcing the AIBQ framework’s emphasis on moat durability over capability snapshots. The release of Sonnet 4.6 compounds Anthropic’s enterprise advantage, as Anthropic now controls both the premium and mid-tier segments of the API market—a pricing architecture no competitor currently matches. However, Google’s Gemini 3.1 Pro at \$2/\$12 with benchmark leadership on 13 of 16 metrics creates significant pressure on the middle tier, where Sonnet 4.6’s \$3/\$15 must now justify a 50% premium over a broadly superior model on headline benchmarks.

Capital efficiency tells a fundamentally different story. OpenAI’s \$110 billion round collapsed its ratio from 0.31x to approximately 0.11x, inverting the cohort hierarchy. Anthropic now leads on ARR-based efficiency at 0.23x, followed by Databricks at 0.16x—but with largely different quality characteristics. Databricks’ 0.16x coexists with positive FCF and an NRR over 140%, reflecting mature, proven unit economics. Anthropic’s 0.23x reflects both explosive revenue growth (\$1 billion to \$14 billion in ARR in 14 months) and the denominator shock of the \$30 billion Series G: capital raised ahead of deployment, not capital destruction. OpenAI’s 0.11x reflects the most extreme denominator shock in venture history: \$110 billion in new capital against unchanged revenue, compressing efficiency below Databricks’ and into proximity with xAI’s infrastructure-first 0.08x ratio. Anthropic’s gross margins reached an estimated 40% in 2025 (revised downward from an earlier 50% projection in January 2026), with a road map to 63% by 2027 and 77% by 2028, and cash-flow-positive operations expected by 2027—approximately three years ahead of OpenAI, which projects cumulative operating losses of around \$74 billion in 2028 and profitability by 2029 or 2030. OpenAI projects \$280 billion in revenue by 2030 and \$600 billion in total computing spending over the same period—projections that, if achieved, would restore efficiency to approximately 1.6x but require sustained 80%+ annual growth from a \$20 billion base.

Frontier Five competitive assessment summary

Dimension	Leader	Runner-up	Laggard
Capability	Google/Anthropic (domain-dependent)	OpenAI	xAI
Distribution	OpenAI	Anthropic	SSI
Pricing power	Anthropic	Google	xAI
Enterprise trust	Databricks	Anthropic	xAI
Capital efficiency	Anthropic (ratio)/Databricks (FCF)	Databricks (ratio)	OpenAI
Moat durability	Databricks	Anthropic	SSI

Source: PitchBook • As of February 27, 2026

53: "ARC-AGI-2 Leaderboard," ARC Prize, n.d., accessed February 19, 2026.
54: "Text Arena," Arena, February 18, 2026.

Anthropic's \$14 billion ARR as of February 2026, up from \$1 billion just 14 months earlier, represents the fastest B2B revenue ramp-up in software history. The company's internal projections target \$20 billion to \$30 billion in ARR for full-year 2026, with an optimistic scenario reaching \$70 billion in revenue and \$17 billion in cash flow by 2028. The federal use ban is immaterial in dollar terms (around \$200 million, or 1.4% of ARR) but carries cascading risk through the supply chain risk designation, which Anthropic is legally challenging. Commercial enterprise metrics—business subscriptions quadrupling since January 1, over 500 customers at \$1 million+ ARR, eight of the Fortune 10—suggest the pipeline is accelerating independently of the government channel.

OpenAI projects \$29.4 billion in revenue for 2026 and has set an ambitious \$100 billion target for 2027—though at OpenAI's 0.11x capital efficiency and 42x valuation multiple, the Amazon partnership and Pentagon deal expand distribution without addressing the core challenge: Revenue must grow faster than any enterprise software company in history to compress the multiple toward the Anthropic-Databricks range.⁵⁵ xAI is buying market share at unsustainable prices (0.08x capital efficiency). Its Pentagon deal validates the AIBQ revenue quality concern: Political alignment, not product quality, drives the revenue stream. SSI remains a pure option bet with no competitive position until product launch.

55: "OpenAI's Revenue Could Reach \$100 Billion in 2027, Altman Suggests," The Information, Laura Mandaro, November 3, 2025.

Part IIIA: Company profiles: The Frontier Five

Each profile follows an identical structure to enable direct comparison, with sections on the company's investment thesis, trade, business mechanics, financial snapshot, AIBQ analysis, bull/bear cases, and actionable triggers. Companies are presented in order of AIBQ score, from highest to lowest.

Databricks: The enterprise foundation

AIBQ score: 8.7/10

Investment thesis

Databricks is a data infrastructure company with AI capabilities, not an AI lab. While frontier AI labs burn billions of dollars pursuing capability leadership with uncertain economics, Databricks generates positive FCF, with an NRR over 140%, gross margins exceeding 80%, and more than 800 customers paying \$1 million+ annually.^{56, 57} The AI opportunity (over \$1.4 billion in ARR, around 26% of total revenue and accelerating) is genuine upside to an already strong core business. Databricks is the only Frontier Five company entirely unaffected by the Pentagon crisis that reshaped the cohort in late February 2026: no government AI contract exposure, no political risk, and no supply chain designation vulnerability. Its capital efficiency ratio of 0.16x now leads the cohort following OpenAI's ratio compression, reinforcing the misclassification thesis.

The trade

At 24.8x ARR, the valuation is aggressive but defensible given the company's growth (65%+ and accelerating), margins (80%+), FCF trajectory (positive), and quality (140%+ NRR). Databricks is the only Frontier Five company valued using traditional enterprise software frameworks without requiring execution-risk haircuts. Its IPO catalyst provides near-term liquidity, and its AI products provide long-term upside. The 24.8x multiple now represents a 41% discount to OpenAI's 42x—a gap that has widened from near-parity before the \$110 billion round—despite Databricks scoring 3.9 AIBQ points higher.

Position	Core holding: Anchor allocation in AI portfolio
Entry point	Current round (\$134B) or IPO allocation (H2 2026 likely)
Exit trigger	Hold through IPO and lockup; exit at 35x+ forward
Position size	Max 8% of AI allocation; lowest risk supports larger position
Conviction	High: Only Frontier Five company where traditional frameworks apply

Source: PitchBook • As of February 27, 2026

Business mechanics

Databricks provides a unified data analytics platform combining data engineering, data science, and ML. The company pioneered the “data lakehouse” architecture, which features the flexibility of data lakes with the reliability of data warehouses. Databricks' core products include Delta Lake (storage); MLflow (ML lifecycle); Unity Catalog (governance); and, increasingly, AI-native products, including Agent Bricks



8.7/10

AIBQ score

9/10

Capital efficiency

9/10

Revenue quality

8/10

Computing independence

9/10

Governance optionality

8/10

Moat durability

56: “Databricks Grows +55% YoY, Surpasses \$4.8B Revenue Run-Rate, and Is Raising >\$4B Series L at \$134B Valuation,” Databricks, December 16, 2025.

57: “Under the Hood of the AI Economy With Databricks CEO Ali Ghodsi,” CNBC, Deirdre Bosa, February 9, 2026.

DATABRICKS: THE ENTERPRISE FOUNDATION

(multi-agent systems), Lakebase (a serverless Postgres database for AI agents), and Genie (a conversational AI assistant for enterprise data).

The consumption-based enterprise model creates natural expansion: Customers pay for computing and storage used, growing spending as data volumes increase. The revenue split is approximately 74% platform (data engineering and structured query language analytics) and 26% AI products (model serving,

vector search, and Agent Bricks). AI products reached over \$1.4 billion in ARR as of Q4 2025, faster than any product line in company history, with Lakebase adoption at twice the pace of structured query language (SQL) analytics adoption at the same stage and Genie expanding the addressable market to nontechnical users across the enterprise.

Databricks has built a layered business model where each layer reinforces the others:


Databricks business model architecture

Business layer	ARR	Growth rate	Gross margin	Customer base
Data lakehouse (foundation)	\$3.0B	~40%	~85%	Core platform; all customers
Data warehousing (SQL)	\$1.0B+	~100%+	~75%	50% of lakehouse customers
AI & ML products	\$1.4B+	~150%+	~70%	Growing share of lakehouse
Lakebase (new)	Nascent	2x SQL pace	TBD	Thousands (early)

Sources: PitchBook, [Databricks](#) • As of February 27, 2026

Note: Segment-level growth rates and gross margins are PitchBook estimates except where noted. Databricks discloses only aggregate >80% subscription gross margin and total revenue growth. Data warehousing growth is derived from company milestones (\$400 million in ARR in April 2024 to over \$1 billion in December 2025). Customer penetration percentages are estimated; at its June 2025 Data + AI Summit, the Databricks CEO disclosed that 50% of customers use more than six products.

Databricks ARR (\$B) by business layer



Sources: PitchBook, [Databricks](#) • As of February 27, 2026

DATABRICKS: THE ENTERPRISE FOUNDATION

Financial snapshot

Metric	Value	Context
Valuation	\$134B	24.8x ARR; 41% discount to OpenAI's 42x despite scoring 3.9 AIBQ points higher
ARR	\$5.4B	>65% YoY; accelerating
AI products ARR	\$1.4B+	~26% of total; fastest-growing segment
Total customers	20,000+	Including 60% of Fortune 500
Capital raised	\$33.1B	Includes equity and debt facilities
Gross margin	>80%	Software economics—highest in Frontier Five
FCF	Positive	Only FCF-positive company in Frontier Five
NRR	>140%	Best in class; exceeds Snowflake (125%)
Capital efficiency	0.16x	Now #2 in cohort; leads OpenAI (0.11x) following \$110B round
\$1M+ customers	>800	Deep enterprise penetration
\$10M+ customers	>70	New disclosure; signals enterprise concentration health
Government exposure	None	Only Frontier Five company insulated from Pentagon crisis

Sources: PitchBook, [Databricks](#) • As of February 27, 2026

58: ARR by industry are estimates based on PitchBook data and public company disclosures; they are not disclosed by Databricks.

AIBQ deep dive

Capital efficiency: 9/10

Databricks is the only Frontier Five company generating positive FCF. The 0.16x ARR/capital raised ratio (\$5.4 billion/\$33.1 billion) now ranks second in the cohort, behind Anthropic (0.23x), following OpenAI's compression from 0.31x to 0.11x after the \$110 billion round. The critical distinction: Databricks converts capital into sustainable, self-funding economics rather than growth subsidized by future rounds. Anthropic's higher 0.23x ratio reflects explosive growth (\$1 billion to \$14 billion in ARR in 14 months) but margins turning positive only in 2025; Databricks' lower ratio coexists

with gross margins over 80% and positive FCF—a fundamentally different economic quality.

The gross margins over 80% are software economics, not AI lab economics. Unlike OpenAI (40% to 45%) or xAI (15% to 25%), Databricks operates with a traditional software margin structure. A score of 9 reflects exceptional efficiency with demonstrated profitability capability. Databricks falls short of 10 only because growth investment constrains stated margins below the potential maximum.

Revenue quality: 9/10

Every quality metric for Databricks is exceptional: an NRR over 140% (negative net churn), enterprise revenue of more than 95%, common multiyear contracts, a top-10 concentration of less than 15%, and more than 800 customers at over \$1 million in ARR. The company's customer base spans industries: financial services (over \$800 million in ARR from J.P. Morgan, Goldman Sachs, and Citi), healthcare/pharma (over \$500 million from Moderna, Pfizer, and UnitedHealth), retail (over \$400 million from

Walmart, Target, and Home Depot), and technology and energy (over \$600 million from Adobe, Comcast, Shell, and others).⁵⁸ Integration depth creates genuine switching costs—migrating from Databricks requires moving data governance, access controls, and compliance frameworks. Databricks' revenue base has zero government AI contract exposure—a quality characteristic that gained significance in late February 2026 as the Pentagon crisis introduced political risk to competitors' revenue streams.

Computing independence: 8/10

Databricks has a multicloud strategy (AWS, Azure, and Google Cloud Platform) with customer choice. Unlike AI labs, Databricks' business model is computing light;

customers bring their own cloud infrastructure. No single hyperscaler can constrain the business.

Governance optionality: 9/10

Databricks is the most IPO ready in the Frontier Five. It is a standard Delaware C corporation with audit-ready financials and an experienced board. CEO Ali Ghodsi has been preparing for an IPO since 2021. An S-1 filing is likely in H1 or H2 2026, with pricing likely in H2 2026. The additional \$2 billion in debt that Databricks secured in January 2026 signals accelerated IPO

preparation. The Pentagon deal has no bearing on Databricks' IPO timeline or investor appetite—a structural advantage over Anthropic, whose IPO readiness is now complicated by the federal use ban, and OpenAI, whose nonprofit-to-PBC conversion and \$390 billion in combined hyperscaler obligations add governance complexity.

Moat durability: 8/10

Data relationships, not AI capability, create Databricks' moat. Deep integration with enterprise data stacks—governance, access controls, and compliance frameworks—creates switching costs that pure AI capability cannot replicate. AI products (\$1.4 billion in ARR, growing over 150% YoY) prove expansion capability. Recent strategic integrations with OpenAI (\$100 million multiyear spending commitment in September 2025), Anthropic (\$100 million over five years, as of March 2025), and Google Gemini 3.1 Pro (launched natively on Databricks) strengthen platform positioning as the model-agnostic enterprise

AI layer. Databricks is now the only enterprise platform with native access to every major frontier model family: OpenAI, Anthropic, Google, and Meta. Model agnosticism also insulates Databricks from the benchmark churn now defining the frontier: Whether Google leads on ARC-AGI-2 or Anthropic leads on SWE-bench, Databricks customers access both. A key concern is AI executive Naveen Rao departing in September 2025 to launch hardware startup Unconventional AI (in which Databricks is an investor); no successor to the head of AI role has been publicly named.

DATABRICKS: THE ENTERPRISE FOUNDATION

Bull and bear cases

Bull case (\$200B+ at IPO) = 40%	Bear case (\$80B-\$100B at IPO) = 20%
IPO prices at 25x+ forward ARR amid strong institutional demand	AI narrative fades; multiple compresses to Snowflake levels (14x to 16x)
AI products accelerate to 30%+ of ARR, justifying higher multiple	Growth decelerates to 35% to 40% as enterprise spending slows
Growth sustained at 50%+ through 2027, exceeding expectations	IPO market closes; liquidity delays to 2027
Only Frontier Five company insulated from political/government risk; safe-haven premium emerges	Snowflake AI competition intensifies
Hyperscaler acquisition interest emerges, providing valuation floor	Gemini 3.1 Pro's benchmark leadership at \$2/\$12 pricing commoditizes the AI model layer, reducing willingness to pay for Databricks' model-serving products
Returns: 50%+ from current round	Returns: 25%-40% downside from current round

Source: PitchBook • As of February 27, 2026

Actionable triggers

Signal	Trigger event	Context
Positive	S-1 filing shows 60%+ growth sustained	Add to full allocation
Positive	AI products exceed 30% of ARR	Upgrade moat durability score to 9/10
Positive	Pentagon deal drives safe-haven reallocation toward Databricks in institutional AI portfolios	Monitor secondary market pricing for premium expansion
Negative	IPO delayed beyond Q4 2026	Reduce position; catalyst delayed
Negative	Growth decelerates below 45%	Reduce position; multiple-compression risk
Neutral	IPO prices at 22x-25x forward	Hold through lockup; fundamentals matter more than entry price

Source: PitchBook • As of February 27, 2026



PitchBook provides actionable insights across the global capital markets.

Additional research:



SpaceX Initiation Report: When Venus and Jupiter Align

Download the report [here](#)



Q4 2025 AI VC Trends

Download the report [here](#)

PitchBook Insights is an online compendium of in-depth data, news, analysis, and perspectives that shape the private capital markets.

PitchBook subscribers enjoy exclusive access to a comprehensive suite of private market insights, including proprietary research, news, data, tools, and more on the [PitchBook platform](#).

Nizar Tarhuni

Executive Vice President of Research and Market Intelligence

Paul Condra

Senior Director, Global Head of Private Markets Research

Report created by:

Harrison Rolfes

Senior Research Analyst, Private Company Coverage

Megan Woodard

Senior Graphic Designer

Mara Potter

Principal Graphic Designer, Creative

Learn more about [PitchBook's Institutional Research team](#).

Click [here](#) for PitchBook's report methodologies.

COPYRIGHT © 2026 by PitchBook Data, Inc. All rights reserved. No part of this publication may be reproduced in any form or by any means—graphic, electronic, or mechanical, including photocopying, recording, taping, and information storage and retrieval systems—without the express written permission of PitchBook Data, Inc. Contents are based on information from sources believed to be reliable, but accuracy and completeness cannot be guaranteed. Nothing herein should be construed as any past, current or future recommendation to buy or sell any security or an offer to sell, or a solicitation of an offer to buy any security. This material does not purport to contain all of the information that a prospective investor may wish to consider and is not to be relied upon as such or used in substitution for the exercise of independent judgment.